

Sensor-Based Air Pollutant Concentration Regression

Machine Learning · Assignment #1 · Air Quality Monitoring (UCI Air Quality dataset)

Predict a ground-truth pollutant concentration from low-cost metal-oxide sensor signals and weather variables, using Linear / Polynomial regression (from scratch + scikit-learn) and Decision Tree regressors.

Stage	Contents	Status
1. Data Preparation	Load, EDA, missing values, target & feature selection, scaling, 80/20 split	✔ this section
2. Model Training	Linear & polynomial regression from scratch, sklearn cross-check, Decision Trees (<code>max_depth</code> sweep), CV tuning, ablations	✔
3. Model Evaluation	MSE, RMSE, R^2 ; metric vs. epoch and vs. tree depth	✔
4. Model Visualization	Tree plot (<code>max_depth=3</code>), feature importances	✔
5. Baseline Comparison	Linear vs. Decision Tree	✔
6. Discussion	Short written answers	✔

Stage 1 — Data Preparation

1.0 Setup

```
import warnings
warnings.filterwarnings("ignore")

import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from scipy import stats
from IPython.display import display
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from statsmodels.stats.outliers_influence import variance_inflation_factor

SEED = 42 # single seed -> reproducible split
DATA_PATH = "../data/air+quality/AirQualityUCI.csv" # relative to src/notebooks/

# Plot style: clean, recessive grid, colour-blind-safe (Okabe-Ito) palette
BLUE, ORANGE, GREEN, RED, GREY = "#0072B2", "#E69F00", "#009E73", "#D55E00", "#7F7F7F"
sns.set_theme(style="whitegrid", context="notebook")
plt.rcParams.update({"figure.dpi": 110, "axes.spines.top": False, "axes.spines.right": False,
                    "grid.alpha": 0.35, "axes.titleweight": "bold", "axes.titlesize": 12})
pd.set_option("display.max_columns", 30, "display.width", 140, "display.float_format", "{:,.3f}".format)
```

1.1 Load the data

The UCI CSV is **semicolon-separated with decimal commas**, and has two trailing empty columns and trailing blank rows (an Excel export artefact). We parse it accordingly.

```
raw = pd.read_csv(DATA_PATH, sep=";", decimal=",")
print(f"Raw file shape: {raw.shape[0]:,} rows x {raw.shape[1]} columns")
display(raw.head())
display(raw.tail(3))
```

Raw file shape: 9,471 rows x 17 columns

	Date	Time	CO(GT)	PT08.S1(CO)	NMHC(GT)	C6H6(GT)	PT08.S2(NMHC)	NOx(GT)	PT08.S3(NOx)	NO2(GT)	PT08.S4(NO2)	PT08.S5(O3)	T	RH	AH	Unnamed: 15	Unnamed: 16
0	10/03/2004	18.00.00	2.600	1,360.000	150.000	11.900	1,046.000	166.000	1,056.000	113.000	1,692.000	1,268.000	13.600	48.900	0.758	NaN	NaN
1	10/03/2004	19.00.00	2.000	1,292.000	112.000	9.400	955.000	103.000	1,174.000	92.000	1,559.000	972.000	13.300	47.700	0.726	NaN	NaN
2	10/03/2004	20.00.00	2.200	1,402.000	88.000	9.000	939.000	131.000	1,140.000	114.000	1,555.000	1,074.000	11.900	54.000	0.750	NaN	NaN
3	10/03/2004	21.00.00	2.200	1,376.000	80.000	9.200	948.000	172.000	1,092.000	122.000	1,584.000	1,203.000	11.000	60.000	0.787	NaN	NaN
4	10/03/2004	22.00.00	1.600	1,272.000	51.000	6.500	836.000	131.000	1,205.000	116.000	1,490.000	1,110.000	11.200	59.600	0.789	NaN	NaN

	Date	Time	CO(GT)	PT08.S1(CO)	NMHC(GT)	C6H6(GT)	PT08.S2(NMHC)	NOx(GT)	PT08.S3(NOx)	NO2(GT)	PT08.S4(NO2)	PT08.S5(O3)	T	RH	AH	Unnamed: 15	Unnamed: 16
9468	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN
9469	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN
9470	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN

1.2 Structural cleaning (artefacts of the file format)

```
df = raw.dropna(how="all").dropna(axis=1, how="all").copy() # drop blank rows / the 2 empty 'Unnamed' columns
print(f"After removing blank rows/columns: {df.shape[0]:,} rows x {df.shape[1]} columns "
      f"(removed {raw.shape[0]-df.shape[0]} blank rows, {raw.shape[1]-df.shape[1]} empty columns)")

# Build a proper timestamp and index the data by it (dates are dd/mm/yyyy, times use '.' separators)
df["Datetime"] = pd.to_datetime(df["Date"] + " " + df["Time"], format="%d/%m/%Y %H.%M.%S")
df = df.drop(columns=["Date", "Time"]).set_index("Datetime").sort_index()

GT_COLS = ["CO(GT)", "NMHC(GT)", "C6H6(GT)", "NOx(GT)", "NO2(GT)"] # candidate targets
SENSOR_COLS = ["PT08.S1(CO)", "PT08.S2(NMHC)", "PT08.S3(NOx)", "PT08.S4(NO2)", "PT08.S5(O3)"]
WEATHER_COLS = ["T", "RH", "AH"]
FEATURE_CANDIDATES = SENSOR_COLS + WEATHER_COLS

print("\nTime coverage :", df.index.min(), "->", df.index.max())
print("Duplicate rows:", df.duplicated().sum(), "| duplicate timestamps:", df.index.duplicated().sum())
gaps = df.index.to_series().diff().value_counts()
print("Gaps between consecutive timestamps:\n", gaps.to_string())
```

After removing blank rows/columns: 9,357 rows x 15 columns (removed 114 blank rows, 2 empty columns)

Time coverage : 2004-03-10 18:00:00 -> 2005-04-04 14:00:00

Duplicate rows: 31 | duplicate timestamps: 0

Gaps between consecutive timestamps:

```
Datetime
0 days 01:00:00    9356
```

Observations. The series is a clean, gap-free hourly grid (every step is exactly 1 h) with no duplicates. Note the data really ends in **April 2005**, slightly beyond the "March 2004 – February 2005" in the brief — it is just over one year, so every season (and weather regime) appears at least once.

1.3 Schema, dtypes and raw summary statistics

```
schema = pd.DataFrame({
    "dtype": df.dtypes.astype(str),
    "role": ["target candidate" if c in GT_COLS else "sensor feature" if c in SENSOR_COLS else "weather feature" for c in df.columns],
    "n_unique": df.nunique(),
})
display(schema)
print("RAW describe() (before treating -200 as missing):")
display(df.describe().T)
```

	dtype	role	n_unique
CO(GT)	float64	target candidate	97
PT08.S1(CO)	float64	sensor feature	1042
NMHC(GT)	float64	target candidate	430
C6H6(GT)	float64	target candidate	408
PT08.S2(NMHC)	float64	sensor feature	1246
NOx(GT)	float64	target candidate	926
PT08.S3(NOx)	float64	sensor feature	1222
NO2(GT)	float64	target candidate	284
PT08.S4(NO2)	float64	sensor feature	1604
PT08.S5(O3)	float64	sensor feature	1744
T	float64	weather feature	437
RH	float64	weather feature	754
AH	float64	weather feature	6684

RAW describe() (before treating -200 as missing):

	count	mean	std	min	25%	50%	75%	max
CO(GT)	9,357.000	-34.208	77.657	-200.000	0.600	1.500	2.600	11.900
PT08.S1(CO)	9,357.000	1,048.990	329.833	-200.000	921.000	1,053.000	1,221.000	2,040.000
NMHC(GT)	9,357.000	-159.090	139.789	-200.000	-200.000	-200.000	-200.000	1,189.000
C6H6(GT)	9,357.000	1.866	41.380	-200.000	4.000	7.900	13.600	63.700
PT08.S2(NMHC)	9,357.000	894.595	342.333	-200.000	711.000	895.000	1,105.000	2,214.000
NOx(GT)	9,357.000	168.617	257.434	-200.000	50.000	141.000	284.000	1,479.000
PT08.S3(NOx)	9,357.000	794.990	321.994	-200.000	637.000	794.000	960.000	2,683.000
NO2(GT)	9,357.000	58.149	126.940	-200.000	53.000	96.000	133.000	340.000
PT08.S4(NO2)	9,357.000	1,391.480	467.210	-200.000	1,185.000	1,446.000	1,662.000	2,775.000
PT08.S5(O3)	9,357.000	975.072	456.938	-200.000	700.000	942.000	1,255.000	2,523.000
T	9,357.000	9.778	43.204	-200.000	10.900	17.200	24.100	44.600
RH	9,357.000	39.485	51.216	-200.000	34.100	48.600	61.900	88.700
AH	9,357.000	-6.838	38.977	-200.000	0.692	0.977	1.296	2.231

Observation — hidden missing values. min is -200 in every column. Concentrations, humidity and sensor voltages cannot be negative, and T = -200 °C is physically impossible. The UCI documentation states that **-200 is the missing-value sentinel**. Left untreated it would badly distort every statistic, scaler and regression fit, so we convert it to NaN first.

1.4 Missing-value diagnosis



```
df = df.replace(-200, np.nan)

miss = pd.DataFrame({"missing": df.isna().sum(), "missing_%": df.isna().mean() * 100,
                    "valid": df.notna().sum()}).sort_values("missing_", ascending=False)
display(miss.style.background_gradient(subset=["missing_%"], cmap="Oranges").format({"missing_%": "{:.1f}%"}))

fig, axes = plt.subplots(1, 2, figsize=(14, 4.5), gridspec_kw={"width_ratios": [1, 1.6]})
m = miss["missing_%"].sort_values()
axes[0].barh(m.index, m.values, color=[RED if v > 50 else ORANGE if v > 10 else BLUE for v in m.values])
for i, v in enumerate(m.values): axes[0].text(v + 1, i, f"{v:.1f}%", va="center", fontsize=9)
axes[0].set_xlim(0, 105); axes[0].set_xlabel("% missing"); axes[0].set_title("Missing share per column")

# Missingness over time: one row per column, dark = missing
weekly_missing = df.isna().astype(int).resample("W").mean()
sns.heatmap(weekly_missing.T, cmap="Greys", cbar_kws={"label": "weekly share missing"}, ax=axes[1], xticklabels=False)
wk = weekly_missing.index
axes[1].set_xticks(np.arange(0, len(wk), 8)); axes[1].set_xticklabels([d.strftime("%b %y") for d in wk[::8]], rotation=0)
axes[1].set_title("Missingness over time (weekly)"); axes[1].set_xlabel("")
plt.tight_layout(); plt.show()
```

	missing	missing_%	valid
NMHC(GT)	8443	90.2%	914
CO(GT)	1683	18.0%	7674
NO2(GT)	1642	17.5%	7715
NOx(GT)	1639	17.5%	7718
PT08.S1(CO)	366	3.9%	8991
C6H6(GT)	366	3.9%	8991
PT08.S2(NMHC)	366	3.9%	8991
PT08.S3(NOx)	366	3.9%	8991
PT08.S4(NO2)	366	3.9%	8991
PT08.S5(O3)	366	3.9%	8991
T	366	3.9%	8991
RH	366	3.9%	8991
AH	366	3.9%	8991




```
# Is the missingness structured? (a) rows where ALL sensors+weather are missing, (b) how GT-missingness overlaps
all_feat_missing = df[FEATURE_CANDIDATES].isna().all(axis=1)
any_feat_missing = df[FEATURE_CANDIDATES].isna().any(axis=1)
print(f"Rows with ALL 8 sensor/weather features missing : {all_feat_missing.sum()}")
print(f"Rows with ANY sensor/weather feature missing : {any_feat_missing.sum()} "
      f" -> partial-missing rows = {(any_feat_missing & ~all_feat_missing).sum()}")
print(f"Rows where C6H6(GT) is missing ALSO have all features missing: "
      f"{(df['C6H6(GT)'].isna() == all_feat_missing).all()}")

# Missingness co-occurrence matrix (Jaccard-style: P(col B missing | col A missing))
M = df.isna().astype(int)
cooc = pd.DataFrame({b: {a: (M[a] & M[b]).sum() / max(M[a].sum(), 1) for a in M.columns} for b in M.columns})
plt.figure(figsize=(8, 6))
sns.heatmap(cooc, annot=True, fmt=".2f", cmap="Blues", vmin=0, vmax=1, cbar_kws={"label": "P(column on x missing | column on y missing)"})
plt.title("Missingness co-occurrence"); plt.tight_layout(); plt.show()

# Longest consecutive outage of target candidates (hours)
def longest_run(s):
    g = s.ne(s.shift()).cumsum(); r = s.groupby(g).sum(); return int(r.max())
print("\nLongest consecutive outage (hours):")
print({c: longest_run(df[c].isna()) for c in ["CO(GT)", "NOx(GT)", "NO2(GT)", "C6H6(GT)", "PT08.S1(CO)"]})
```

Rows with ALL 8 sensor/weather features missing : 366
 Rows with ANY sensor/weather feature missing : 366 -> partial-missing rows = 0
 Rows where C6H6(GT) is missing ALSO have all features missing: True



Longest consecutive outage (hours):
 {'CO(GT)': 173, 'NOx(GT)': 173, 'NO2(GT)': 173, 'C6H6(GT)': 76, 'PT08.S1(CO)': 76}

Diagnosis

- **NMHC(GT) is ~90 % missing** (only 914 valid hours). It is unusable as a target and cannot be imputed credibly → **excluded**.
- All five sensors, T, RH, AH and C6H6(GT) share the **same 366 missing hours**: whole rows are blank (instrument outages), there are *no* partially-missing feature rows. The missingness co-occurrence matrix confirms the block structure. There is nothing to impute for the features – those rows hold **zero** information and are dropped.
- CO(GT), NOx(GT), NO2(GT) are additionally missing for ~1.3–1.7 k hours (~17–18 %), as the reference analyser was offline at times, often for long consecutive stretches (see the outage lengths above) — time-interpolating across multi-day gaps would fabricate data.
- **Never impute the target**: a regression label filled in by a heuristic makes the model learn the heuristic. Rows without a target are dropped.

1.5 Choosing the target variable

```
cand = pd.DataFrame(index=GT_COLS)
cand["valid_rows"] = df[GT_COLS].notna().sum()
cand["valid_%"] = df[GT_COLS].notna().mean() * 100
cand["usable_rows (target & features valid)"] = [(df[c].notna() & ~any_feat_missing).sum() for c in GT_COLS]
cand["skew"] = df[GT_COLS].skew()

# strongest single-sensor linear relation (Pearson r) and the sensor it comes from
best = {c: df[[c] + FEATURE_CANDIDATES].corr()[c].drop(c).abs().sort_values(ascending=False).head(1) for c in GT_COLS}
cand["best |r| with a feature"] = [best[c].iloc[0] for c in GT_COLS]
cand["... which feature"] = [best[c].index[0] for c in GT_COLS]
display(cand)
```

	valid_rows	valid_%	usable_rows (target & features valid)	skew	best r with a feature	... which feature
CO(GT)	7674	82.013	7344	1.370	0.916	PT08.S2(NMHC)
NMHC(GT)	914	9.768	887	1.557	0.878	PT08.S2(NMHC)
C6H6(GT)	8991	96.088	8991	1.362	0.982	PT08.S2(NMHC)
NOx(GT)	7718	82.484	7396	1.716	0.787	PT08.S5(O3)
NO2(GT)	7715	82.452	7393	0.622	0.708	PT08.S5(O3)

Decision: target = CO(GT) (true hourly CO concentration, mg/m³).

| Candidate | Verdict | |---|---| | NMHC (GT) | ~90 % missing → rejected. | | C6H6 (GT) | Most complete, but benzene in this dataset is almost a deterministic function of the PT08.S2(NMHC) sensor ($|r| \approx 0.98$) — the task would be trivial and not discriminate between models. | | NOx (GT) / NO2 (GT) | Valid, but ~17 % missing and NOx is extremely right-skewed. | | **CO (GT)** | ~7.7 k valid rows (plenty), has a **strong but imperfect** relation ($|r| \approx 0.9$ with PT08.S2) with several sensors, a visible non-linear component, and a diurnal traffic signal — ideal for contrasting Linear/Polynomial regression with Decision Trees. |

Features are restricted to the **8 sensors + weather variables** as the brief specifies; the other GT pollutants are *not* used as predictors (they would be unavailable at inference time for a low-cost sensor and would leak the answer).

```
TARGET = "CO(GT)"
FEATURES = FEATURE_CANDIDATES.copy()
print("Target :", TARGET)
print("Features:", FEATURES)

# Drop: (i) rows with no sensor/weather readings, (ii) rows with no target. -> Never impute the target.
data = df[FEATURES + [TARGET]].dropna()
flow = pd.DataFrame({"step": ["Raw rows after structural cleaning", "- rows with all features missing (sensor outage)",
                             f"- rows with missing target {TARGET}"], "= Modelling dataset"],
                    "rows": [len(df), -int(any_feat_missing.sum()),
                             -int((df[FEATURES].notna().all(axis=1) & df[TARGET].isna()).sum()), len(data)]})

display(flow)
assert data.isna().sum().sum() == 0
print(f"Retained {len(data)/len(df):.1%} of hourly records; remaining NaNs: {int(data.isna().sum().sum())}")
```

Target : CO(GT)
Features: ['PT08.S1(CO)', 'PT08.S2(NMHC)', 'PT08.S3(NOx)', 'PT08.S4(NO2)', 'PT08.S5(O3)', 'T', 'RH', 'AH']

	step	rows
0	Raw rows after structural cleaning	9357
1	- rows with all features missing (sensor outage)	-366
2	- rows with missing target CO(GT)	-1647
3	= Modelling dataset	7344

Retained 78.5% of hourly records; remaining NaNs: 0

1.6 Target analysis — CO (GT) : distribution, class-wise balance

```
y = data[TARGET]
desc = y.describe(percentiles=[.01, .05, .25, .5, .75, .95, .99]).to_frame("CO(GT)").T
desc["skew"], desc["kurtosis"] = y.skew(), y.kurt()
desc["CV (std/mean)"] = y.std() / y.mean()
display(desc)


# Resolution of the label: CO is reported in 0.1 mg/m3 steps
print("Distinct target values:", y.nunique(), "| smallest step:", np.diff(np.sort(y.unique()))[0].round(2))

fig, ax = plt.subplots(1, 4, figsize=(18, 4))
ax[0].hist(y, bins=60, color=BLUE, edgecolor="white"); ax[0].axvline(y.mean(), color=RED, ls="--", label=f"mean {y.mean():.2f}")
ax[0].axvline(y.median(), color=ORANGE, ls="--", label=f"median {y.median():.2f}"); ax[0].legend(); ax[0].set_title("CO(GT) histogram");
ax[0].set_xlabel("mg/m³")
ax[1].hist(np.log1p(y), bins=60, color=GREEN, edgecolor="white"); ax[1].set_title("log(1 + CO(GT))"); ax[1].set_xlabel("log(1+mg/m³)")
sns.boxplot(x=y, ax=ax[2], color=BLUE, fliersize=2); ax[2].set_title("Boxplot (outliers by 1.5·IQR)")
stats.probplot(y, dist="norm", plot=ax[3]); ax[3].set_title("Normal Q-Q plot")
ax[3].get_lines()[0].set_color(BLUE); ax[3].get_lines()[1].set_color(RED)
plt.tight_layout(); plt.show()

sh_stat, sh_p = stats.shapiro(y.sample(5000, random_state=SEED))
print(f"Shapiro-Wilk (5k sample): W={sh_stat:.3f}, p={sh_p:.1e} -> normality {'rejected' if sh_p < .05 else 'not rejected'}; "
      f"skew {y.skew():.2f} -> log-transform skew {np.log1p(y).skew():.2f}")
```

	count	mean	std	min	1%	5%	25%	50%	75%	95%	99%	max	skew	kurtosis	CV (std/mean)
CO(GT)	7,344,000	2.130	1.436	0.100	0.200	0.500	1.100	1.800	2.800	4.900	6.700	11.900	1.357	2.658	0.674

Distinct target values: 94 | smallest step: 0.1



Shapiro-Wilk (5k sample): $W=0.901$, $p=5.4e-49$ → normality rejected; skew 1.36 → log-transform skew 0.22

Observations. CO(GT) is **right-skewed** (skew ≈ 1.3 , long tail up to ~ 12 mg/m³ with a few high-pollution episodes); mean > median; Shapiro-Wilk rejects normality; the Q-Q plot bends upward at the upper tail. The values are quantised in 0.1 mg/m³ steps. Implications: MSE/RMSE will be dominated by the rare high-concentration hours, and a plain linear fit will tend to under-predict the peaks. We keep the target in **original units** so the reported RMSE is physically interpretable (mg/m³) and comparable across all models; a log-transform is noted as an optional experiment, not applied by default.

Class-wise balance

CO(GT) is a continuous target, so "classes" are defined by binning it into **air-quality regimes**. We use the quartile bins (equal-count, to check how the split is balanced) and the EU/WHO-motivated bands (8-h limit of ~ 10 mg/m³ is hit only extremely rarely, so we use practical bands: *Low* < 1.5, *Moderate* 1.5–3, *High* 3–5, *Very high* ≥ 5 mg/m³).

```
bands_edges = [0, 1.5, 3, 5, np.inf]
bands_labels = ["Low (<1.5)", "Moderate (1.5-3)", "High (3-5)", "Very high (>=5)"]
band = pd.cut(y, bins=bands_edges, labels=bands_labels, right=False)
quart = pd.qcut(y, 4, duplicates="drop")


bal = pd.DataFrame({"count": band.value_counts().sort_index()})
bal["share_%"] = bal["count"] / bal["count"].sum() * 100
bal["mean_CO"] = y.groupby(band).mean(); bal["std_CO"] = y.groupby(band).std()
display(bal)
imbalance_ratio = bal["count"].max() / bal["count"].min()
print(f"Imbalance ratio (largest : smallest regime) = {imbalance_ratio:.1f} : 1")
print("Quartile bins:", quart.value_counts().sort_index().to_dict())

fig, ax = plt.subplots(1, 2, figsize=(13, 4))
colors = [BLUE, GREEN, ORANGE, RED]
ax[0].bar(bal.index, bal["count"], color=colors)
for i, (c, s) in enumerate(zip(bal["count"], bal["share_%"])): ax[0].text(i, c + 40, f"{c:,}\n({s:.1f}%)", ha="center", fontsize=9)
ax[0].set_ylim(0, bal["count"].max() * 1.2); ax[0].set_title("Class-wise balance of CO(GT) regimes"); ax[0].set_ylabel("hours")
ecdf = np.sort(y); ax[1].plot(ecdf, np.arange(1, len(ecdf) + 1) / len(ecdf), color=BLUE)
for e in bands_edges[1:-1]: ax[1].axvline(e, color=GREY, ls=":")
ax[1].set_title("Empirical CDF with regime thresholds"); ax[1].set_xlabel("CO(GT) mg/m³"); ax[1].set_ylabel("cumulative share")
plt.tight_layout(); plt.show()
```

	count	share_%	mean_CO	std_CO
CO(GT)				
Low (<1.5)	2820	38.399	0.884	0.341
Moderate (1.5-3)	2808	38.235	2.102	0.423
High (3-5)	1369	18.641	3.740	0.548
Very high (>=5)	347	4.725	6.124	1.141

Imbalance ratio (largest : smallest regime) = 8.1 : 1

Quartile bins: {Interval(0.099, 1.1, closed='right'): 2064, Interval(1.1, 1.8, closed='right'): 1698, Interval(1.8, 2.8, closed='right'): 1753, Interval(2.8, 11.9, closed='right'): 1829}



Diagnosis — imbalance. The regimes are clearly **imbalanced**: most hours are *Low/Moderate* and only a few percent are *Very high*. For regression this matters in three ways:

1. A random split may by chance put too few extreme hours in the test set → unstable R^2 . We therefore make the **80/20 split stratified on these regime bins**, so train and test have the same target distribution (verified in §1.12).
2. Error metrics are dominated by the rare high values; we will also report errors **per regime** during evaluation (Stage 3).
3. We do **not** re-sample/oversample (that is a classification remedy and would distort the continuous target); stratification is the appropriate control.

1.7 Feature analysis — distributions & descriptive statistics


```
X = data[FEATURES]
fs = X.describe(percentiles=[.01, .25, .5, .75, .99]).T
fs["skew"], fs["kurtosis"], fs["CV"] = X.skew(), X.kurt(), X.std() / X.mean()
fs["range"] = fs["max"] - fs["min"]
display(fs)
print("Scale mismatch: largest/smallest feature std =", f"{X.std().max() / X.std().min():.0f}x",
      "| largest/smallest mean magnitude =", f"{X.mean().abs().max() / X.mean().abs().min():.0f}x")
```

	count	mean	std	min	1%	25%	50%	75%	99%	max	skew	kurtosis	CV	range
PT08.S1(CO)	7,344.000	1,110.581	218.681	647.000	746.000	946.000	1,075.000	1,246.000	1,722.570	2,040.000	0.727	0.294	0.197	1,393.000
PT08.S2(NMHC)	7,344.000	947.198	265.472	387.000	471.430	743.000	919.000	1,125.250	1,642.140	2,214.000	0.532	0.033	0.280	1,827.000
PT08.S3(NOx)	7,344.000	826.920	256.648	322.000	401.000	649.000	795.000	960.000	1,664.570	2,683.000	1.152	2.917	0.310	2,361.000
PT08.S4(NO2)	7,344.000	1,444.753	350.344	551.000	742.000	1,203.000	1,447.000	1,673.000	2,342.420	2,775.000	0.219	-0.031	0.242	2,224.000
PT08.S5(O3)	7,344.000	1,043.513	405.570	221.000	347.000	744.750	990.000	1,305.000	2,109.140	2,523.000	0.577	0.003	0.389	2,302.000
T	7,344.000	17.770	8.863	-1.900	2.043	11.200	16.900	23.800	39.300	44.600	0.377	-0.417	0.499	46.500
RH	7,344.000	49.060	17.452	9.200	14.800	35.400	49.300	62.500	83.057	88.700	-0.015	-0.844	0.356	79.500
AH	7,344.000	0.989	0.400	0.185	0.273	0.698	0.960	1.259	1.940	2.181	0.340	-0.460	0.404	1.996

Scale mismatch: largest/smallest feature std = 1,014x | largest/smallest mean magnitude = 1,460x

```
fig, axes = plt.subplots(2, 4, figsize=(18, 7))
for a, c in zip(axes.ravel(), FEATURES):
    sns.histplot(X[c], bins=50, kde=True, color=BLUE, ax=a, edgecolor="white")
    a.axvline(X[c].median(), color=ORANGE, ls="---", lw=1); a.set_title(f"{c} (skew {X[c].skew():.2f})"); a.set_xlabel("")
plt.suptitle("Feature distributions (dashed = median)", y=1.01, fontweight="bold"); plt.tight_layout(); plt.show()

fig, ax = plt.subplots(figsize=(14, 4.2))
sns.boxplot(data=(X - X.mean()) / X.std(), orient="v", ax=ax, color=BLUE, fliersize=2)
ax.set_title("Boxplots of z-scored features (shows relative spread & outliers on one scale)"); ax.set_ylabel("z-score")
plt.tight_layout(); plt.show()
```



Observations.

- **Wildly different scales:** sensors are in the hundreds–thousands ($\approx 200\text{--}2800$), $T \approx -2\text{...}45\text{ }^{\circ}\text{C}$, $RH \approx 10\text{--}90\%$, $AH \approx 0.2\text{--}2.2$. Gradient-descent regression (Stage 2, "metric vs. epoch") and polynomial features need **standardisation**; otherwise the large-scale sensors dominate the gradient and the high-degree terms explode numerically. Decision Trees are scale-invariant, so they will be fed the same scaled data purely for a fair like-for-like comparison.
- PT08.S3(NOx) is clearly right-skewed with a long tail (skew ≈ 1.15 , kurtosis ≈ 2.9); the other sensors are mildly skewed; T, RH, AH are broad and roughly symmetric (RH and AH platykurtic).
- No physically impossible values remain after the -200 treatment (checked below).

1.8 Outlier & validity audit


```
Z = (X - X.mean()) / X.std()
q1, q3 = X.quantile(.25), X.quantile(.75); iqr = q3 - q1
out = pd.DataFrame({
    "IQR outliers (n)": ((X < q1 - 1.5 * iqr) | (X > q3 + 1.5 * iqr)).sum(),
    "IQR outliers (%)": ((X < q1 - 1.5 * iqr) | (X > q3 + 1.5 * iqr)).mean() * 100,
    "|z|>3 (n)": (Z.abs() > 3).sum(),
    "max |z|": Z.abs().max()})
yq1, yq3 = y.quantile([.25, .75]); yi = yq3 - yq1
out.loc[TARGET] = [int(((y < yq1 - 1.5*yi) | (y > yq3 + 1.5*yi)).sum()), ((y < yq1 - 1.5*yi) | (y > yq3 + 1.5*yi)).mean()*100,
    int((((y - y.mean()) / y.std()).abs() > 3).sum()), ((y - y.mean()) / y.std()).abs().max()])
display(out)

# Physical-validity checks
checks = {"CO(GT) >= 0": (y >= 0).all(), "RH within [0,100]": X["RH"].between(0, 100).all(),
    "AH > 0": (X["AH"] > 0).all(), "sensor signals > 0": (X[SENSOR_COLS] > 0).all().all(),
    "T in plausible range (-10..50 C)": X["T"].between(-10, 50).all()}
display(pd.Series(checks, name="passes").to_frame())

# Are the outliers real events or glitches? Look at the extreme hours in time.
fig, ax = plt.subplots(figsize=(14, 3.8))
ax.plot(y.index, y, lw=.5, color=BLUE); hi = y[y > yq3 + 1.5 * yi]
ax.scatter(hi.index, hi, s=6, color=RED, label=f"IQR outliers ({len(hi)})")
ax.set_title("CO(GT) over time — outliers are clustered pollution episodes, not isolated glitches"); ax.set_ylabel("mg/m³"); ax.legend()
plt.tight_layout(); plt.show()
```

	IQR outliers (n)	IQR outliers (%)	z >3 (n)	max z
PT08.S1(CO)	90.000	1.225	46.000	4.250
PT08.S2(NMHC)	47.000	0.640	29.000	4.772
PT08.S3(NOx)	191.000	2.601	91.000	7.232
PT08.S4(NO2)	65.000	0.885	29.000	3.797
PT08.S5(O3)	65.000	0.885	27.000	3.648
T	6.000	0.082	1.000	3.027
RH	0.000	0.000	0.000	2.284
AH	18.000	0.245	0.000	2.979
CO(GT)	257.000	3.499	98.000	6.802

	passes
CO(GT) >= 0	True
RH within [0,100]	True
AH > 0	True
sensor signals > 0	True
T in plausible range (-10..50 C)	True




Decision — keep the outliers. They are not random sensor glitches: ~70 % of the flagged CO hours have another flagged hour right next to them (contiguous *episodes*), and they concentrate in the **morning (08–09 h) and evening (17–20 h) rush hours**, consistent with traffic emissions and with high sensor readings. Removing them would delete exactly the hours a pollution model should learn to predict, and make the test score optimistic. Standardisation (mean/std) is sensitive to outliers but, since none are impossible values, it is retained; no clipping/winsorising is applied.

1.9 Relationships — correlations, linearity, multicollinearity

```
fig, ax = plt.subplots(1, 2, figsize=(16, 6))
corr_p = data.corr(method="pearson"); corr_s = data.corr(method="spearman")
mask = np.triu(np.ones_like(corr_p, dtype=bool), k=1)
sns.heatmap(corr_p, mask=mask, annot=True, fmt=".2f", cmap="RdBu_r", center=0, vmin=-1, vmax=1, ax=ax[0], cbar=False, annot_kws={"size": 8})
ax[0].set_title("Pearson (linear)")

sns.heatmap(corr_s, mask=mask, annot=True, fmt=".2f", cmap="RdBu_r", center=0, vmin=-1, vmax=1, ax=ax[1], annot_kws={"size": 8})
ax[1].set_title("Spearman (monotonic)"); plt.tight_layout(); plt.show()


rel = pd.DataFrame({"Pearson r": corr_p[TARGET], "Spearman rho": corr_s[TARGET]}).drop(TARGET)
rel["|gap|"] = (rel["Spearman rho"].abs() - rel["Pearson r"].abs())
display(rel.sort_values("Pearson r", key=abs, ascending=False))
```



	Pearson r	Spearman rho	gap
PT08.S2(NMHC)	0.916	0.927	0.011
PT08.S1(CO)	0.879	0.879	0.000
PT08.S5(O3)	0.854	0.856	0.002
PT08.S3(NOx)	-0.703	-0.811	0.108
PT08.S4(NO2)	0.631	0.594	-0.036
RH	0.049	-0.006	-0.042
AH	0.049	0.053	0.005
T	0.022	0.069	0.047

```
# Scatter of each feature vs target with linear (red) and LOWESS (orange) fits -> reveals non-linearity
from statsmodels.nonparametric.smoothers_lowess import lowess
fig, axes = plt.subplots(2, 4, figsize=(18, 7.5))
samp = data.sample(2500, random_state=SEED)
for a, c in zip(axes.ravel(), FEATURES):
    a.scatter(samp[c], samp[TARGET], s=5, alpha=.3, color=BLUE)
    b = np.polyfit(data[c], y, 1); xs = np.linspace(data[c].min(), data[c].max(), 100)
    a.plot(xs, np.polyval(b, xs), color=RED, lw=2, label="linear")
    lw = lowess(samp[TARGET], samp[c], frac=.3); a.plot(lw[:, 0], lw[:, 1], color=ORANGE, lw=2, label="LOWESS")
    a.set_title(f"{c} (r={corr_p.loc[c, TARGET]:.2f})"); a.set_xlabel(c); a.set_ylabel("CO(GT)")
axes[0, 0].legend(); plt.suptitle("Feature vs target (2,500-point sample)", y=1.01, fontweight="bold"); plt.tight_layout(); plt.show()
```

Feature vs target (2,500-point sample)



```
# Multicollinearity: VIF (>5 moderate, >10 severe) and condition number of the standardised feature matrix
Xs = (X - X.mean()) / X.std()
vif = pd.Series([variance_inflation_factor(Xs.values, i) for i in range(Xs.shape[1])], index=FEATURES, name="VIF").sort_values(ascending=False)
display(vif.to_frame().style.background_gradient(cmap="Reds").format("{:.1f}"))
print(f"Condition number (standardised X): {np.linalg.cond(Xs.values):.1f}")


# Most correlated feature pairs
pairs = corr_p.loc[FEATURES, FEATURES].where(np.triu(np.ones((8, 8), bool), 1)).stack().rename("r")
display(pairs.reindex(pairs.abs().sort_values(ascending=False).index).head(6).to_frame())
sns.pairplot(samp[["PT08.S1(CO)", "PT08.S2(NMHC)", "PT08.S5(O3)", "T", "AH", TARGET]].sample(1200, random_state=SEED),
             corner=True, plot_kws={"s": 8, "alpha": .4, "color": BLUE}, diag_kws={"color": BLUE})
plt.suptitle("Pairplot of the key variables", y=1.02, fontweight="bold"); plt.show()
```

	VIF
PT08.S2(NMHC)	15.2
T	13.6
PT08.S4(NO2)	10.3
AH	10.0
PT08.S5(O3)	8.2
PT08.S1(CO)	8.1
RH	7.6
PT08.S3(NOx)	4.1

Condition number (standardised X): 10.8

		r
PT08.S1(CO)	PT08.S5(O3)	0.898
	PT08.S2(NMHC)	0.888
PT08.S2(NMHC)	PT08.S5(O3)	0.877
PT08.S3(NOx)	PT08.S5(O3)	-0.794
PT08.S2(NMHC)	PT08.S3(NOx)	-0.785
	PT08.S4(NO2)	0.772

Pairplot of the key variables



Observations.



- The CO target is most strongly related to **PT08.S2(NMHC)** , **PT08.S1(CO)** , **PT08.S5(O3)** (and negatively to **PT08.S3(NOx)**); the weather variables show weak linear association. **PT08.S3(NOx)** has $|\text{Spearman}| \approx 0.81$ vs $|\text{Pearson}| \approx 0.70$, i.e. a monotonic but clearly non-linear relation, and the LOWESS curves bend away from the straight line (see plots) → there is **some genuine non-linearity**, which motivates the polynomial and tree models.
- **Severe multicollinearity**: the CO, NMHC and O₃ sensors are strongly inter-correlated, AH is tied to T / RH , and four VIFs exceed 10 (**PT08.S2** ≈ 15 , **T** ≈ 14 , **PT08.S4** ≈ 10 , **AH** ≈ 10 ; only **PT08.S3** is below 5), although the condition number of the standardised matrix is a moderate ≈ 11 . Consequences for the next stages: (a) a degree-1 OLS fit is numerically fine after standardisation, but polynomial expansions will multiply the collinearity → use a stable `lstsq / pinv` solve or a small L2 term; (b) individual **linear coefficients will not be reliably interpretable**; (c) tree **feature importances will be split among correlated sensors** – interpret groups, not single features.
- All 8 features are kept: the assignment asks for sensors + weather, no feature is constant, and VIF-driven dropping would hurt tree performance for no benefit. Regularisation (Ridge) is *available* as a remedy if the polynomial fit proves unstable (decided in Stage 2 with validation data, not on the test set).

1.10 Temporal structure (seasonality, diurnal cycle, sensor drift)

```
d = data.copy(); d["hour"] = d.index.hour; d["dow"] = d.index.dayofweek; d["month"] = d.index.to_period("M")
fig, ax = plt.subplots(1, 3, figsize=(18, 4))
h = d.groupby("hour")[TARGET].agg(["mean", "std"]); ax[0].plot(h.index, h["mean"], color=BLUE, lw=2)
ax[0].fill_between(h.index, h["mean"] - h["std"], h["mean"] + h["std"], color=BLUE, alpha=.15); ax[0].set_title("CO(GT) by hour of day (mean ± 1 sd)");
ax[0].set_xlabel("hour")
wd = d.groupby("dow")[TARGET].mean(); ax[1].bar(["Mon", "Tue", "Wed", "Thu", "Fri", "Sat", "Sun"], wd.values, color=[BLUE]*5 + [ORANGE]*2);
ax[1].set_title("CO(GT) by day of week")
mo = d.groupby("month")[TARGET].mean(); ax[2].plot(mo.index.astype(str), mo.values, marker="o", color=BLUE); ax[2].set_title("Monthly mean CO(GT)");
ax[2].tick_params(axis="x", rotation=60)
plt.tight_layout(); plt.show()

# Sensor drift / seasonal shift: monthly mean of each (z-scored) feature
mz = (d[FEATURES].groupby(d["month"]).mean() - X.mean()) / X.std()
plt.figure(figsize=(12, 3.8)); sns.heatmap(mz.T, cmap="RdBu_r", center=0, annot=True, fmt=".1f", annot_kws={"size": 7}, cbar_kws={"label": "z of monthly mean"})
plt.title("Monthly mean of each feature (z-scored) – seasonal shift & sensor drift"); plt.xlabel(""); plt.tight_layout(); plt.show()

# Autocorrelation of the target: how dependent are neighbouring hours?
ac = [y.autocorr(l) for l in (1, 2, 3, 6, 12, 24, 48)]
display(pd.DataFrame({"lag (h)": (1, 2, 3, 6, 12, 24, 48), "autocorr": ac}).set_index("lag (h)").T)
```



lag (h)	1	2	3	6	12	24	48
autocorr	0.831	0.595	0.405	0.136	0.137	0.521	0.282

Observations.

- Clear **diurnal cycle** (peaks ≈ 2.9 mg/m³ at 09 h and ≈ 3.7 at 19 h, night minimum ≈ 0.7 at 05 h) and a **weekend dip** (Sat ≈ 1.9 , Sun ≈ 1.4 vs ≈ 2.4 on weekdays) $\rightarrow$ CO is traffic-driven. The brief restricts predictors to sensors + weather; the sensors implicitly carry the diurnal signal, so we do **not** add hour/day features (a possible extension).
- **Seasonality & drift**: monthly mean CO is highest in **Oct-Dec** (≈ 2.6 – 2.7) and lowest in **Aug** (≈ 1.25 , holiday period); T swings from ≈ 30 °C (Aug) to ≈ 7 °C (Feb). The PT08.S4(NO2) sensor mean falls steadily from $\approx 1,600$ – $1,700$ (spring/summer 2004) to $\approx 1,000$ – $1,200$ (winter 2005) — a mix of seasonal effect and metal-oxide sensor drift, which matters for any out-of-time evaluation.
- **Strong autocorrelation** (lag-1 ≈ 0.83 , lag-24 ≈ 0.52): neighbouring hours are near duplicates. A *random* split therefore yields a test set interleaved with training hours (optimistic, "interpolation" scores), whereas a *chronological* split measures true forecasting/generalisation to a new season. Evaluated in §1.11.

1.11 Train/test split strategy — random vs. chronological

```
# What would a chronological 80/20 split look like? (first 80% of time -> train, last 20% -> test)
cut = int(len(data) * .8); tr_t, te_t = data.iloc[:cut], data.iloc[cut:]
print(f"Chronological split: train {tr_t.index.min():%d %b %Y} -> {tr_t.index.max():%d %b %Y} | test {te_t.index.min():%d %b %Y} -> {te_t.index.max():%d %b %Y}")
cmp = pd.DataFrame({"train mean": tr_t[FEATURES + [TARGET]].mean(), "test mean": te_t[FEATURES + [TARGET]].mean()})
cmp["shift (in train std)"] = (cmp["test mean"] - cmp["train mean"]) / tr_t[FEATURES + [TARGET]].std()
display(cmp)
ks = {c: stats.ks_2samp(tr_t[c], te_t[c]).statistic for c in FEATURES + [TARGET]}
print("KS statistic train-vs-test (chronological):", {k: round(v, 2) for k, v in ks.items()})
```

Chronological split: train 10 Mar 2004 → 28 Jan 2005 | test 28 Jan 2005 → 04 Apr 2005

	train mean	test mean	shift (in train std)
PT08.S1(CO)	1,112.152	1,104.295	-0.035
PT08.S2(NMHC)	970.195	855.228	-0.433
PT08.S3(NOx)	842.590	764.251	-0.300
PT08.S4(NO2)	1,523.574	1,129.521	-1.220
PT08.S5(O3)	1,047.547	1,027.378	-0.051
T	19.556	10.631	-1.049
RH	48.894	49.724	0.047
AH	1.076	0.645	-1.134
CO(GT)	2.175	1.950	-0.154

KS statistic train-vs-test (chronological): {'PT08.S1(CO)': np.float64(0.04), 'PT08.S2(NMHC)': np.float64(0.2), 'PT08.S3(NOx)': np.float64(0.14), 'PT08.S4(NO2)': np.float64(0.53), 'PT08.S5(O3)': np.float64(0.07), 'T': np.float64(0.43), 'RH': np.float64(0.05), 'AH': np.float64(0.52), 'CO(GT)': np.float64(0.09)}

Decision. A chronological split trains on Mar 2004 – Jan 2005 (including the hot summer) and tests on the colder, drier late-Jan – Apr 2005 period, creating a large **covariate shift**: T −1.0 sd, AH −1.1 sd, PT08.S4 −1.2 sd (KS statistics ≈ 0.4–0.5), while CO itself shifts only −0.15 sd. Such a test would mostly measure season extrapolation rather than model quality, and the polynomial/tree comparison the assignment asks for would be dominated by that shift.

The assignment simply specifies an **80/20 train/test split**, so we use a **shuffled, target-stratified 80/20 split** (random_state=42) as the primary protocol. We state its limitation openly: because of hourly autocorrelation, test scores will be somewhat optimistic relative to true forecasting. The chronological split above is kept as a documented *robustness check* option for the Discussion.

1.12 Final preprocessing: stratified 80/20 split → scaling (fit on train only)

```
X_all, y_all = data[FEATURES], data[TARGET]
strata = pd.cut(y_all, bins=bands_edges, labels=bands_labels, right=False) # regimes from §1.6

X_train, X_test, y_train, y_test = train_test_split(
    X_all, y_all, test_size=0.20, random_state=SEED, shuffle=True, stratify=strata)

print(f"Train: {X_train.shape} | Test: {X_test.shape} | test share = {len(X_test)/len(X_all):.1%}")

# Standardisation: statistics learned ONLY from the training set (no test-set leakage), then applied to both
scaler = StandardScaler().fit(X_train)
X_train_s = pd.DataFrame(scaler.transform(X_train), columns=FEATURES, index=X_train.index)
X_test_s = pd.DataFrame(scaler.transform(X_test), columns=FEATURES, index=X_test.index)

chk = pd.DataFrame({"train mean": X_train_s.mean(), "train std": X_train_s.std(ddof=0),
                    "test mean": X_test_s.mean(), "test std": X_test_s.std(ddof=0)})
display(chk)
assert np.allclose(X_train_s.mean(), 0, atol=1e-9) and np.allclose(X_train_s.std(ddof=0), 1)
```

Train: (5875, 8) | Test: (1469, 8) | test share = 20.0%

	train mean	train std	test mean	test std
PT08.S1(CO)	-0.000	1.000	0.003	1.004
PT08.S2(NMHC)	0.000	1.000	-0.023	0.980
PT08.S3(NOx)	0.000	1.000	0.014	1.009
PT08.S4(NO2)	-0.000	1.000	-0.004	0.985
PT08.S5(O3)	0.000	1.000	-0.004	1.008
T	-0.000	1.000	0.011	1.018
RH	-0.000	1.000	-0.008	1.010
AH	-0.000	1.000	0.000	0.987

```
# Verify the split is representative: class balance, target distribution, feature distributions
sb = pd.DataFrame({"train_n": pd.cut(y_train, bands_edges, labels=bands_labels, right=False).value_counts().sort_index(),
                  "test_n": pd.cut(y_test, bands_edges, labels=bands_labels, right=False).value_counts().sort_index()})
sb["train_%"] = sb["train_n"] / sb["train_n"].sum() * 100; sb["test_%"] = sb["test_n"] / sb["test_n"].sum() * 100
display(sb)

tt = pd.DataFrame({"train": y_train.describe(), "test": y_test.describe()}); display(tt.T)
print(f"KS test on target train vs test: stat={stats.ks_2samp(y_train, y_test).statistic:.3f}, p={stats.ks_2samp(y_train, y_test).pvalue:.2f}")


fig, ax = plt.subplots(1, 3, figsize=(18, 4))
w = .38; xi = np.arange(len(sb)); ax[0].bar(xi - w/2, sb["train_%"], w, color=BLUE, label="train"); ax[0].bar(xi + w/2, sb["test_%"], w, color=ORANGE,
label="test")
ax[0].set_xticks(xi); ax[0].set_xticklabels(sb.index, rotation=15); ax[0].set_ylabel("% of split"); ax[0].set_title("Regime balance preserved by
stratification"); ax[0].legend()
sns.kdeplot(y_train, ax=ax[1], color=BLUE, label="train", fill=True, alpha=.25); sns.kdeplot(y_test, ax=ax[1], color=ORANGE, label="test", fill=True,
alpha=.25)
ax[1].set_title("Target density: train vs test"); ax[1].legend()
for c, col in zip(["PT08.S2(NMHC)", "T"], [BLUE, GREEN]):
    sns.kdeplot(X_train_s[c], ax=ax[2], color=col, label=f"{c} train"); sns.kdeplot(X_test_s[c], ax=ax[2], color=col, ls="--", label=f"{c} test")
ax[2].set_title("Scaled feature density: train vs test"); ax[2].legend(fontsize=8)
plt.tight_layout(); plt.show()

# Scaling effect
fig, ax = plt.subplots(1, 2, figsize=(14, 3.8))
sns.boxplot(data=X_train, ax=ax[0], color=BLUE, fliersize=1); ax[0].set_yscale("symlog"); ax[0].set_title("Before scaling (symlog axis!)");
ax[0].tick_params(axis="x", rotation=40)
sns.boxplot(data=X_train_s, ax=ax[1], color=GREEN, fliersize=1); ax[1].set_title("After standardisation (train statistics)");
ax[1].tick_params(axis="x", rotation=40)
plt.tight_layout(); plt.show()
```

	train_n	test_n	train_%	test_%
CO(GT)				
Low (<1.5)	2256	564	38.400	38.393
Moderate (1.5-3)	2246	562	38.230	38.257
High (3-5)	1095	274	18.638	18.652
Very high (>=5)	278	69	4.732	4.697

	count	mean	std	min	25%	50%	75%	max
train	5,875,000	2.132	1.440	0.100	1.100	1.800	2.850	11.900
test	1,469,000	2.120	1.424	0.100	1.000	1.800	2.800	8.700

KS test on target train vs test: stat=0.015, p=0.94



Scaling / normalisation / regularisation plan — decisions made from the EDA

Aspect	Decision	Evidence from EDA
Missing sentinel	-200 → NaN	min = -200 in every column

Aspect	Decision		Evidence from EDA
Missing rows	drop outage rows & rows without target; no imputation	366 whole-row outages; multi-day target gaps; target must not be imputed	
Dropped column	NMHC (GT)	~90 % missing	
Target	CO (GT) , original units (mg/m³)	skewed but interpretable RMSE; log is an optional experiment	
Feature scaling	Z-score standardisation, fitted on train only	scale mismatch of up to ~10³x; needed for gradient-descent "epochs" & polynomial terms; harmless for trees	
Normalisation (min-max)	not used	outliers would squash the bulk of the data; z-score is preferred	
Outliers	kept	real pollution episodes, no impossible values	
Split	80/20 shuffled, stratified on CO regimes, seed=42	class imbalance; stratification keeps test representative	
Regularisation	none by default; Ridge/L2 held in reserve for polynomial regression; trees controlled via max_depth	VIF shows strong multicollinearity; high-degree polynomials may overfit	
Polynomial degree	to be chosen by validation/CV on the training set only	non-linearity observed, but test set stays untouched	
Hand-off to Stage 2	X_train_s, X_test_s, y_train, y_test, FEATURES, TARGET, scaler	—	

```
# Final hand-off summary for the model-training stage
summary = pd.DataFrame({
    "object": ["X_train_s", "X_test_s", "y_train", "y_test"],
    "shape": [X_train_s.shape, X_test_s.shape, y_train.shape, y_test.shape],
    "NaNs": [int(X_train_s.isna().sum().sum()), int(X_test_s.isna().sum().sum()), int(y_train.isna().sum()), int(y_test.isna().sum())],
    "scaled": [True, True, False, False]}.set_index("object"))
display(summary)
print("Features :", FEATURES); print("Target  :", TARGET, "(mg/m³)")
print("Index overlap train/test:", len(set(X_train.index) & set(X_test.index)), "(must be 0)")
assert len(set(X_train.index) & set(X_test.index)) == 0
```

	shape	NaNs	scaled
object			
X_train_s	(5875, 8)	0	True
X_test_s	(1469, 8)	0	True
y_train	(5875,)	0	False
y_test	(1469,)	0	False

```
Features : ['PT08.S1(CO)', 'PT08.S2(NMHC)', 'PT08.S3(NOx)', 'PT08.S4(NO2)', 'PT08.S5(O3)', 'T', 'RH', 'AH']
Target   : CO(GT) (mg/m³)
Index overlap train/test: 0 (must be 0)
```

 **Stage 1 complete.** The data are cleaned, the target and features are fixed, the scaler is fitted on the training data only, and the stratified train/test split is verified.

Stage 2 — Model Training

Design principles carried over from Stage 1

Stage 1 finding	Consequence for training
Hourly data are strongly autocorrelated (lag-1 ≈ 0.83)	Hyper-parameters (polynomial degree, learning rate, tree depth) are chosen with group-aware 5-fold CV on the training set (whole calendar weeks are kept together), never on the test set.
Features are collinear (VIF > 10)	Ridge (L2) is studied as an ablation for the polynomial model; coefficients are reported with bootstrap CIs.
Features differ ~10³x in scale	Gradient descent is run on standardised inputs; an ablation demonstrates why.
CO is right-skewed, ~8:1 regime imbalance	Besides MSE/RMSE/R² we also store residuals and per-regime errors ; log-target is tested as an ablation.
Shuffled split is optimistic	An out-of-time (chronological) robustness check is run for the final models.

What this stage produces for later stages (all kept in memory): the model registry RESULTS , fitted models FITTED , train/test predictions PRED , per-epoch histories HISTORY , the tree depth sweep TREE_SWEEP , feature-importance tables, learning curves and all ablation tables.

2.0 Setup, from-scratch metrics and the model registry

```
import time
from itertools import combinations_with_replacement
import matplotlib
import sklearn
from sklearn.base import BaseEstimator, RegressorMixin, clone
from sklearn.pipeline import Pipeline, make_pipeline
from sklearn.linear_model import LinearRegression, Ridge
from sklearn.preprocessing import PolynomialFeatures
from sklearn.tree import DecisionTreeRegressor
from sklearn.dummy import DummyRegressor
from sklearn.compose import TransformedTargetRegressor
from sklearn.model_selection import GroupKFold, cross_validate, learning_curve, GridSearchCV
from sklearn.metrics import mean_squared_error, r2_score, mean_absolute_error
from sklearn.inspection import permutation_importance

print({m.__name__: m.__version__ for m in (np, pd, sklearn)})

# ---- evaluation metrics implemented from scratch (verified against sklearn below) ----
def mse(y_true, y_pred): return float(np.mean((np.asarray(y_true, float) - np.asarray(y_pred, float)) ** 2))
def rmse(y_true, y_pred): return float(np.sqrt(mse(y_true, y_pred)))
def mae(y_true, y_pred): return float(np.mean(np.abs(np.asarray(y_true, float) - np.asarray(y_pred, float))))
def r2(y_true, y_pred):
    y_true, y_pred = np.asarray(y_true, float), np.asarray(y_pred, float)
    return float(1 - np.sum((y_true - y_pred) ** 2) / np.sum((y_true - y_true.mean()) ** 2))
def regression_metrics(y_true, y_pred):
    return {"mse": mse(y_true, y_pred), "rmse": rmse(y_true, y_pred), "r2": r2(y_true, y_pred), "mae": mae(y_true, y_pred)}

_a, _b = y_test.values, y_test.values * 0.9 + 0.1
assert np.isclose(mse(_a, _b), mean_squared_error(_a, _b)) and np.isclose(r2(_a, _b), r2_score(_a, _b)) and np.isclose(mae(_a, _b),
mean_absolute_error(_a, _b))
print("From-scratch MSE / RMSE / R² / MAE match sklearn.metrics ✓")

# ---- registries shared with Stages 3-6 ----
RESULTS, FITTED, PRED, HISTORY = {}, {}, {}, {}

def timed_fit(model, X, y, **kw):
    t = time.perf_counter(); model.fit(X, y, **kw); return model, time.perf_counter() - t

def evaluate(name, model, family, config, fit_time, complexity=None):
    """Score a fitted model on the scaled train/test sets and store everything needed later."""
    p_tr, p_te = np.asarray(model.predict(X_train_s)), np.asarray(model.predict(X_test_s))
    row = {"model": name, "family": family, "config": config, "complexity": complexity, "fit_time_s": fit_time}
    for split, yt, yp in (("train", y_train, p_tr), ("test", y_test, p_te)):
        row.update({f"{split}_{k}": v for k, v in regression_metrics(yt, yp).items()})
    row["rmse_gap (test-train)"] = row["test_rmse"] - row["train_rmse"]
    RESULTS[name] = row; FITTED[name] = model; PRED[name] = {"train": p_tr, "test": p_te}
    return row

def epochs_to_tol(hist, col="train_mse", tol=0.01):
    """First epoch at which the loss is within `tol` (relative) of its final value."""
    h = hist[col].values; ok = np.where(h <= h[-1] * (1 + tol))[0]
    return int(hist["epoch"].values[ok[0]]) if len(ok) else None

{'numpy': '2.4.6', 'pandas': '3.0.6', 'sklearn': '1.9.1'}
From-scratch MSE / RMSE / R² / MAE match sklearn.metrics ✓
```

2.1 Validation protocol (used for every modelling choice)

The test set is touched only to **report** results. All selection (learning rate, polynomial degree, ridge strength, tree depth/leaf size, ablations) uses 5-fold cross-validation **within the training set**. Because neighbouring hours are near-duplicates, folds are formed from **whole calendar weeks** (`GroupKFold`), so a validation fold never contains hours adjacent to training hours. This gives a more honest (more pessimistic) estimate than a plain shuffled K-fold.

```

groups = pd.factorize(X_train.index.to_period("W").astype(str))[0]      # one group per calendar week
cv = GroupKFold(n_splits=5)
folds = list(cv.split(X_train_s, y_train, groups))

fold_tbl = pd.DataFrame([{"fold": i + 1, "train_n": len(tr), "val_n": len(va), "val_weeks": len(np.unique(groups[va])),
                          "val_CO_mean": y_train.values[va].mean(), "val_CO_std": y_train.values[va].std()}
                          for i, (tr, va) in enumerate(folds)]).set_index("fold")

display(fold_tbl)
assert all(set(groups[tr]).isdisjoint(groups[va]) for tr, va in folds)
print(f"{len(np.unique(groups))} weekly groups; no week appears in both a training and a validation fold ✓")

def cv_eval(est, X=None, y=None, **kw):
    """Group-CV summary of an sklearn-compatible estimator (RMSE / R², mean ± sd over folds)."""
    X = X_train_s if X is None else X; y = y_train if y is None else y
    r = cross_validate(est, X, y, groups=groups, cv=cv, return_train_score=True,
                      scoring={"rmse": "neg_root_mean_squared_error", "r2": "r2"}, **kw)
    return {"cv_rmse": -r["test_rmse"].mean(), "cv_rmse_sd": r["test_rmse"].std(), "cv_r2": r["test_r2"].mean(),
           "cv_train_rmse": -r["train_rmse"].mean()}

```

	train_n	val_n	val_weeks	val_CO_mean	val_CO_std
fold					
1	4706	1169	11	2.068	1.354
2	4704	1171	11	2.019	1.369
3	4695	1180	12	2.045	1.316
4	4705	1170	11	2.186	1.512
5	4690	1185	12	2.342	1.600

57 weekly groups; no week appears in both a training and a validation fold ✓

Observation. 57 weekly groups give five folds of ~1,170 validation hours (11–12 weeks each) with no week shared between training and validation. Fold mean CO ranges from 2.02 to 2.34 mg/m³, so folds are reasonably comparable, and the spread of fold scores (reported as ± sd below) honestly reflects seasonal variation.

2.2 Linear regression from scratch (NumPy only)

One estimator, three solvers, all minimising the same objective $\|y - X\theta\|^2 / n + (\alpha/n) \|\theta\|^2$ (the intercept is never penalised; $\alpha = 0$ is ordinary least squares and α is directly comparable with scikit-learn's `Ridge(alpha)`):

- `solver="normal"` – exact least-squares solution (`np.linalg.lstsq`, SVD-based and therefore stable under collinearity; with $\alpha > 0$ the ridge normal equations are solved);
- `solver="gd"` – batch gradient descent, $\theta \leftarrow \theta - \text{lr} \cdot \nabla$, with **per-epoch recording** of train/test MSE, RMSE and R^2 (needed for the *metric-vs-epoch* plots);
- `solver="adam"` – Adam, used later for the polynomial model where plain GD is badly conditioned.

The class follows the scikit-learn `fit/predict` API purely so that it can plug into the same CV / pipeline tooling – **no scikit-learn algorithm is used inside it**.

```

class LinearRegressionScratch(RegressorMixin, BaseEstimator):
    def __init__(self, solver="normal", alpha=0.0, lr=0.1, epochs=1000, beta1=0.9, beta2=0.999, track=False):
        self.solver, self.alpha, self.lr, self.epochs = solver, alpha, lr, epochs
        self.beta1, self.beta2, self.track = beta1, beta2, track

    @staticmethod
    def _design(X): return np.c_[np.ones(len(X)), np.asarray(X, float)]

    def _grad(self, Xb, y, theta):
        g = (2.0 / len(y)) * Xb.T @ (Xb @ theta - y)
        g[1:] += (2.0 * self.alpha / len(y)) * theta[1:] # L2 term, intercept excluded
        return g

    def fit(self, X, y, eval_sets=None):
        Xb, y = self._design(X), np.asarray(y, float).ravel(); d = Xb.shape[1]
        if self.solver == "normal":
            if self.alpha == 0:
                theta = np.linalg.lstsq(Xb, y, rcond=None)[0]
            else:
                P = self.alpha * np.eye(d); P[0, 0] = 0
                theta = np.linalg.solve(Xb.T @ Xb + P, Xb.T @ y)
        elif self.solver in ("gd", "adam"):
            theta = self._descend(Xb, y, eval_sets)
        else:
            raise ValueError(self.solver)
        self.theta_, self.intercept_, self.coef_ = theta, theta[0], theta[1:]
        return self

    def _descend(self, Xb, y, eval_sets):
        theta, m, v = np.zeros(Xb.shape[1]), np.zeros(Xb.shape[1]), np.zeros(Xb.shape[1])
        ev = [(k: (self._design(Xe), np.asarray(ye, float).ravel()) for k, (Xe, ye) in (eval_sets or {}).items())]
        rows, self.diverged_ = [], False

        def record(ep):
            r = {"epoch": ep}
            for nm, (A, t) in {"train": (Xb, y), **ev}.items():
                p = A @ theta
                r[f"{nm}_mse"], r[f"{nm}_rmse"], r[f"{nm}_r2"] = mse(t, p), rmse(t, p), r2(t, p)
            rows.append(r)

        with np.errstate(all="ignore"):
            if self.track: record(0)
            for t in range(1, self.epochs + 1):
                g = self._grad(Xb, y, theta)
                if self.solver == "gd":
                    theta = theta - self.lr * g
                else: # Adam (bias-corrected)
                    m = self.beta1 * m + (1 - self.beta1) * g
                    v = self.beta2 * v + (1 - self.beta2) * g ** 2
                    theta = theta - self.lr * (m / (1 - self.beta1 ** t)) / (np.sqrt(v / (1 - self.beta2 ** t)) + 1e-8)
                if not np.all(np.isfinite(theta)) or np.abs(theta).max() > 1e12:
                    self.diverged_ = True; break
                if self.track: record(t)
            self.history_ = pd.DataFrame(rows) if self.track else None
            return theta

    def predict(self, X): return self._design(X) @ self.theta_

# ---- unit tests of the implementation ----
# (1) analytic gradient vs central finite differences
rs = np.random.default_rng(0); Xt, yt_, th = rs.normal(size=(60, 4)), rs.normal(size=60), rs.normal(size=5)
lm = LinearRegressionScratch(alpha=0.7); Xb_t = lm._design(Xt)
loss = lambda t: np.mean((Xb_t @ t - yt_) ** 2) + 0.7 / len(yt_) * np.sum(t[1:] ** 2)
num = np.array([(loss(th + 1e-6 * e) - loss(th - 1e-6 * e)) / 2e-6 for e in np.eye(5)])
print("gradient check, max |analytic - numeric| =", f"{np.abs(lm._grad(Xb_t, yt_, th) - num).max():.2e}")
assert np.allclose(lm._grad(Xb_t, yt_, th), num, atol=1e-6)
# (2) recovers known coefficients on noiseless synthetic data
w_true = np.array([1.5, -2.0, 0.5, 3.0]); ys = 4.0 + Xt @ w_true
assert np.allclose(LinearRegressionScratch().fit(Xt, ys).theta_, np.r_[4.0, w_true])
assert np.allclose(LinearRegressionScratch(solver="gd", lr=0.1, epochs=2000).fit(Xt, ys).theta_, np.r_[4.0, w_true], atol=1e-4)
print("Closed form and GD both recover the true coefficients on synthetic data ✓")

```

gradient check, max |analytic - numeric| = 5.32e-09
 Closed form and GD both recover the true coefficients on synthetic data ✓

2.2.1 Choosing the gradient-descent learning rate (train loss only)

```
Xb_all = LinearRegressionScratch._design(X_train_s)
lam_max = np.linalg.eigvalsh(Xb_all.T @ Xb_all / len(Xb_all)).max()
lr_stable = 2.0 / (2.0 * lam_max) # GD on MSE is stable iff lr < 2/L, L = 2 * lam_max(X^T X / n)
print(f"λ_max(X^T X / n) = {lam_max:.2f} -> theoretical stability limit lr < {lr_stable:.3f}")

EPOCHS_LIN = 1000
theta_opt = LinearRegressionScratch().fit(X_train_s, y_train).theta_
lr_grid = [0.001, 0.01, 0.05, 0.1, 0.15, 0.25]
lr_rows, lr_hist = [], {}
for lr in lr_grid:
    m_ = LinearRegressionScratch(solver="gd", lr=lr, epochs=EPOCHS_LIN, track=True).fit(X_train_s, y_train)
    h = m_.history_; lr_hist[lr] = h
    lr_rows.append({"lr": lr, "diverged": m_.diverged_, "final_train_rmse": h["train_rmse"].iloc[-1] if not m_.diverged_ else np.nan,
                  "epochs_to_1%": epochs_to_tol(h) if not m_.diverged_ else None,
                  "||theta - theta_OLS||": np.linalg.norm(m_.theta_ - theta_opt) if not m_.diverged_ else np.nan})
lr_tbl = pd.DataFrame(lr_rows).set_index("lr"); display(lr_tbl)


ok = lr_tbl[~lr_tbl["diverged"]] & (lr_tbl["||theta - theta_OLS||"] < 1e-3)
LR_LIN = float(ok["epochs_to_1%"].astype(float).idxmin()); print("Selected learning rate:", LR_LIN)

fig, ax = plt.subplots(1, 2, figsize=(14, 4.2))
for lr, c in zip(lr_grid, [GREY, BLUE, GREEN, ORANGE, RED, "#CC79A7"]):
    h = lr_hist[lr]; ax[0].plot(h["epoch"], h["train_rmse"].clip(upper=10), color=c, label=f"lr={lr}" + (" (diverged)" if h["train_rmse"].iloc[-1] > 1e3 or len(h) < EPOCHS_LIN else ""))
    ax[0].axhline(rmsely_train, X_train_s.values @ theta_opt[1:] + theta_opt[0], color="k", ls=":", label="OLS optimum")
    ax[0].set_xscale("symlog", linthresh=1); ax[0].set_xlim(0, EPOCHS_LIN); ax[0].set_yscale("log"); ax[0].set_xlabel("epoch"); ax[0].set_ylabel("train RMSE (clipped at 10)")
ax[0].set_title("Learning-rate sweep (batch GD)"); ax[0].legend(fontsize=8)
ax[1].bar(["λ_max-based\nstability limit"] + [str(l) for l in lr_grid], [lr_stable] + lr_grid, color=[RED] + [GREEN if l < lr_stable else RED for l in lr_grid])
ax[1].set_yscale("log"); ax[1].set_title("Tested learning rates vs. the theoretical limit (red = unstable)"); ax[1].set_ylabel("learning rate (log)")
plt.tight_layout(); plt.show()
```

$\lambda_{\max}(X^T X / n) = 4.21 \rightarrow$ theoretical stability limit $lr < 0.238$

lr	diverged	final_train_rmse	epochs_to_1%	theta - theta_OLS
0.001	False	0.618	990.000	1.183
0.010	False	0.492	884.000	0.545
0.050	False	0.480	405.000	0.024
0.100	False	0.480	203.000	0.001
0.150	False	0.480	135.000	0.000
0.250	True	NaN	NaN	NaN

Selected learning rate: 0.15



Observations.

- The theoretical stability limit for plain gradient descent is $lr < 0.238$ (from the largest eigenvalue of $X^T X / n$); the experiment agrees: $lr = 0.25$ diverges, 0.15 is the fastest stable rate.
- Small rates are simply too slow: $lr = 0.001$ has not converged after 1,000 epochs (RMSE 0.618 vs. optimum 0.480).

- **Selected $\text{lr} = 0.15$** : it reaches within 1 % of the optimum in ~135 epochs and ends within 1e-3 of the OLS coefficients.
The choice uses *training loss only*.

2.2.2 Final linear models: closed form, gradient descent (with per-epoch tracking) and scikit-learn

```
N_LIN_NE, N_LIN_GD, N_LIN_SK = "Linear | scratch (normal eq.)", "Linear | scratch (GD)", "Linear | sklearn"
N_MEAN = "Baseline | mean predictor"

m, t = timed_fit(LinearRegressionScratch(solver="normal"), X_train_s, y_train); evaluate(N_LIN_NE, m, "Linear", "OLS, lstsq", t,
complexity=len(FEATURES) + 1)
m, t = timed_fit(LinearRegressionScratch(solver="gd", lr=LR_LIN, epochs=EPOCHS_LIN, track=True), X_train_s, y_train,
eval_sets=({"test": (X_test_s, y_test)}))
evaluate(N_LIN_GD, m, "Linear", f"batch GD, lr={LR_LIN}, {EPOCHS_LIN} epochs", t, complexity=len(FEATURES) + 1); HISTORY[N_LIN_GD] = m.history_
m, t = timed_fit(LinearRegression(), X_train_s, y_train); evaluate(N_LIN_SK, m, "Linear", "sklearn.LinearRegression", t, complexity=len(FEATURES) + 1)
m, t = timed_fit(DummyRegressor(strategy="mean"), X_train_s, y_train); evaluate(N_MEAN, m, "Baseline", "predict train mean", t, complexity=1)

def compare_pair(a, b):
    pa, pb = PRED[a]["test"], PRED[b]["test"]
    return {"max |Δ test prediction|": np.abs(pa - pb).max(), "corr(test predictions)": np.corrcoef(pa, pb)[0, 1],
           "Δ test RMSE": abs(RESULTS[a]["test_rmse"] - RESULTS[b]["test_rmse"]), "Δ test R²": abs(RESULTS[a]["test_r2"] - RESULTS[b]["test_r2"])}

cols = ["train_mse", "train_rmse", "train_r2", "test_mse", "test_rmse", "test_r2"]
display(pd.DataFrame(RESULTS).T.loc[[N_MEAN, N_LIN_NE, N_LIN_GD, N_LIN_SK], ["config"] + cols])

coef_cmp = pd.DataFrame({"scratch normal eq.": np.r_[FITTED[N_LIN_NE].intercept_, FITTED[N_LIN_NE].coef_],
                        "scratch GD": np.r_[FITTED[N_LIN_GD].intercept_, FITTED[N_LIN_GD].coef_],
                        "sklearn": np.r_[FITTED[N_LIN_SK].intercept_, FITTED[N_LIN_SK].coef_]}, index=["intercept"] + FEATURES)
coef_cmp["|NE - sklearn|"] = (coef_cmp["scratch normal eq."] - coef_cmp["sklearn"]).abs()
coef_cmp["|GD - sklearn|"] = (coef_cmp["scratch GD"] - coef_cmp["sklearn"]).abs()
display(coef_cmp)
display(pd.DataFrame({"scratch NE vs sklearn": compare_pair(N_LIN_NE, N_LIN_SK), "scratch GD vs sklearn": compare_pair(N_LIN_GD, N_LIN_SK)}).T)
assert np.allclose(coef_cmp["scratch normal eq."], coef_cmp["sklearn"], atol=1e-8)
h = HISTORY[N_LIN_GD]; print(f"GD: within 1% of final train loss after {epochs_to_tol(h)} epochs; final train RMSE {h['train_rmse'].iloc[-1]:.5f}")
```

	config	train_mse	train_rmse	train_r2	test_mse	test_rmse	test_r2
Baseline mean predictor	predict train mean	2.073	1.440	0.000	2.025	1.423	-0.000
Linear scratch (normal eq.)	OLS, lstsq	0.230	0.480	0.889	0.243	0.493	0.880
Linear scratch (GD)	batch GD, lr=0.15, 1000 epochs	0.230	0.480	0.889	0.243	0.493	0.880
Linear sklearn	sklearn.LinearRegression	0.230	0.480	0.889	0.243	0.493	0.880

	scratch normal eq.	scratch GD	sklearn	NE - sklearn	GD - sklearn
intercept	2.132	2.132	2.132	0.000	0.000
PT08.S1(CO)	0.394	0.394	0.394	0.000	0.000
PT08.S2(NMHC)	1.366	1.366	1.366	0.000	0.000
PT08.S3(NOx)	0.199	0.199	0.199	0.000	0.000
PT08.S4(NO2)	-0.174	-0.174	-0.174	0.000	0.000
PT08.S5(O3)	-0.077	-0.077	-0.077	0.000	0.000
T	-0.244	-0.244	-0.244	0.000	0.000
RH	-0.030	-0.030	-0.030	0.000	0.000
AH	0.087	0.087	0.087	0.000	0.000

	max Δ test prediction	corr(test predictions)	Δ test RMSE	Δ test R²
scratch NE vs sklearn	0.000	1.000	0.000	0.000
scratch GD vs sklearn	0.000	1.000	0.000	0.000

GD: within 1% of final train loss after 135 epochs; final train RMSE 0.47988

Observations — from scratch vs. scikit-learn (linear).

- The closed-form solution, the gradient-descent solution and `sklearn.LinearRegression` give the **same coefficients (agreement to < 1e-8 for the closed form)** and the same predictions (max difference ≈ 0 at the 3-decimal level shown), hence identical metrics: **train RMSE 0.480 / R² 0.889, test RMSE 0.493 / R² 0.880**.
- The mean predictor scores $R^2 \approx 0$ (test RMSE 1.423), so the sensors explain ~88 % of the variance linearly.
- Standardised coefficients already show the dominant sensor: PT08.S2(NMHC) (+1.37) $\gg$ PT08.S1(CO) (+0.39) > T (-0.24) > PT08.S3 (+0.20).
- The train–test gap is tiny (RMSE +0.013): the linear model is **bias-limited, not variance-limited**.

2.3 Polynomial regression from scratch

Feature expansion `PolynomialFeaturesScratch` builds every monomial of total degree 1...d (e.g. d=2: x_i , $x_i \cdot x_j$, x_i^2) and is verified against `sklearn.preprocessing.PolynomialFeatures`. Because polynomial columns have very different scales (and are highly collinear), the expanded matrix is re-standardised with **training-fold statistics** inside the estimator before the (closed-form or gradient-based) least-squares fit.

```
class PolynomialFeaturesScratch:
    """All monomials of total degree 1..d (no bias column); column order identical to sklearn's."""
    def __init__(self, degree): self.degree = degree
    def fit(self, X, y=None):
        self.n_features_ = np.asarray(X).shape[1]
        self.combos_ = [c for k in range(1, self.degree + 1) for c in combinations_with_replacement(range(self.n_features_), k)]
        return self
    def transform(self, X):
        X = np.asarray(X, float)
        return np.column_stack([np.prod(X[:, c], axis=1) for c in self.combos_])
    def feature_names(self, names):
        return ["*".join(names[i] for i in c) for c in self.combos_]

class PolyRegressionScratch(RegressorMixin, BaseEstimator):
    def __init__(self, degree=2, solver="normal", alpha=0.0, lr=0.05, epochs=2000, track=False):
        self.degree, self.solver, self.alpha, self.lr, self.epochs, self.track = degree, solver, alpha, lr, epochs, track
    def _expand(self, X): return (self.poly_.transform(X) - self.mu_) / self.sd_
    def fit(self, X, y, eval_sets=None):
        X = np.asarray(X, float); self.poly_ = PolynomialFeaturesScratch(self.degree).fit(X)
        P = self.poly_.transform(X); self.mu_, self.sd_ = P.mean(0), P.std(0); self.sd_[self.sd_ == 0] = 1.0
        ev = {k: (self._expand(Xe), ye) for k, (Xe, ye) in (eval_sets or {}).items()}
        self.lin_ = LinearRegressionScratch(self.solver, self.alpha, self.lr, self.epochs, track=self.track).fit((P - self.mu_) / self.sd_, y, ev)
        self.intercept_, self.coef_, self.history_ = self.lin_.intercept_, self.lin_.coef_, getattr(self.lin_, "history_", None)
        return self
    def predict(self, X): return self.lin_.predict(self._expand(np.asarray(X, float)))

# ---- verify the expansion against sklearn ----
for deg in (1, 2, 3):
    mine = PolynomialFeaturesScratch(deg).fit(X_train_s).transform(X_train_s)
    ref = PolynomialFeatures(deg, include_bias=False).fit_transform(X_train_s)
    assert mine.shape == ref.shape and np.allclose(mine, ref)
print("PolynomialFeaturesScratch == sklearn PolynomialFeatures for d = 1, 2, 3 ✓")
nfeat = {d: PolynomialFeaturesScratch(d).fit(X_train_s).transform(X_train_s[:5]).shape[1] for d in range(1, 7)}
display(pd.DataFrame({"degree": list(nfeat.keys()), "n_terms": list(nfeat.values()), "terms / training rows": [v / len(X_train_s) for v in nfeat.values()])).set_index("degree").T)
```

PolynomialFeaturesScratch == sklearn PolynomialFeatures for d = 1, 2, 3 ✓

degree	1	2	3	4	5	6
n_terms	8.000	44.000	164.000	494.000	1,286.000	3,002.000
terms / training rows	0.001	0.007	0.028	0.084	0.219	0.511

2.3.1 How the "appropriate" polynomial degree is obtained

The degree controls the bias/variance trade-off. We fit degrees 1–5 by closed-form least squares and score each with the **5-fold group CV** (training set only). We then use the **one-standard-error rule**: pick the *smallest* degree whose CV RMSE is within one standard error of the best — i.e. prefer the simpler model unless a more complex one is clearly better. The same expansion is also crossed with a Ridge strength α (an ablation addressing the multicollinearity found in Stage 1).

```

DEGREES = [1, 2, 3, 4, 5]
ALPHAS = [0, 1e-2, 1e-1, 1, 10, 100, 1e3, 1e4]          # alpha = 0 -> plain least squares (lstsq)
Xv, yv = X_train_s.values, y_train.values

def poly_cv_curve(degrees, alphas):
    """CV curve for polynomial (ridge) regression. The expansion is deterministic, so it is done once per degree;
    standardisation, fit and scoring are done per fold with training-fold statistics only (no leakage)."""
    rows = []
    for deg in degrees:
        P = PolynomialFeaturesScratch(deg).fit(Xv).transform(Xv)
        for k, (tr, va) in enumerate(folds):
            mu, sd = P[tr].mean(0), P[tr].std(0); sd[sd == 0] = 1
            A, B = (P[tr] - mu) / sd, (P[va] - mu) / sd
            ym = yv[tr].mean(); G, c = A.T @ A, A.T @ (yv[tr] - ym)      # centred columns -> intercept = mean(y)
            for a in alphas:
                w = np.linalg.lstsq(A, yv[tr] - ym, rcond=None)[0] if a == 0 else np.linalg.solve(G + a * np.eye(G.shape[0]), c)
                rows.append({"degree": deg, "alpha": a, "fold": k, "n_terms": P.shape[1],
                            "train_rmse": rmse(yv[tr], A @ w + ym), "val_rmse": rmse(yv[va], B @ w + ym), "val_r2": r2(yv[va], B @ w + ym)})
    return pd.DataFrame(rows)

t0 = time.perf_counter(); POLY_CV = poly_cv_curve(DEGREES, ALPHAS); print(f"degree x alpha CV finished in {time.perf_counter() - t0:.0f}s")
agg = POLY_CV.groupby(["degree", "alpha"]).agg(n_terms=("n_terms", "first"), train_rmse=("train_rmse", "mean"), cv_rmse=("val_rmse", "mean"),
                                              cv_rmse_sd=("val_rmse", "std"), cv_r2=("val_r2", "mean")).reset_index()
agg["cv_rmse_se"] = agg["cv_rmse_sd"] / np.sqrt(len(folds))

unreg = agg[agg["alpha"] == 0].set_index("degree")
best_d = unreg["cv_rmse"].idxmin(); thr = unreg.loc[best_d, "cv_rmse"] + unreg.loc[best_d, "cv_rmse_se"]
DEG_STAR = int(unreg[unreg["cv_rmse"] <= thr].index.min())
unreg = unreg.assign(selected=["<-- chosen (1-SE)" if d == DEG_STAR else ("best CV" if d == best_d else "") for d in unreg.index])
display(unreg[["n_terms", "train_rmse", "cv_rmse", "cv_rmse_sd", "cv_rmse_se", "cv_r2", "selected"]])
print(f"Best CV degree = {best_d}; 1-SE threshold = {thr:.4f} -> chosen degree = {DEG_STAR}")

best_row = agg.loc[agg["cv_rmse"].idxmin()]; DEG_R, ALPHA_R = int(best_row["degree"]), float(best_row["alpha"])
print(f"Best (degree, alpha) overall with Ridge: degree={DEG_R}, alpha={ALPHA_R:g}, CV RMSE={best_row['cv_rmse']:.4f}")

```

degree x alpha CV finished in 2s

	n_terms	train_rmse	cv_rmse	cv_rmse_sd	cv_rmse_se	cv_r2	selected
degree							
1	8	0.479	0.488	0.061	0.027	0.883	
2	44	0.407	0.453	0.054	0.024	0.899	<-- chosen (1-SE)
3	164	0.380	0.500	0.061	0.027	0.877	
4	494	0.341	1.228	0.809	0.362	0.076	
5	1286	0.274	8.571	13.349	5.970	-93.910	

Best CV degree = 2; 1-SE threshold = 0.4771 -> chosen degree = 2

Best (degree, alpha) overall with Ridge: degree=2, alpha=1, CV RMSE=0.4516

Observations — how the degree was chosen.


- Training RMSE falls monotonically with degree (0.479 → 0.274) — it can never be used for selection.
- The **CV RMSE is U-shaped**: 0.488 (d=1) → **0.453 (d=2)** → 0.500 (d=3) → 1.23 (d=4) → 8.6 (d=5, with enormous fold-to-fold variance). From degree 4 the number of terms (494) is large and strongly collinear, so unregularised least squares extrapolates wildly on unseen weeks.
- Degree 2 is both the CV minimum and the 1-SE choice, so **degree = 2 (44 terms)** is selected.
- Ridge confirms the diagnosis: a large enough α rescues degrees 4–5 (d=5: 8.57 → 0.52 at $\alpha=100$) but **never beats degree 2** (best overall: d=2, $\alpha=1$, CV 0.4516 vs. 0.4530 — a negligible, within-noise gain). Degree 2 needs no regularisation.

```

fig, ax = plt.subplots(1, 3, figsize=(18, 4.4))
u = unreg.reset_index()
ax[0].errorbar(u["degree"], u["cv_rmse"], yerr=u["cv_rmse_se"], color=RED, marker="o", capsize=3, label="CV (validation weeks)")
ax[0].plot(u["degree"], u["train_rmse"], color=BLUE, marker="s", label="train")
ax[0].axvline(DEG_STAR, color=GREEN, ls="--", label=f"chosen degree = {DEG_STAR}"); ax[0].set_yscale("log")
ax[0].set_xticks(DEGREES); ax[0].set_xlabel("polynomial degree"); ax[0].set_ylabel("RMSE (mg/m³, log)"); ax[0].set_title("Bias-variance: degree selection ( $\alpha = 0$ )"); ax[0].legend()

piv = agg.pivot(index="degree", columns="alpha", values="cv_rmse")
sns.heatmap(piv, annot=True, fmt=".3f", cmap="viridis_r", ax=ax[1], cbar_kws={"label": "CV RMSE"}, norm=matplotlib.colors.LogNorm())
ax[1].set_title("CV RMSE: degree  $\times$  ridge  $\alpha$ "); ax[1].set_xlabel(" $\alpha$  ( $0 =$  no regularisation)")
for deg, c in zip(DEGREES, [GREY, BLUE, GREEN, ORANGE, RED]):
    s = agg[agg["degree"] == deg]; ax[2].plot(range(len(s)), s["cv_rmse"], marker="o", color=c, label=f"d={deg}")
ax[2].set_xticks(range(len(ALPHAS))); ax[2].set_xticklabels([f"a:{g}" for a in ALPHAS]); ax[2].set_yscale("log")
ax[2].set_xlabel("ridge  $\alpha$ "); ax[2].set_ylabel("CV RMSE (log)"); ax[2].set_title("Regularisation tames high degrees"); ax[2].legend()
plt.tight_layout(); plt.show()

```



2.3.2 Final polynomial models: scratch (closed form), scratch (Adam, tracked per epoch), scikit-learn, and the ridge variant

```

N_POLY_NE, N_POLY_GD, N_POLY_SK = f"Poly(d={DEG_STAR}) | scratch (normal eq.)", f"Poly(d={DEG_STAR}) | scratch (Adam GD)", f"Poly(d={DEG_STAR}) | sklearn"
N_PR_SC, N_PR_SK = f"Poly(d={DEG_R})+Ridge(alpha={ALPHA_R:g}) | scratch", f"Poly(d={DEG_R})+Ridge(alpha={ALPHA_R:g}) | sklearn"
nterms = lambda d: sum(1 for k in range(1, d + 1) for _ in combinations_with_replacement(range(len(FEATURES)), k)) + 1

# --- gradient-based poly fit: pick Adam's learning rate by TRAIN loss only (same principle as for the linear model)
EPOCHS_POLY = 3000
adam_rows = []
for lr in [0.003, 0.01, 0.03, 0.1]:
    mm = PolyRegressionScratch(DEG_STAR, solver="adam", lr=lr, epochs=EPOCHS_POLY).fit(X_train_s, y_train)
    adam_rows.append({"lr": lr, "diverged": mm.lin_.diverged_, "train_rmse": rmse(y_train, mm.predict(X_train_s))})
adam_tbl = pd.DataFrame(adam_rows).set_index("lr"); display(adam_tbl)
LR_POLY = float(adam_tbl[~adam_tbl["diverged"]]["train_rmse"].idxmin()); print("Selected Adam lr:", LR_POLY)

m, t = timed_fit(PolyRegressionScratch(DEG_STAR, solver="normal"), X_train_s, y_train); evaluate(N_POLY_NE, m, "Polynomial", f"degree {DEG_STAR}, OLS lstsq", t, nterms(DEG_STAR))
m, t = timed_fit(PolyRegressionScratch(DEG_STAR, solver="adam", lr=LR_POLY, epochs=EPOCHS_POLY, track=True), X_train_s, y_train, eval_sets={"test": (X_test_s, y_test)})
evaluate(N_POLY_GD, m, "Polynomial", f"degree {DEG_STAR}, Adam lr={LR_POLY}, {EPOCHS_POLY} epochs", t, nterms(DEG_STAR)); HISTORY[N_POLY_GD] = m.history_

m, t = timed_fit(make_pipeline(PolynomialFeatures(DEG_STAR, include_bias=False), StandardScaler(), LinearRegression()), X_train_s, y_train)
evaluate(N_POLY_SK, m, "Polynomial", f"degree {DEG_STAR}, sklearn pipeline", t, nterms(DEG_STAR))
m, t = timed_fit(PolyRegressionScratch(DEG_R, solver="normal", alpha=ALPHA_R), X_train_s, y_train); evaluate(N_PR_SC, m, "Polynomial+Ridge", f"degree {DEG_R}, alpha {ALPHA_R:g}", t, nterms(DEG_R))
m, t = timed_fit(make_pipeline(PolynomialFeatures(DEG_R, include_bias=False), StandardScaler(), Ridge(alpha=ALPHA_R)), X_train_s, y_train)
evaluate(N_PR_SK, m, "Polynomial+Ridge", f"degree {DEG_R}, alpha {ALPHA_R:g}", t, nterms(DEG_R))

display(pd.DataFrame(RESULTS).T.loc[[N_POLY_NE, N_POLY_GD, N_POLY_SK, N_PR_SC, N_PR_SK], ["config"] + cols])
display(pd.DataFrame({"scratch NE vs sklearn": compare_pair(N_POLY_NE, N_POLY_SK), "scratch Adam vs sklearn": compare_pair(N_POLY_GD, N_POLY_SK),
    "ridge scratch vs sklearn": compare_pair(N_PR_SC, N_PR_SK)}).T)
h = HISTORY[N_POLY_GD]; print(f"Adam: within 1% of its final train loss after {epochs_to_tol(h)} epochs; final train RMSE {h['train_rmse'].iloc[-1]:.4f} "
    f"(closed-form optimum {RESULTS[N_POLY_NE]['train_rmse']:.4f})")

```

	diverged	train_rmse
lr		
0.003	False	0.415
0.010	False	0.411
0.030	False	0.411
0.100	False	0.420

Selected Adam lr: 0.03

	config	train_rmse	train_rmse	train_r2	test_rmse	test_rmse	test_r2
Poly(d=2) scratch (normal eq.)	degree 2, OLS lstsq	0.169	0.411	0.918	0.200	0.447	0.901
Poly(d=2) scratch (Adam GD)	degree 2, Adam lr=0.03, 3000 epochs	0.169	0.411	0.918	0.200	0.447	0.901
Poly(d=2) sklearn	degree 2, sklearn pipeline	0.169	0.411	0.918	0.200	0.447	0.901
Poly(d=2)+Ridge(α=1) scratch	degree 2, alpha 1	0.169	0.411	0.918	0.200	0.447	0.901
Poly(d=2)+Ridge(α=1) sklearn	degree 2, alpha 1	0.169	0.411	0.918	0.200	0.447	0.901

	max Δ test prediction	corr(test predictions)	Δ test RMSE	Δ test R²
scratch NE vs sklearn	0.000	1.000	0.000	0.000
scratch Adam vs sklearn	0.001	1.000	0.000	0.000
ridge scratch vs sklearn	0.000	1.000	0.000	0.000

Adam: within 1% of its final train loss after 638 epochs; final train RMSE 0.4112 (closed-form optimum 0.4112)

Observations — from scratch vs. scikit-learn (polynomial, d = 2).

- Closed-form scratch, Adam-trained scratch and the sklearn pipeline agree on every metric: **train RMSE 0.411 / R² 0.918, test RMSE 0.447 / R² 0.901** (maximum prediction difference for Adam: 0.0013 mg/m³ after 3,000 epochs).
- Adam converges (within 1 % of its final loss) in ~640 epochs; the expanded features are highly correlated, which is why an adaptive optimiser was used instead of plain GD.
- The ridge variant (d=2, α=1) is practically the same model (test RMSE 0.4474 vs. 0.4470); scratch and sklearn ridge also agree.
- Versus the linear model the polynomial lowers test RMSE from 0.493 to **0.447 (–9 %)** and raises R² from 0.880 to 0.901, at the price of a somewhat larger train–test gap (0.036 vs. 0.013).

2.4 Decision Tree Regressor

Trees are scale-invariant, so they are fed the same standardised matrix as the other models (verified below), which keeps the comparison like-for-like. We (i) run the **required max_depth experiment** (3, 5, 10, None), (ii) a fine **depth sweep 1–20** with both CV and test curves for the metric-vs-depth plots, then (iii) tune **max_depth × min_samples_leaf** by group-CV with the 1-SE rule, and (iv) try cost-complexity pruning as an alternative regulariser.

```
# scale-invariance check: same tree on raw vs standardised features gives identical predictions
t_raw = DecisionTreeRegressor(max_depth=5, random_state=SEED).fit(X_train, y_train)
t_std = DecisionTreeRegressor(max_depth=5, random_state=SEED).fit(X_train_s, y_train)
assert np.allclose(t_raw.predict(X_test), t_std.predict(X_test_s)); print("Tree predictions identical on raw and standardised features ✓\n")

REQ_DEPTHS = [3, 5, 10, None]
tree_names = {}
for d in REQ_DEPTHS:
    nm = f"Tree | max_depth={d}"; tree_names[d] = nm
    m, t = timed_fit(DecisionTreeRegressor(max_depth=d, random_state=SEED), X_train_s, y_train)
    evaluate(nm, m, "Decision Tree", f"max_depth={d}", t, complexity=m.get_n_leaves())
    RESULTS[nm].update({"depth_actual": m.get_depth(), "n_leaves": m.get_n_leaves(), **dict(cv_eval(DecisionTreeRegressor(max_depth=d,
    random_state=SEED))))})
req = pd.DataFrame(RESULTS).T.loc[list(tree_names.values()), ["depth_actual", "n_leaves"] + cols + ["cv_rmse", "cv_rmse_sd", "cv_r2"]]
display(req)
```

Tree predictions identical on raw and standardised features ✓

	depth_actual	n_leaves	train_rmse	train_rmse	train_r2	test_rmse	test_rmse	test_r2	cv_rmse	cv_rmse_sd	cv_r2
Tree max_depth=3	3	8	0.304	0.552	0.853	0.343	0.585	0.831	0.577	0.038	0.835
Tree max_depth=5	5	32	0.194	0.440	0.906	0.263	0.513	0.870	0.509	0.051	0.873
Tree max_depth=10	10	693	0.068	0.260	0.967	0.296	0.544	0.854	0.598	0.041	0.824
Tree max_depth=None	28	4245	0.000	0.000	1.000	0.372	0.610	0.817	0.659	0.055	0.786

Observations — the required max_depth experiment.

max_depth	leaves	train RMSE	test RMSE	test R ²
3	8	0.552	0.585	0.831
5	32	0.440	0.513	0.870
10	693	0.260	0.544	0.854
None (grew to depth 28)	4,245	0.000	0.610	0.817



Depth 3 **under-fits** (high bias), depth 5 balances bias and variance, and depth 10 / None **over-fit**: the unrestricted tree memorises the training set ($R^2 = 1.000$, one leaf per ~ 1.4 samples) while test error is 19 % worse than at depth 5. CV agrees with the test ordering, i.e. the conclusion does not rely on the test set.

```
DEPTH_GRID = list(range(1, 21)) + [None]
rows = []
for d in DEPTH_GRID:
    m = DecisionTreeRegressor(max_depth=d, random_state=SEED).fit(X_train_s, y_train)
    r = {"max_depth": "None" if d is None else d, "depth_actual": m.get_depth(), "n_leaves": m.get_n_leaves()}
    for split, X_, y_ in (("train", X_train_s, y_train), ("test", X_test_s, y_test)):
        r.update({f"{split}_{k}": v for k, v in regression_metrics(y_, m.predict(X_)).items()})
    r.update(dict(cv_eval(DecisionTreeRegressor(max_depth=d, random_state=SEED))))
    rows.append(r)
TREE_SWEEP = pd.DataFrame(rows)
display(TREE_SWEEP[["max_depth", "depth_actual", "n_leaves", "train_rmse", "cv_rmse", "test_rmse", "train_r2", "cv_r2",
"test_r2"]].style.format(precision=4)
        .highlight_min(subset=["cv_rmse"], color="#cfe8cf"))
d_best_cv = TREE_SWEEP.loc[TREE_SWEEP["cv_rmse"].idxmin(), "depth_actual"]
print("Depth with the lowest CV RMSE:", d_best_cv)

fig, ax = plt.subplots(1, 2, figsize=(14, 4.2)); x = TREE_SWEEP["depth_actual"]
ax[0].plot(x, TREE_SWEEP["train_rmse"], color=BLUE, marker="o", ms=4, label="train"); ax[0].plot(x, TREE_SWEEP["cv_rmse"], color=RED, marker="o", ms=4,
label="CV (validation weeks)")
ax[0].plot(x, TREE_SWEEP["test_rmse"], color=ORANGE, marker="o", ms=4, label="test (report only)"); ax[0].axvline(d_best_cv, color=GREEN, ls="--",
label=f"CV-best depth = {d_best_cv}")
ax[0].set_xlabel("effective tree depth (last point = max_depth=None)"); ax[0].set_ylabel("RMSE (mg/m³)"); ax[0].set_title("Decision tree: RMSE vs
depth"); ax[0].legend()
ax[1].plot(x, TREE_SWEEP["n_leaves"], color=GREY, marker="o", ms=4); ax[1].set_yscale("log"); ax[1].set_xlabel("effective tree depth");
ax[1].set_ylabel("number of leaves (log)")
ax[1].set_title("Model complexity vs depth"); plt.tight_layout(); plt.show()
```

	max_depth	depth_actual	n_leaves	train_rmse	cv_rmse	test_rmse	train_r2	cv_r2	test_r2
0	1	1	2	0.9351	0.9423	0.9336	0.5781	0.5658	0.5696
1	2	2	4	0.6629	0.6739	0.6757	0.7880	0.7769	0.7746
2	3	3	8	0.5515	0.5773	0.5854	0.8532	0.8347	0.8308
3	4	4	16	0.4864	0.5263	0.5366	0.8859	0.8643	0.8578
4	5	5	32	0.4403	0.5092	0.5131	0.9065	0.8730	0.8700
5	6	6	64	0.4094	0.4998	0.5049	0.9191	0.8771	0.8741
6	7	7	125	0.3802	0.5470	0.5014	0.9303	0.8527	0.8758
7	8	8	237	0.3450	0.5643	0.5015	0.9426	0.8420	0.8758
8	9	9	417	0.3032	0.5793	0.5399	0.9556	0.8338	0.8561
9	10	10	693	0.2603	0.5981	0.5442	0.9673	0.8236	0.8538
10	11	11	1069	0.2184	0.6321	0.5648	0.9770	0.8033	0.8425
11	12	12	1522	0.1786	0.6398	0.5853	0.9846	0.7990	0.8308
12	13	13	2004	0.1444	0.6380	0.5828	0.9899	0.7986	0.8323
13	14	14	2470	0.1165	0.6436	0.5869	0.9934	0.7956	0.8299
14	15	15	2933	0.0920	0.6505	0.5798	0.9959	0.7906	0.8340
15	16	16	3324	0.0722	0.6552	0.6015	0.9975	0.7875	0.8213
16	17	17	3584	0.0567	0.6544	0.6020	0.9984	0.7899	0.8210
17	18	18	3823	0.0428	0.6443	0.5776	0.9991	0.7941	0.8353
18	19	19	3968	0.0317	0.6633	0.5928	0.9995	0.7837	0.8265
19	20	20	4076	0.0231	0.6620	0.5842	0.9997	0.7848	0.8315
20	None	28	4245	0.0000	0.6593	0.6096	1.0000	0.7859	0.8165

Depth with the lowest CV RMSE: 6



Observations. The fine sweep shows the classic U-shape: CV RMSE is lowest at **depth 6 (0.500)** and rises steadily afterwards, while training RMSE keeps falling to 0 and the leaf count explodes past 10^3 . The test curve is flat around depths 6–8 (≈ 0.50) and degrades beyond. (The test curve is displayed for the later evaluation plots only; selection uses the CV curve.)

```

# --- joint tuning of max_depth x min_samples_leaf with group CV, then the 1-SE rule for a parsimonious tree
grid = {"max_depth": [3, 4, 5, 6, 7, 8, 10, 12, 15, None], "min_samples_leaf": [1, 3, 5, 10, 20, 50]}
gs = GridSearchCV(DecisionTreeRegressor(random_state=SEED), grid, cv=cv, scoring="neg_root_mean_squared_error", n_jobs=-1, return_train_score=True)
gs.fit(X_train_s, y_train, groups=groups)
P_ = gs.cv_results_["params"]
G = pd.DataFrame({"max_depth": pd.Series([p["max_depth"] for p in P_], dtype=object), "min_samples_leaf": [p["min_samples_leaf"] for p in P_],
                  "train_rmse": -gs.cv_results_["mean_train_score"], "cv_rmse": -gs.cv_results_["mean_test_score"], "cv_rmse_sd":
                  gs.cv_results_["std_test_score"]})
G["cv_rmse_se"] = G["cv_rmse_sd"] / np.sqrt(len(folds)); G["_d"] = G["max_depth"].map(lambda v: 99 if v is None else int(v))
G["depth_lbl"] = G["max_depth"].map(lambda v: "None" if v is None else str(int(v)))
best = G.loc[G["cv_rmse"].idxmin()]
simple = G[G["cv_rmse"] <= best["cv_rmse"] + best["cv_rmse_se"]].sort_values(["_d", "min_samples_leaf"], ascending=[True, False]).iloc[0]
TREE_BEST = {"max_depth": None if simple["_d"] == 99 else int(simple["_d"]), "min_samples_leaf": int(simple["min_samples_leaf"])}
print(f"Lowest CV RMSE : depth={best['depth_lbl']}, min_samples_leaf={int(best['min_samples_leaf'])} CV RMSE={best['cv_rmse']:.4f}")
print(f"1-SE choice : {TREE_BEST} CV RMSE={simple['cv_rmse']:.4f} (threshold {best['cv_rmse'] + best['cv_rmse_se']:.4f})")

piv = G.pivot(index="depth_lbl", columns="min_samples_leaf", values="cv_rmse").loc[["None" if v is None else str(v) for v in grid["max_depth"]]]
plt.figure(figsize=(8, 4.5)); sns.heatmap(piv, annot=True, fmt=".3f", cmap="viridis_r", cbar_kws={"label": "CV RMSE"}); plt.title("Tree tuning grid (group-CV RMSE)")
plt.xlabel("min_samples_leaf"); plt.ylabel("max_depth"); plt.tight_layout(); plt.show()


N_TREE_T = "Tree | tuned (CV, 1-SE)"
m, t = timed_fit(DecisionTreeRegressor(*TREE_BEST, random_state=SEED), X_train_s, y_train)
evaluate(N_TREE_T, m, "Decision Tree", str(TREE_BEST), t, complexity=m.get_n_leaves())
RESULTS[N_TREE_T].update({"depth_actual": m.get_depth(), "n_leaves": m.get_n_leaves(), **dict(cv_eval(DecisionTreeRegressor(*TREE_BEST,
random_state=SEED))))})

# --- cost-complexity pruning as an alternative regulariser
full = DecisionTreeRegressor(random_state=SEED).fit(X_train_s, y_train)
path = full.cost_complexity_pruning_path(X_train_s, y_train)
cand = np.unique(np.quantile(path.ccp_alphas[:-1], np.linspace(0, 0.995, 30)))
ccp = pd.DataFrame([{"ccp_alpha": a, **cv_eval(DecisionTreeRegressor(ccp_alpha=a, random_state=SEED)),
                    "n_leaves": DecisionTreeRegressor(ccp_alpha=a, random_state=SEED).fit(X_train_s, y_train).get_n_leaves()} for a in cand])
best_c = ccp.loc[ccp["cv_rmse"].idxmin()]; CCP_BEST = float(best_c["ccp_alpha"])
print(f"Best ccp_alpha = {CCP_BEST:.5f} (leaves {int(best_c['n_leaves'])}, CV RMSE {best_c['cv_rmse']:.4f})")
fig, ax = plt.subplots(figsize=(7, 4)); ax.semilogx(ccp["ccp_alpha"].clip(lower=1e-6), ccp["cv_rmse"], marker="o", color=RED, label="CV RMSE")
ax.semilogx(ccp["ccp_alpha"].clip(lower=1e-6), ccp["cv_train_rmse"], marker="s", color=BLUE, label="train RMSE"); ax.axvline(max(CCP_BEST, 1e-6),
color=GREEN, ls="—")
ax.set_xlabel("ccp_alpha"); ax.set_ylabel("RMSE"); ax.set_title("Cost-complexity pruning path"); ax.legend(); plt.tight_layout(); plt.show()
N_TREE_P = "Tree | pruned (ccp, CV)"
m, t = timed_fit(DecisionTreeRegressor(ccp_alpha=CCP_BEST, random_state=SEED), X_train_s, y_train)
evaluate(N_TREE_P, m, "Decision Tree", f"ccp_alpha={CCP_BEST:.5f}", t, complexity=m.get_n_leaves())
RESULTS[N_TREE_P].update({"depth_actual": m.get_depth(), "n_leaves": m.get_n_leaves(), **dict(cv_eval(DecisionTreeRegressor(ccp_alpha=CCP_BEST,
random_state=SEED))))})


```

Lowest CV RMSE : depth=6, min_samples_leaf=20 CV RMSE=0.4926

1-SE choice : {'max_depth': 5, 'min_samples_leaf': 50} CV RMSE=0.5134 (threshold 0.5143)



Best ccp_alpha = 0.00281 (leaves 21, CV RMSE 0.5212)



Observations.

- The lowest CV RMSE (0.4926) is at depth 6 with min_samples_leaf = 20. The **1-SE rule** picks the simpler **depth 5**, **min_samples_leaf = 50** (CV 0.5134), giving 29 leaves and a better-regularised tree (train-test gap 0.046 vs. 0.073 for the plain depth-5 tree).
- Cost-complexity pruning reaches a similar optimum (21 leaves, CV 0.521): the two regularisers agree that the useful structure of this problem is captured by roughly 20–30 leaves.
- Tuned and pruned trees do **not** beat the plain depth-5 tree on test RMSE (0.516 / 0.530 vs. 0.513) — differences of this size are within the fold noise (CV sd ≈ 0.05) — but they are more stable (smaller gap) and cheaper to interpret.

2.5 Master results table (all models, train and test) and the from-scratch vs scikit-learn verdict

CV scores are added for the models whose hyper-parameters were chosen by CV, so that Stage 5 can compare on *selection-safe* numbers as well as on the test set.

```
# add group-CV scores for the remaining headline models (cheap closed-form / sklearn fits)
for nm, est in {N_MEAN: DummyRegressor(), N_LIN_NE: LinearRegressionScratch(), N_POLY_NE: PolyRegressionScratch(DEG_STAR),
               N_PR_SC: PolyRegressionScratch(DEG_R, alpha=ALPHA_R)}.items():
    RESULTS[nm].update(dict(cv_eval(est)))

ORDER = [N_MEAN, N_LIN_NE, N_LIN_GD, N_LIN_SK, N_POLY_NE, N_POLY_GD, N_POLY_SK, N_PR_SC, N_PR_SK,
         tree_names[3], tree_names[5], tree_names[10], tree_names[None], N_TREE_T, N_TREE_P]
MASTER = pd.DataFrame(RESULTS).T.loc[ORDER].drop(columns=["depth_actual"], errors="ignore")
num = [c for c in MASTER.columns if c not in ("model", "family", "config", "complexity")]
MASTER[num] = MASTER[num].astype(float)
show = ["family", "config", "complexity", "train_mse", "train_rmse", "train_r2", "test_mse", "test_rmse", "test_r2", "rmse_gap (test-train)",
        "cv_rmse", "cv_r2"]
display(MASTER[show].style.format(precision=4, na_rep="--").background_gradient(subset=["test_rmse"], cmap="RdYlGn_r")
        .background_gradient(subset=["test_r2"], cmap="RdYlGn"))

# scratch-vs-sklearn verdict
pairs = [(("Linear", N_LIN_NE, N_LIN_SK), ("Linear (GD)", N_LIN_GD, N_LIN_SK), (f"Poly d={DEG_STAR}", N_POLY_NE, N_POLY_SK),
         (f"Poly d={DEG_STAR} (Adam)", N_POLY_GD, N_POLY_SK), (f"Poly d={DEG_R}+Ridge", N_PR_SC, N_PR_SK))]
verdict = pd.DataFrame({lbl: {**compare_pair(a, b), "scratch test RMSE": RESULTS[a]["test_rmse"], "sklearn test RMSE": RESULTS[b]["test_rmse"],
                                "scratch test R^2": RESULTS[a]["test_r2"], "sklearn test R^2": RESULTS[b]["test_r2"]} for lbl, a, b in pairs}).T
display(verdict.style.format("{:.3e}", subset=["max |Δ test prediction|", "Δ test RMSE", "Δ test R^2"]).format("{:.5f}", subset=["corr(test
predictions)", "scratch test RMSE", "sklearn test RMSE", "scratch test R^2", "sklearn test R^2"])))
```

	family	config	complexity	train_mse	train_rmse	train_r2	test_mse	test_rmse	test_r2	rmse_gap (test-train)	cv_rmse	cv_r2
Baseline mean predictor	Baseline	predict train mean	1	2.0726	1.4397	0.0000	2.0255	1.4232	0.0000	-0.0165	1.4378	-0.0100
Linear scratch (normal eq.)	Linear	OLS, lstsq	9	0.2303	0.4799	0.8889	0.2433	0.4833	0.8799	0.0134	0.4882	0.8827
Linear scratch (GD)	Linear	batch GD, lr=0.15, 1000 epochs	9	0.2303	0.4799	0.8889	0.2433	0.4833	0.8799	0.0134	--	--
Linear sklearn	Linear	sklearn.LinearRegression	9	0.2303	0.4799	0.8889	0.2433	0.4833	0.8799	0.0134	--	--
Poly(d=2) scratch (normal eq.)	Polynomial	degree 2, OLS lstsq	45	0.1691	0.4112	0.9184	0.1998	0.4470	0.9013	0.0359	0.4530	0.8988
Poly(d=2) scratch (Adam GD)	Polynomial	degree 2, Adam lr=0.03, 3000 epochs	45	0.1691	0.4112	0.9184	0.1998	0.4470	0.9013	0.0359	--	--
Poly(d=2) sklearn	Polynomial	degree 2, sklearn pipeline	45	0.1691	0.4112	0.9184	0.1998	0.4470	0.9013	0.0359	--	--
Poly(d=2)+Ridge(alpha=1) scratch	Polynomial+Ridge	degree 2, alpha 1	45	0.1692	0.4113	0.9184	0.2001	0.4474	0.9012	0.0360	0.4516	0.8994
Poly(d=2)+Ridge(alpha=1) sklearn	Polynomial+Ridge	degree 2, alpha 1	45	0.1692	0.4113	0.9184	0.2001	0.4474	0.9012	0.0360	--	--
Tree max_depth=3	Decision Tree	max_depth=3	8	0.3042	0.5515	0.8532	0.3427	0.5854	0.8308	0.0339	0.5773	0.8347
Tree max_depth=5	Decision Tree	max_depth=5	32	0.1938	0.4403	0.9065	0.2633	0.5131	0.8700	0.0728	0.5092	0.8730
Tree max_depth=10	Decision Tree	max_depth=10	693	0.0678	0.2603	0.9673	0.2961	0.5442	0.8538	0.2838	0.5981	0.8236
Tree max_depth=None	Decision Tree	max_depth=None	4245	0.0000	0.0000	1.0000	0.3716	0.5096	0.8185	0.6096	0.6593	0.7859
Tree tuned (CV, 1-SE)	Decision Tree	{'max_depth': 5, 'min_samples_leaf': 50}	29	0.2207	0.4698	0.8935	0.2659	0.5157	0.8687	0.0459	0.5134	0.8709
Tree pruned (ccp, CV)	Decision Tree	ccp_alpha=0.00281	21	0.2139	0.4625	0.8968	0.2813	0.5304	0.8611	0.0679	0.5212	0.8670

	max Δ test prediction	corr(test predictions)	Δ test RMSE	Δ test R²	scratch test RMSE	sklearn test RMSE	scratch test R²	sklearn test R²
Linear	7.105e-15	1.00000	1.110e-16	0.000e+00	0.49327	0.49327	0.87986	0.87986
Linear (GD)	1.091e-05	1.00000	1.248e-07	6.079e-08	0.49327	0.49327	0.87986	0.87986
Poly d=2	2.220e-14	1.00000	1.110e-16	1.110e-16	0.44702	0.44702	0.90134	0.90134
Poly d=2 (Adam)	1.294e-03	1.00000	2.980e-06	1.316e-06	0.44702	0.44702	0.90134	0.90134
Poly d=2+Ridge	3.144e-13	1.00000	1.221e-15	5.551e-16	0.44737	0.44737	0.90118	0.90118

Observations.

- **Scratch ≡ scikit-learn.** For every model family and solver, the from-scratch implementations reproduce scikit-learn to numerical precision (Δ test RMSE $\approx$ 0, correlation of predictions = 1.0), which validates the from-scratch code.
- **Ranking by test RMSE:** Poly d=2 (0.447) < Linear (0.493) < Tree depth 5 (0.513) $\approx$ tuned tree (0.516) < pruned tree (0.530) < depth 10 (0.544) < depth 3 (0.585) < unrestricted tree (0.610) $\ll$ mean baseline (1.423). The **CV ranking is identical in order for the headline models**, so the ordering is not a test-set artefact.
- The polynomial model is the most accurate; the linear model is a strong, simple baseline (R^2 0.88); trees are competitive only when depth-limited.

2.6 Model-interpretation artifacts (for Stages 3–4)

Three complementary importance views are stored, because impurity-based importances are known to be biased toward high-cardinality / correlated features and Stage 1 showed strong collinearity among the sensors:

1. **tree impurity (MDI) importances** for every required depth;
2. **permutation importance** on the test set (model-agnostic, in RMSE units) for the linear, polynomial and tree models;
3. **standardised linear coefficients with bootstrap 95 % CIs** (500 resamples of the training set).

```
# 1) impurity importances
IMP_TREE = pd.DataFrame({nm: FITTED[nm].feature_importances_ for nm in [tree_names[3], tree_names[5], tree_names[10], tree_names[None], N_TREE_T,
N_TREE_P]}, index=FEATURES)

# 2) permutation importance on the test set (increase in RMSE when a feature is shuffled)
PERM = {}
for nm in [N_LIN_NE, N_POLY_NE, tree_names[3], tree_names[5], N_TREE_T]:
    pi = permutation_importance(FITTED[nm], X_test_s, y_test, scoring="neg_root_mean_squared_error", n_repeats=20, random_state=SEED)
    PERM[nm] = pd.DataFrame({"mean": pi.importances_mean, "std": pi.importances_std}, index=FEATURES)

# 3) bootstrap CIs of the standardised linear coefficients
rb = np.random.default_rng(SEED); B = []
for _ in range(500):
    idx = rb.integers(0, len(X_train_s), len(X_train_s)); B.append(LinearRegressionScratch().fit(X_train_s.values[idx], y_train.values[idx]).coef_)
B = np.array(B)
COEF_BOOT = pd.DataFrame({"coef": FITTED[N_LIN_NE].coef_, "ci_low": np.percentile(B, 2.5, 0), "ci_high": np.percentile(B, 97.5, 0)}, index=FEATURES)
COEF_BOOT["significant (CI excludes 0)"] = (COEF_BOOT["ci_low"] > 0) | (COEF_BOOT["ci_high"] < 0)


display(IMP_TREE.style.format("{:.3f}").background_gradient(cmap="Blues", axis=0))
display(PERM[N_TREE_T].sort_values("mean", ascending=False).T.style.format("{:.3f}"))
display(COEF_BOOT.style.format({"coef": "{:.3f}", "ci_low": "{:.3f}", "ci_high": "{:.3f}"}))

fig, ax = plt.subplots(1, 3, figsize=(18, 4.4))
IMP_TREE[[tree_names[3], tree_names[5], N_TREE_T]].sort_values(tree_names[5]).plot.barh(ax=ax[0], color=[BLUE, ORANGE, GREEN]); ax[0].set_title("Tree
impurity importance"); ax[0].legend(fontsize=7)
pn = PERM[N_TREE_T]["mean"].sort_values(); ax[1].barh(pn.index, pn.values, xerr=PERM[N_TREE_T].loc[pn.index, "std"], color=GREEN);
ax[1].set_title("Permutation importance (tuned tree, test)"); ax[1].set_xlabel("ΔRMSE when shuffled")
cb = COEF_BOOT.sort_values("coef"); ax[2].barh(cb.index, cb["coef"], xerr=[cb["coef"] - cb["ci_low"], cb["ci_high"] - cb["coef"]], color=BLUE);
ax[2].axvline(0, color="k", lw=.8)
ax[2].set_title("Linear coefficients (standardised) ± 95% bootstrap CI"); plt.tight_layout(); plt.show()
```

	Tree max_depth=3	Tree max_depth=5	Tree max_depth=10	Tree max_depth=None	Tree tuned (CV, 1-SE)	Tree pruned (ccp, CV)
PT08.S1(CO)	0.000	0.003	0.013	0.017	0.002	0.002
PT08.S2(NMHC)	1.000	0.958	0.904	0.880	0.970	0.958
PT08.S3(NOx)	0.000	0.007	0.014	0.016	0.000	0.007
PT08.S4(NO2)	0.000	0.004	0.011	0.014	0.000	0.004
PT08.S5(O3)	0.000	0.000	0.008	0.012	0.000	0.000
T	0.000	0.027	0.036	0.039	0.028	0.027
RH	0.000	0.000	0.007	0.011	0.000	0.000
AH	0.000	0.000	0.007	0.011	0.000	0.000

	PT08.S2(NMHC)	T	PT08.S1(CO)	PT08.S3(NOx)	PT08.S4(NO2)	PT08.S5(O3)	RH	AH
mean	1.418	0.092	0.003	0.000	0.000	0.000	0.000	0.000
std	0.029	0.007	0.001	0.000	0.000	0.000	0.000	0.000

	coef	ci_low	ci_high	significant (CI excludes 0)
PT08.S1(CO)	0.394	0.354	0.437	True
PT08.S2(NMHC)	1.366	1.242	1.484	True
PT08.S3(NOx)	0.199	0.156	0.244	True
PT08.S4(NO2)	-0.174	-0.241	-0.101	True
PT08.S5(O3)	-0.077	-0.121	-0.031	True
T	-0.244	-0.281	-0.206	True
RH	-0.030	-0.061	0.003	False
AH	0.087	0.047	0.125	True



Observations — feature importance (three views agree).

- PT08.S2(NMHC) dominates everywhere:** 97 % of the tuned tree's impurity importance, a permutation importance of +1.42 mg/m³ RMSE when shuffled (the next feature, T, is only +0.09), and the largest linear coefficient (1.37, CI 1.24–1.48). The depth-3 tree splits **only** on this sensor.
- T is the clear second feature for the tree (≈ 3 %) and the third-largest linear effect (-0.24 , CI excludes 0). PT08.S1(CO) matters linearly (+0.39) but almost never in the tree: it is redundant with PT08.S2 ($r = 0.89$, Stage 1), so the tree gives it ~ 0 importance even though it is individually predictive — a collinearity effect, not irrelevance.
- Only RH has a coefficient whose 95 % bootstrap CI includes zero.
- Caveat: with highly collinear sensors, importance is shared/arbitrarily assigned among them; the drop-one ablation below is the more reliable check.

2.7 Ablation studies

All ablations are scored with the **group-CV on the training set** (the test set is not used to make any choice here). They feed the Stage 5 baseline comparison.

2.7.1 Feature ablation — which inputs does each model family actually need? (all features, sensors only, weather only, one best sensor, and leave-one-feature-out.)

```


def make_linear(): return LinearRegressionScratch()
def make_poly(): return PolyRegressionScratch(DEG_STAR)
def make_tree(): return DecisionTreeRegressor(**TREE_BEST, random_state=SEED)
FACT = {"Linear": make_linear, f"Poly(d={DEG_STAR})": make_poly, "Tree (tuned)": make_tree}

subsets = {"All 8 features": FEATURES, "Sensors only (5)": SENSOR_COLS, "Weather only (T, RH, AH)": WEATHER_COLS, "Best single sensor (PT08.S2)":
["PT08.S2(NMHC)"]}
subsets.update({f"Drop {f}": [c for c in FEATURES if c != f] for f in FEATURES})
rows = []
for sname, cols_ in subsets.items():
    for mname, mk in FACT.items():
        r = cv_eval(mk(), X=X_train_s[cols_]); rows.append({"feature set": sname, "model": mname, "n_features": len(cols_), "cv_rmse": r["cv_rmse"],
"cv_rmse_sd": r["cv_rmse_sd"], "cv_r2": r["cv_r2"]})
ABL_FEAT = pd.DataFrame(rows)
tbl = ABL_FEAT.pivot(index="feature set", columns="model", values="cv_rmse").loc[list(subsets)]
base = tbl.loc["All 8 features"]; delta = (tbl - base)
display(pd.concat({"CV RMSE": tbl, "Δ vs all features": delta}, axis=1).style.format(precision=4).background_gradient(cmap="RdYlGn_r",
subset=pd.IndexSlice[:, "Δ vs all features"])))

fig, ax = plt.subplots(figsize=(12, 4.2)); w = 0.27; xi = np.arange(len(tbl))
for i, (mname, c) in enumerate(zip(tbl.columns, [BLUE, ORANGE, GREEN])): ax.bar(xi + (i - 1) * w, tbl[mname], w, color=c, label=mname)
ax.set_xticks(xi); ax.set_xticklabels(tbl.index, rotation=35, ha="right"); ax.set_ylabel("CV RMSE (mg/m³)"); ax.set_title("Feature ablation (group-CV,
training set)"); ax.legend()
plt.tight_layout(); plt.show()

```

feature set	CV RMSE			Δ vs all features		
	Linear	Poly(d=2)	Tree (tuned)	Linear	Poly(d=2)	Tree (tuned)
All 8 features	0.4882	0.4530	0.5134	0.0000	0.0000	0.0000
Sensors only (5)	0.5025	0.4687	0.5276	0.0143	0.0157	0.0143
Weather only (T, RH, AH)	1.4423	1.4557	1.4709	0.9541	1.0027	0.9575
Best single sensor (PT08.S2)	0.5750	0.5196	0.5341	0.0868	0.0666	0.0208
Drop PT08.S1(CO)	0.5101	0.4581	0.5046	0.0220	0.0051	-0.0088
Drop PT08.S2(NMHC)	0.6072	0.5241	0.6639	0.1191	0.0710	0.1505
Drop PT08.S3(NOx)	0.4949	0.4524	0.5108	0.0068	-0.0006	-0.0026
Drop PT08.S4(NO2)	0.4895	0.4596	0.5134	0.0013	0.0065	0.0000
Drop PT08.S5(O3)	0.4862	0.4491	0.5126	-0.0019	-0.0040	-0.0008
Drop T	0.4924	0.4502	0.5304	0.0043	-0.0029	0.0170
Drop RH	0.4877	0.4483	0.5134	-0.0004	-0.0048	0.0000
Drop AH	0.4875	0.4525	0.5134	-0.0006	-0.0005	0.0000



Observations — feature ablation (CV RMSE).

- Weather alone is useless** (≈ 1.44 – 1.47 , i.e. the same as the mean predictor, 1.44): the weather variables carry almost no *direct* information about CO.
- Sensors alone are almost as good as everything** ($+0.014$ – 0.016 RMSE): weather adds only a small correction (mostly via T, as the linear coefficient suggests).
- PT08.S2 is the one indispensable feature**: dropping it hurts every model ($+0.12$ linear, $+0.07$ polynomial, **$+0.15$ tree**), whereas dropping any other single feature changes CV RMSE by ≤ 0.02 (and sometimes improves it slightly): the sensors are mutually redundant.
- A single feature (PT08.S2) already gets RMSE 0.52–0.58**, and the tree is least dependent on the other features (0.534 with S2 only).

- Dropping T hurts the tree (+0.017) but not the polynomial (−0.003): a possible reason (not tested) is that the polynomial can recover part of that information through sensor interaction terms.


2.7.2 Preprocessing ablations — (a) is standardisation really needed for gradient descent? (b) does a log-transformed target help (Stage 1 flagged the skew)? (c) does the tree split criterion matter?

```
# (a) GD on RAW (unscaled) features vs standardised features, same epoch budget; lr = half the stability limit of each design matrix
Xr = LinearRegressionScratch._design(X_train.values); lam_r = np.linalg.eigvalsh(Xr.T @ Xr / len(Xr)).max(); lr_raw = 0.5 * 2 / (2 * lam_r)
raw = LinearRegressionScratch(solver="gd", lr=lr_raw, epochs=EPOCHS_LIN, track=True).fit(X_train, y_train)
sc = LinearRegressionScratch(solver="gd", lr=LR_LIN, epochs=EPOCHS_LIN, track=True).fit(X_train_s, y_train)
opt_rmse = RESULTS[N_LIN_NE]["train_rmse"]
ABL_SCALE = pd.DataFrame({"raw features": {"stable lr": lr_raw, "train RMSE @ end": raw.history_["train_rmse"].iloc[-1], "gap to OLS optimum":
raw.history_["train_rmse"].iloc[-1] - opt_rmse},
                          "standardised": {"stable lr": LR_LIN, "train RMSE @ end": sc.history_["train_rmse"].iloc[-1], "gap to OLS optimum":
sc.history_["train_rmse"].iloc[-1] - opt_rmse}}).T
display(ABL_SCALE.style.format("{:.6g}"))
fig, ax = plt.subplots(figsize=(7, 3.8)); ax.plot(raw.history_["epoch"], raw.history_["train_rmse"], color=RED, label=f"raw (lr={lr_raw:.1e})")
ax.plot(sc.history_["epoch"], sc.history_["train_rmse"], color=GREEN, label=f"standardised (lr={LR_LIN})"); ax.axhline(opt_rmse, color="k", ls=":",
label="OLS optimum")
ax.set_yscale("log"); ax.set_xscale("symlog", linthresh=1); ax.set_xlim(0, EPOCHS_LIN); ax.set_xlabel("epoch"); ax.set_ylabel("train RMSE");
ax.set_title("Why standardise: GD convergence"); ax.legend(); plt.tight_layout(); plt.show()

# (b) log-target
log_models = {"Linear": make_linear, f"Poly(d={DEG_STAR})": make_poly, "Tree (tuned)": make_tree}
rows = []
for mname, mk in log_models.items():
    base_ = cv_eval(mk()); logd = cv_eval(TransformedTargetRegressor(regressor=mk(), func=np.log1p, inverse_func=np.expm1))
    rows.append({"model": mname, "CV RMSE (raw target)": base_["cv_rmse"], "CV RMSE (log1p target)": logd["cv_rmse"], "CV R2 raw": base_["cv_r2"], "CV
R2 log": logd["cv_r2"]})
ABL_LOG = pd.DataFrame(rows).set_index("model"); display(ABL_LOG.style.format(precision=4))

# (c) split criterion (tuned tree, CV)
rows = [{"criterion": c, **cv_eval(DecisionTreeRegressor(*TREE_BEST, criterion=c, random_state=SEED))}] for c in ["squared_error", "friedman_mse",
"absolute_error", "poisson"]
ABL_CRIT = pd.DataFrame(rows).set_index("criterion")[["cv_rmse", "cv_rmse_sd", "cv_r2", "cv_train_rmse"]]; display(ABL_CRIT.style.format(precision=4))
```

	stable lr	train RMSE @ end	gap to OLS optimum
raw features	8.04578e-08	0.620901	0.141017
standardised	0.15	0.479884	6.93362e-12



	CV RMSE (raw target)	CV RMSE (log1p target)	CV R ² raw	CV R ² log
model				
Linear	0.4882	0.5237	0.8827	0.8623
Poly(d=2)	0.4530	0.4488	0.8988	0.9002
Tree (tuned)	0.5134	0.5121	0.8709	0.8707

	cv_rmse	cv_rmse_sd	cv_r2	cv_train_rmse
criterion				
squared_error	0.5134	0.0478	0.8709	0.4687
friedman_mse	0.5134	0.0478	0.8709	0.4687
absolute_error	0.5019	0.0472	0.8762	0.4711
poisson	0.5060	0.0438	0.8744	0.4631

Observations — preprocessing ablations.

- **Scaling is essential for gradient descent.** On raw features the largest stable learning rate is tiny, and after the same 1,000 epochs the training RMSE is 0.621 versus 0.480 for standardised inputs (the optimum): the badly scaled directions of the loss surface converge very slowly. This empirically supports the Stage 1 decision.
- **Log-target:** worse for the linear model (CV RMSE 0.488 $\rightarrow$ 0.524; a plausible cause is that the back-transform amplifies errors on the peaks, though this was not tested separately), neutral for the polynomial (0.453 $\rightarrow$ 0.449, within noise) and no effect for the tree. The raw target is retained.
- **Split criterion:** `absolute_error` (0.502) and `poisson` (0.506) are about 0.01 better than `squared_error` / `friedman_mse` (0.513, which are identical here), i.e. within one standard error (≈ 0.021) of each other. Since the assignment's metrics are MSE-based, `squared_error` is kept.


2.7.3 Regularisation path for the linear model — how do the coefficients (and the CV error) respond to L2 shrinkage under multicollinearity?

```

RIDGE_ALPHAS = np.logspace(-2, 5, 22)
path_rows, cv_rows = [], []
for a in RIDGE_ALPHAS:
    mdl = LinearRegressionScratch(alpha=a).fit(X_train_s, y_train); path_rows.append(pd.Series(mdl.coef_, index=FEATURES, name=a))
    cv_rows.append({"alpha": a, **cv_eval(LinearRegressionScratch(alpha=a))})
RIDGE_PATH, RIDGE_CV = pd.DataFrame(path_rows), pd.DataFrame(cv_rows).set_index("alpha")
a_best = RIDGE_CV["cv_rmse"].idxmin(); print(f"Best ridge alpha (linear) = {a_best:.3g}; CV RMSE {RIDGE_CV.loc[a_best, 'cv_rmse']:.4f} vs OLS
{RIDGE_CV.iloc[0]['cv_rmse']:.4f}")
fig, ax = plt.subplots(1, 2, figsize=(14, 4.2))
for f in FEATURES: ax[0].plot(RIDGE_PATH.index, RIDGE_PATH[f], label=f)
ax[0].set_xscale("log"); ax[0].axhline(0, color="k", lw=.6); ax[0].set_xlabel("ridge  $\alpha$ "); ax[0].set_ylabel("coefficient"); ax[0].set_title("Ridge
coefficient path (linear model)"); ax[0].legend(fontsize=7, ncol=2)
ax[1].errorbar(RIDGE_CV.index, RIDGE_CV["cv_rmse"], yerr=RIDGE_CV["cv_rmse_sd"] / np.sqrt(len(folds)), color=RED, marker="o", ms=3);
ax[1].set_xscale("log")
ax[1].set_xlabel("ridge  $\alpha$ "); ax[1].set_ylabel("CV RMSE"); ax[1].set_title("Does shrinkage help the linear model?"); plt.tight_layout(); plt.show()

```

Best ridge alpha (linear) = 21.5; CV RMSE 0.4877 vs OLS 0.4882



Observation. Ridge shrinkage does **not** improve the linear model (best $\alpha \approx 21$: CV RMSE 0.4877 vs. 0.4882 for OLS) and large α quickly hurts. The multicollinearity found in Stage 1 therefore affects the *interpretation* of individual coefficients (the coefficient paths of the correlated sensors cross and change sign as α grows), but **not the predictive accuracy**.

2.8 Learning curves, per-regime errors and out-of-time robustness


- **Learning curves** (group-CV): is more data likely to help each model, or is it bias-limited?
- **Per-regime errors** on the test set: where (Low ... Very-high CO) does each model fail? (Stage 1 flagged the 8:1 imbalance.)
- **Out-of-time check:** refit with *fixed* hyper-parameters on the first 80 % of the timeline and test on the last 20 % — the chronological protocol rejected as the default in Stage 1.

```

LC = {}
sizes = np.linspace(0.1, 1.0, 8)
for mname, mk in FACT.items():
    ts, trs, vas = learning_curve(mk(), X_train_s, y_train, groups=groups, cv=cv, train_sizes=sizes, scoring="neg_root_mean_squared_error",
    shuffle=True, random_state=SEED)[0:3]
    LC[mname] = pd.DataFrame({"n_train": ts, "train_rmse": -trs.mean(1), "val_rmse": -vas.mean(1), "val_rmse_sd": vas.std(1)}).set_index("n_train")
fig, ax = plt.subplots(1, 3, figsize=(17, 4), sharey=True)
for a, (mname, df_) in zip(ax, LC.items()):
    a.plot(df_.index, df_["train_rmse"], color=BLUE, marker="o", label="train"); a.errorbar(df_.index, df_["val_rmse"], yerr=df_["val_rmse_sd"],
    color=RED, marker="o", capsize=2, label="CV validation")
    a.set_title(mname); a.set_xlabel("training samples"); a.legend()
ax[0].set_ylabel("RMSE (mg/m³)"); plt.suptitle("Learning curves (group-CV)", y=1.02, fontweight="bold"); plt.tight_layout(); plt.show()
display(pd.concat([k: v.iloc[[0, 3, -1]] for k, v in LC.items()]).style.format(precision=4))

```

Learning curves (group-CV)



	train_rmse	val_rmse	val_rmse_sd
n_train			
Linear	470	0.4517	0.5011
	2285	0.4688	0.4906
	4706	0.4785	0.4882
Poly(d=2)	470	0.3792	0.5014
	2285	0.3980	0.4585
	4706	0.4073	0.4530
Tree (tuned)	470	0.6241	0.6570
	2285	0.4699	0.5301
	4706	0.4687	0.5134

Observations.

- **Linear:** training and validation RMSE are flat and almost equal (0.478 vs. 0.488) already from ~500 samples → high bias; **more data will not help**, a more flexible model might.
- **Poly d=2:** a small, shrinking gap (0.407 vs. 0.453) and a validation curve that is nearly flat from ~3,000 samples: close to its capacity.
- **Tuned tree:** the largest data hunger — validation error falls from 0.657 (n=470) to ≈ 0.52 at n ≈ 2,900 and then plateaus.
- Fold-to-fold spread (error bars, ≈ ±0.05) is of the same order as the differences between the models, so differences below ~0.03 RMSE should not be over-interpreted.

```
FINAL = [N_MEAN, N_LIN_NE, N_POLY_NE, N_PR_SC, tree_names[3], tree_names[5], tree_names[10], tree_names[None], N_TREE_T]
bt_test = pd.cut(y_test, bands_edges, labels=bands_labels, right=False)
rows = []
for nm in FINAL:
    err = PRED[nm]["test"] - y_test.values
    for b in bands_labels:
        mk_ = (bt_test == b).values; rows.append({"model": nm, "regime": b, "n": int(mk_.sum()), "rmse": np.sqrt(np.mean(err[mk_] ** 2)), "mae":
np.mean(np.abs(err[mk_])), "bias (pred-true)": err[mk_].mean()})
REGIME = pd.DataFrame(rows)
display(REGIME.pivot(index="model", columns="regime", values="rmse"))
[bands_labels].loc[FINAL].style.format(precision=3).background_gradient(cmap="RdYlGn_r", axis=None).set_caption("Test RMSE by C0 regime")
display(REGIME.pivot(index="model", columns="regime", values="bias (pred-true)"))
[bands_labels].loc[FINAL].style.format(precision=3).background_gradient(cmap="coolwarm", axis=None, vmin=-3, vmax=3).set_caption("Mean error (pred -
true) by C0 regime: negative = under-prediction")

# residual artefacts for Stage 3
RESID = {nm: {"train": y_train.values - PRED[nm]["train"], "test": y_test.values - PRED[nm]["test"]} for nm in PRED}
```

Test RMSE by CO regime

	regime	Low (<1.5)	Moderate (1.5-3)	High (3-5)	Very high (>=5)
model					
Baseline mean predictor		1.302	0.427	1.731	3.984
Linear scratch (normal eq.)		0.384	0.433	0.519	1.173
Poly(d=2) scratch (normal eq.)		0.326	0.391	0.559	0.946
Poly(d=2)+Ridge(α=1) scratch		0.326	0.391	0.561	0.949
Tree max_depth=3		0.405	0.508	0.725	1.327
Tree max_depth=5		0.361	0.437	0.646	1.152
Tree max_depth=10		0.356	0.441	0.771	1.149
Tree max_depth=None		0.386	0.515	0.886	1.190
Tree tuned (CV, 1-SE)		0.361	0.437	0.653	1.161

Mean error (pred – true) by CO regime: negative = under-prediction

	regime	Low (<1.5)	Moderate (1.5-3)	High (3-5)	Very high (>=5)
model					
Baseline mean predictor		1.256	0.047	-1.637	-3.878
Linear scratch (normal eq.)		0.012	0.101	-0.085	-0.927
Poly(d=2) scratch (normal eq.)		0.089	-0.007	-0.114	-0.555
Poly(d=2)+Ridge(α=1) scratch		0.090	-0.007	-0.114	-0.557
Tree max_depth=3		0.166	-0.008	-0.220	-0.710
Tree max_depth=5		0.119	-0.016	-0.133	-0.660
Tree max_depth=10		0.089	-0.011	-0.190	-0.503
Tree max_depth=None		0.076	-0.009	-0.197	-0.414
Tree tuned (CV, 1-SE)		0.119	-0.016	-0.126	-0.556

Observations — where do the models fail?

- **All models under-predict the rare *Very high* hours** (mean error –0.93 linear, –0.56 polynomial, –0.56 tuned tree, –0.41 unrestricted tree) and their RMSE is 2–3× larger there than in the other regimes (0.95–1.19 vs. 0.33–0.65). This is the consequence of the 8:1 imbalance flagged in Stage 1: the squared-error objective is dominated by the bulk of the data.
- The polynomial model is best in the Low (0.326), Moderate (0.391) and Very-high (0.946) regimes; the **linear model is best in the High regime** (0.519 vs. 0.559 polynomial, 0.653 tuned tree).
- Overall RMSE therefore hides substantial regime-dependent differences; Stage 3 will use these tables.

```
cut = int(len(data) * 0.8); Xo_tr, Xo_te, yo_tr, yo_te = data[FEATURES].iloc[:cut], data[FEATURES].iloc[cut:], data[TARGET].iloc[:cut], data[TARGET].iloc[cut:]
OOT_MODELS = {"Mean baseline": lambda: DummyRegressor(), "Linear": make_linear, f"Poly(d={DEG_STAR})": make_poly,
              f"Poly(d={DEG_R})+Ridge": lambda: PolyRegressionScratch(DEG_R, alpha=ALPHA_R), "Tree max_depth=5": lambda:
DecisionTreeRegressor(max_depth=5, random_state=SEED),
              "Tree (tuned)": make_tree}

rows = []
for mname, mk in OOT_MODELS.items():
    pipe = Pipeline([("scale", StandardScaler()), ("m", mk)]).fit(Xo_tr, yo_tr)
    rows.append({"model": mname, "oot_train_rmse": rmse(yo_tr, pipe.predict(Xo_tr)), "oot_test_rmse": rmse(yo_te, pipe.predict(Xo_te)), "oot_test_r2":
r2(yo_te, pipe.predict(Xo_te)), "oot_test_mae": mae(yo_te, pipe.predict(Xo_te))})
OOT = pd.DataFrame(rows).set_index("model")
OOT["random-split test RMSE"] = [RESULTS[n]["test_rmse"] for n in [N_MEAN, N_LIN_NE, N_POLY_NE, N_PR_SC, tree_names[5], N_TREE_T]]
OOT["degradation (oot - random)"] = OOT["oot_test_rmse"] - OOT["random-split test RMSE"]
print(f"Out-of-time split: train {Xo_tr.index.min():%d %b %Y} - {Xo_tr.index.max():%d %b %Y}, test {Xo_te.index.min():%d %b %Y} - {Xo_te.index.max():%d %b %Y}")
display(OOT.style.format(precision=4))
```

Out-of-time split: train 10 Mar 2004 – 28 Jan 2005, test 28 Jan 2005 – 04 Apr 2005

	oot_train_rmse	oot_test_rmse	oot_test_r2	oot_test_mae	random-split test RMSE	degradation (oot - random)
model						
Mean baseline	1.4650	1.3197	-0.0299	1.0677	1.4232	-0.1035
Linear	0.4845	0.4836	0.8617	0.3437	0.4933	-0.0097
Poly(d=2)	0.4140	0.4460	0.8824	0.3020	0.4470	-0.0011
Poly(d=2)+Ridge	0.4144	0.4427	0.8841	0.2980	0.4474	-0.0047
Tree max_depth=5	0.4493	0.5276	0.8354	0.3600	0.5131	0.0145
Tree (tuned)	0.4651	0.5126	0.8446	0.3544	0.5157	-0.0031

Observations — out-of-time robustness (hyper-parameters fixed, *not* re-tuned).

- Contrary to the concern raised in Stage 1, the chronological split is **not materially harder** in RMSE: linear 0.484 (vs. 0.493 random), polynomial 0.446 (0.447), tuned tree 0.513 (0.516). Despite the clear covariate shift in T , AH and $PT08.S4$, the $CO \leftrightarrow$ sensor relationship (dominated by $PT08.S2$) is stable across seasons.
- The ranking is unchanged: **polynomial < linear < trees**; the plain depth-5 tree degrades slightly ($0.513 \rightarrow 0.528$) while the tuned tree does not.
- R^2 is lower out-of-time for the linear model (0.862 vs. 0.880) because the later test period has lower CO variance (the mean predictor scores RMSE 1.32 there).
- This is one split, so it supports the conclusion but is not a precise estimate of forecasting error.

2.9 Hand-off to Stages 3–6 and sanity checks

```
art = pd.DataFrame({"artifact": ["RESULTS / MASTER", "FITTED", "PRED", "RESID", "HISTORY", "TREE_SWEEP", "POLY_CV (agg)", "IMP_TREE", "PERM",
                                "COEF_BOOT",
                                "ABL_FEAT", "ABL_SCALE / ABL_LOG / ABL_CRIT", "RIDGE_PATH / RIDGE_CV", "LC", "REGIME", "OOT"],
                    "content": ["train/test MSE, RMSE, R², MAE (+CV) for 15 models", "fitted estimators (trees can be plotted)", "train/test
                                predictions per model",
                                "train/test residuals per model", "per-epoch train/test metrics for the 2 gradient-based models",
                                "train/CV/test metrics for depth 1-20 and None", "degree × α CV grid", "impurity importances per tree",
                                "permutation importances on test", "linear coefficients with bootstrap CIs", "feature-group and drop-one ablation",
                                "scaling, log-target and criterion ablations", "ridge coefficient path and CV", "learning curves", "per-regime
                                RMSE/MAE/bias", "out-of-time results"]}).set_index("artifact")
display(art)
assert set(HISTORY) == {N_LIN_GD, N_POLY_GD} and all(np.isfinite(HISTORY[k].values).all() for k in HISTORY)
assert all(np.isfinite(MASTER[["train_rmse", "test_rmse"]].values).ravel())
print("Registry holds", len(RESULTS), "models; histories:", {k: len(v) for k, v in HISTORY.items()}, "epochs")
```

artifact	content
RESULTS / MASTER	train/test MSE, RMSE, R², MAE (+CV) for 15 models
FITTED	fitted estimators (trees can be plotted)
PRED	train/test predictions per model
RESID	train/test residuals per model
HISTORY	per-epoch train/test metrics for the 2 gradien...
TREE_SWEEP	train/CV/test metrics for depth 1-20 and None
POLY_CV (agg)	degree × α CV grid
IMP_TREE	impurity importances per tree
PERM	permutation importances on test
COEF_BOOT	linear coefficients with bootstrap CIs
ABL_FEAT	feature-group and drop-one ablation
ABL_SCALE / ABL_LOG / ABL_CRIT	scaling, log-target and criterion ablations
RIDGE_PATH / RIDGE_CV	ridge coefficient path and CV
LC	learning curves
REGIME	per-regime RMSE/MAE/bias
OOT	out-of-time results

Registry holds 15 models; histories: {'Linear | scratch (GD)': 1001, 'Poly(d=2) | scratch (Adam GD)': 3001} epochs

Stage 2 take-aways (feeding the later stages)

- All from-scratch models match scikit-learn; selection used leakage-safe group-CV only.
- Polynomial degree 2 is the CV-optimal and 1-SE degree; linear is bias-limited; trees are best at depth $\approx 5\text{--}6$ with $\approx 20\text{--}30$ leaves.
- $PT08.S2$ (NMHC) is the dominant feature in every model and ablation; weather alone is uninformative but adds a small correction.
- All models under-predict the rare high- CO episodes.
- Results are stable under an out-of-time split.

 **Stage 2 complete** — models trained, verified against scikit-learn, tuned with leakage-safe CV, and every artifact needed for the evaluation, visualisation, comparison and discussion stages is stored.

Stage 3 — Model Evaluation

This stage reports the three required metrics — **MSE, RMSE and R²** — for *every* model, and then goes further than point estimates because Stage 1–2 showed that (i) hourly samples are autocorrelated and (ii) many model differences are small compared with the fold-to-fold noise ($\approx \pm 0.05$ RMSE).

Section	Content
3.1	Required metrics table (train and test) for all 15 models
3.2	Uncertainty: week-clustered bootstrap CIs on test metrics and paired model differences; fold-wise CV comparison
3.3	Performance metrics vs. epoch (gradient-trained regression models)
3.4	Performance metrics vs. tree depth
3.5	Residual diagnostics (fit quality, bias, heteroscedasticity, normality)
3.6	Where do errors occur? CO regime, hour of day, month
3.7	Evaluation summary: accuracy vs. complexity vs. cost

3.1 Required metrics for all models (MSE, RMSE, R²)

```
REQ_COLS = ["train_mse", "train_rmse", "train_r2", "test_mse", "test_rmse", "test_r2"]
REQUIRED = MASTER[REQ_COLS].astype(float).copy()
assert REQUIRED.notna().all().all() and len(REQUIRED) == 15
REQUIRED.columns = pd.MultiIndex.from_tuples([("Train", "MSE"), ("Train", "RMSE"), ("Train", "R²"), ("Test", "MSE"), ("Test", "RMSE"), ("Test", "R²")])
display(REQUIRED.style.format(precision=4).background_gradient(subset=[("Test", "RMSE")], cmap="RdYlGn_r")
        .background_gradient(subset=[("Test", "R²")], cmap="RdYlGn").set_caption("MSE and RMSE in (mg/m³)² and mg/m³; R² unitless"))
print("Extra reported measures (not required): MAE, CV scores and the train-test RMSE gap:")
display(MASTER[["train_mae", "test_mae", "cv_rmse", "cv_r2", "rmse_gap (test-train)"]].astype(float).style.format(precision=4, na_rep="-"))
```

MSE and RMSE in (mg/m³)² and mg/m³; R² unitless

	Train			Test		
	MSE	RMSE	R²	MSE	RMSE	R²
Baseline mean predictor	2.0726	1.4397	0.0000	2.0255	1.4232	-0.0001
Linear scratch (normal eq.)	0.2303	0.4799	0.8889	0.2433	0.4933	0.8799
Linear scratch (GD)	0.2303	0.4799	0.8889	0.2433	0.4933	0.8799
Linear sklearn	0.2303	0.4799	0.8889	0.2433	0.4933	0.8799
Poly(d=2) scratch (normal eq.)	0.1691	0.4112	0.9184	0.1998	0.4470	0.9013
Poly(d=2) scratch (Adam GD)	0.1691	0.4112	0.9184	0.1998	0.4470	0.9013
Poly(d=2) sklearn	0.1691	0.4112	0.9184	0.1998	0.4470	0.9013
Poly(d=2)+Ridge(α=1) scratch	0.1692	0.4113	0.9184	0.2001	0.4474	0.9012
Poly(d=2)+Ridge(α=1) sklearn	0.1692	0.4113	0.9184	0.2001	0.4474	0.9012
Tree max_depth=3	0.3042	0.5515	0.8532	0.3427	0.5854	0.8308
Tree max_depth=5	0.1938	0.4403	0.9065	0.2633	0.5131	0.8700
Tree max_depth=10	0.0678	0.2603	0.9673	0.2961	0.5442	0.8538
Tree max_depth=None	0.0000	0.0000	1.0000	0.3716	0.6096	0.8165
Tree tuned (CV, 1-SE)	0.2207	0.4698	0.8935	0.2659	0.5157	0.8887
Tree pruned (ccp, CV)	0.2139	0.4625	0.8968	0.2813	0.5304	0.8611

Extra reported measures (not required): MAE, CV scores and the train-test RMSE gap:

	train_mae	test_mae	cv_rmse	cv_r2	rmse_gap (test-train)
Baseline mean predictor	1.1058	1.1125	1.4378	-0.0100	-0.0165
Linear scratch (normal eq.)	0.3199	0.3334	0.4882	0.8827	0.0134
Linear scratch (GD)	0.3199	0.3334	–	–	0.0134
Linear sklearn	0.3199	0.3334	–	–	0.0134
Poly(d=2) scratch (normal eq.)	0.2669	0.2860	0.4530	0.8988	0.0359
Poly(d=2) scratch (Adam GD)	0.2669	0.2860	–	–	0.0359
Poly(d=2) sklearn	0.2669	0.2860	–	–	0.0359
Poly(d=2)+Ridge(α=1) scratch	0.2666	0.2861	0.4516	0.8994	0.0360
Poly(d=2)+Ridge(α=1) sklearn	0.2666	0.2861	–	–	0.0360
Tree max_depth=3	0.3949	0.4228	0.5773	0.8347	0.0339
Tree max_depth=5	0.3023	0.3464	0.5092	0.8730	0.0728
Tree max_depth=10	0.1764	0.3380	0.5981	0.8236	0.2838
Tree max_depth=None	0.0000	0.3788	0.6593	0.7859	0.6096
Tree tuned (CV, 1-SE)	0.3105	0.3480	0.5134	0.8709	0.0459
Tree pruned (ccp, CV)	0.3228	0.3660	0.5212	0.8670	0.0679

Observations.

- The required metrics are reported for all 15 models. The scikit-learn and from-scratch versions of the same model are identical to 4 decimals, so the table really contains **9 distinct models** (mean baseline, linear, polynomial d=2, polynomial d=2 + ridge, and 5 trees).
- Best test scores: **Poly d=2 (RMSE 0.447, R² 0.901)**, then Linear (0.493, 0.880), then Tree depth 5 (0.513, 0.870). The mean predictor has R² ≈ 0 and RMSE 1.423.
- Overfitting is visible only in the trees: the train–test RMSE gap grows from 0.034 (depth 3) to 0.073 (depth 5), 0.284 (depth 10) and 0.610 (depth None), while for linear it is 0.013.
- MAE shows the same ordering as RMSE (polynomial 0.286 < linear 0.333 ≈ depth-10 tree 0.338 < depth-5 tree 0.346); the unrestricted tree has the largest MAE among the trees (0.379).

3.2 How certain are these numbers?

A single test RMSE hides sampling noise. We therefore resample the **test set by calendar week** (cluster bootstrap, 1,000 replicates): hours inside a week are strongly correlated, so resampling individual hours would give intervals that are too narrow. All models are evaluated on the *same* resamples, so **paired differences** between models are meaningful.

```
test_weeks = pd.factorize(X_test.index.to_period("W").astype(str))[0]
members = [np.where(test_weeks == w)[0] for w in np.unique(test_weeks)]
rb = np.random.default_rng(SEED)
BOOT_IDX = [np.concatenate([members[i] for i in rb.integers(0, len(members), len(members))]) for _ in range(1000)]
yt_all = y_test.values


def boot_dist(name):
    e = yt_all - PRED[name]["test"]; rm, rr = [], []
    for idx in BOOT_IDX:
        s_res = np.sum(e[idx] ** 2); rm.append(np.sqrt(s_res / len(idx))); rr.append(1 - s_res / np.sum((yt_all[idx] - yt_all[idx].mean()) ** 2))
    return np.array(rm), np.array(rr)

BOOT = {nm: boot_dist(nm) for nm in ORDER}
ci = lambda a: np.percentile(a, [2.5, 97.5])
BOOT_TBL = pd.DataFrame({nm: {"test RMSE": RESULTS[nm]["test_rmse"], "RMSE 95% CI low": ci(BOOT[nm][0])[0], "RMSE 95% CI high": ci(BOOT[nm][0])[1],
                             "test R²": RESULTS[nm]["test_r2"], "R² 95% CI low": ci(BOOT[nm][1])[0], "R² 95% CI high": ci(BOOT[nm][1])[1]} for nm in
ORDER}).T
print(f"{len(members)} test weeks resampled with replacement (cluster bootstrap)")
display(BOOT_TBL.style.format(precision=4))

# forest plot (mean baseline omitted: it is ~3x worse and would compress the scale)
fam_col = {"Linear": BLUE, "Polynomial": ORANGE, "Polynomial+Ridge": ORANGE, "Decision Tree": GREEN}
sub = BOOT_TBL.drop(index=N_MEAN).sort_values("test RMSE", ascending=False)
fig, ax = plt.subplots(1, 2, figsize=(15, 5.5), sharey=True)
for a, (m_, lo, hi, ttl) in zip(ax, [("test RMSE", "RMSE 95% CI low", "RMSE 95% CI high", "Test RMSE (mg/m³) with 95% week-bootstrap CI"),
                                   ("test R²", "R² 95% CI low", "R² 95% CI high", "Test R² with 95% week-bootstrap CI")]):
    for i, (nm, r) in enumerate(sub.iterrows()):
        a.errorbar(r[m_], i, xerr=[r[lo], r[hi] - r[m_]], fmt="o", color=fam_col[RESULTS[nm]["family"]], capsize=3)
    a.set_yticks(range(len(sub))); a.set_yticklabels(sub.index, fontsize=8); a.set_title(ttl); a.grid(axis="y", alpha=.2)
plt.tight_layout(); plt.show()
```

57 test weeks resampled with replacement (cluster bootstrap)

	test RMSE	RMSE 95% CI low	RMSE 95% CI high	test R ²	R ² 95% CI low	R ² 95% CI high
Baseline mean predictor	1.4232	1.3116	1.5342	-0.0001	-0.0163	-0.0000
Linear scratch (normal eq.)	0.4933	0.4154	0.5774	0.8799	0.8479	0.9084
Linear scratch (GD)	0.4933	0.4154	0.5774	0.8799	0.8479	0.9084
Linear sklearn	0.4933	0.4154	0.5774	0.8799	0.8479	0.9084
Poly(d=2) scratch (normal eq.)	0.4470	0.3679	0.5320	0.9013	0.8707	0.9292
Poly(d=2) scratch (Adam GD)	0.4470	0.3679	0.5320	0.9013	0.8707	0.9292
Poly(d=2) sklearn	0.4470	0.3679	0.5320	0.9013	0.8707	0.9292
Poly(d=2)+Ridge($\alpha=1$) scratch	0.4474	0.3685	0.5326	0.9012	0.8705	0.9293
Poly(d=2)+Ridge($\alpha=1$) sklearn	0.4474	0.3685	0.5326	0.9012	0.8705	0.9293
Tree max_depth=3	0.5854	0.5103	0.6690	0.8308	0.7944	0.8630
Tree max_depth=5	0.5131	0.4398	0.5906	0.8700	0.8405	0.8971
Tree max_depth=10	0.5442	0.4725	0.6165	0.8538	0.8251	0.8810
Tree max_depth=None	0.6096	0.5230	0.6986	0.8165	0.7753	0.8542
Tree tuned (CV, 1-SE)	0.5157	0.4439	0.5919	0.8687	0.8400	0.8961
Tree pruned (ccp, CV)	0.5304	0.4611	0.6052	0.8611	0.8320	0.8873



Observations.

- The 95 % week-bootstrap intervals are wide ($\approx \pm 0.08$ RMSE) and overlap heavily across models, so **individual intervals cannot separate the models** — the paired differences below are the proper comparison.
- Even so, the intervals confirm that every model is clearly better than the mean baseline (RMSE 1.31–1.53) and that the unrestricted tree (RMSE up to 0.70) and depth 3 (up to 0.67) are the weakest.

```
# paired differences on identical bootstrap resamples: delta = RMSE(A) - RMSE(B) (negative -> A better)
cmp_pairs = [("Poly d=2 vs Linear", N_POLY_NE, N_LIN_NE), ("Linear vs Tree depth 5", N_LIN_NE, tree_names[5]), ("Linear vs Tree tuned", N_LIN_NE,
N_TREE_T),

              ("Poly d=2 vs Tree tuned", N_POLY_NE, N_TREE_T), ("Tree depth 5 vs depth 3", tree_names[5], tree_names[3]),
              ("Tree depth 5 vs depth 10", tree_names[5], tree_names[10]), ("Tree depth 5 vs None", tree_names[5], tree_names[None]),
              ("Linear vs Poly(d=2)+Ridge", N_LIN_NE, N_PR_SC)]

rows = []
for lbl, a, b in cmp_pairs:
    d = BOOT[a][0] - BOOT[b][0]; lo, hi = ci(d)
    rows.append({"comparison (A vs B)": lbl, "ARMSE point": RESULTS[a]["test_rmse"] - RESULTS[b]["test_rmse"], "ARMSE 95% CI low": lo, "ARMSE 95% CI
high": hi,

                "P(A better)": np.mean(d < 0), "CI excludes 0": bool(lo > 0 or hi < 0)})

PAIRED = pd.DataFrame(rows).set_index("comparison (A vs B)"); display(PAIRED.style.format(precision=4))

# fold-wise CV RMSE (identical group folds) -> paired comparison independent of the test set
FOLD_FACT = {"Mean baseline": lambda: DummyRegressor(), "Linear": make_linear, f"Poly(d={DEG_STAR})": make_poly,
             "Tree depth 3": lambda: DecisionTreeRegressor(max_depth=3, random_state=SEED), "Tree depth 5": lambda: DecisionTreeRegressor(max_depth=5,
random_state=SEED),
             "Tree depth 10": lambda: DecisionTreeRegressor(max_depth=10, random_state=SEED), "Tree depth None": lambda:
DecisionTreeRegressor(random_state=SEED), "Tree (tuned)": make_tree}
FOLD_RMSE = pd.DataFrame({nm: [rmse(y_train.values[va], mk().fit(X_train_s.values[tr], y_train.values[tr]).predict(X_train_s.values[va])) for tr, va in
folds]

                        for nm, mk in FOLD_FACT.items(), index=[f"fold {i + 1}" for i in range(len(folds))])
FOLD_RMSE.loc["mean"] = FOLD_RMSE.mean(); FOLD_RMSE.loc["sd"] = FOLD_RMSE.iloc[:5].std()
display(FOLD_RMSE.style.format(precision=4))
wins = {c: int((FOLD_RMSE.iloc[:5][c] < FOLD_RMSE.iloc[:5][c]).sum()) for c in FOLD_RMSE.columns if c.startswith("Tree")}
print("Folds (out of 5) in which Linear beats each tree:", wins)
print("Folds in which Poly beats Linear:", int((FOLD_RMSE.iloc[:5][f"Poly(d={DEG_STAR})"] < FOLD_RMSE.iloc[:5]['Linear']).sum()))
```

	ARMSE point	ARMSE 95% CI low	ARMSE 95% CI high	P(A better)	CI excludes 0
comparison (A vs B)					
Poly d=2 vs Linear	-0.0463	-0.0622	-0.0323	1.0000	True
Linear vs Tree depth 5	-0.0198	-0.0393	0.0006	0.9730	False
Linear vs Tree tuned	-0.0224	-0.0443	-0.0001	0.9750	True
Poly d=2 vs Tree tuned	-0.0686	-0.0914	-0.0437	1.0000	True
Tree depth 5 vs depth 3	-0.0723	-0.1011	-0.0427	1.0000	True
Tree depth 5 vs depth 10	-0.0311	-0.0675	0.0018	0.9650	False
Tree depth 5 vs None	-0.0965	-0.1389	-0.0559	1.0000	True
Linear vs Poly(d=2)+Ridge	0.0459	0.0321	0.0620	0.0000	True

	Mean baseline	Linear	Poly(d=2)	Tree depth 3	Tree depth 5	Tree depth 10	Tree depth None	Tree (tuned)
fold 1	1.3566	0.4331	0.3876	0.5289	0.4426	0.5273	0.5646	0.4576
fold 2	1.3765	0.5244	0.5017	0.6302	0.5272	0.6104	0.6715	0.5349
fold 3	1.3205	0.4161	0.4020	0.5448	0.4548	0.5866	0.6432	0.4544
fold 4	1.5136	0.5581	0.4913	0.6106	0.5665	0.6500	0.7300	0.5568
fold 5	1.6215	0.5090	0.4826	0.5719	0.5548	0.6160	0.6872	0.5631
mean	1.4378	0.4882	0.4530	0.5773	0.5092	0.5981	0.6593	0.5134
sd	0.1261	0.0610	0.0538	0.0428	0.0572	0.0456	0.0616	0.0534

Folds (out of 5) in which Linear beats each tree: {'Tree depth 3': 5, 'Tree depth 5': 5, 'Tree depth 10': 5, 'Tree depth None': 5, 'Tree (tuned)': 4}
Folds in which Poly beats Linear: 5

Observations — paired comparisons (negative $\Delta \Rightarrow$ A better).

- **Polynomial d=2 beats Linear:** Δ RMSE = -0.046 , 95 % CI $[-0.062, -0.032]$, better in 100 % of resamples, and in **5 / 5 CV folds**.
- **Linear beats the Decision Tree, but only modestly:** vs. depth 5 Δ RMSE = -0.020 (CI $[-0.039, +0.001]$, $P = 0.97$ — borderline, CI just touches 0); vs. the tuned tree -0.022 (CI $[-0.044, -0.0001]$). In CV, linear is better in **5 / 5 folds against every fixed-depth tree and 4 / 5 against the tuned tree**.
- **Polynomial d=2 beats the tuned tree clearly** ($\Delta = -0.069$, CI excludes 0).
- Within the trees, depth 5 is clearly better than depth 3 and than None , but **not significantly better than depth 10** (CI $[-0.068, +0.002]$).
- The ridge-regularised polynomial is statistically indistinguishable from the unregularised one (Δ test RMSE 0.0004).

3.3 Performance metrics vs. epoch (regression models)


Per-epoch train/test **MSE**, **RMSE** and **R²** of the two gradient-trained models. Dotted lines mark the closed-form optimum. The number of epochs was fixed beforehand from the *training* loss; the test curve is displayed purely as a diagnostic (it would reveal over-training, i.e. early stopping, if there were any).

```
fig, axes = plt.subplots(2, 3, figsize=(18, 8))
for r, (nm, ttl, opt) in enumerate([(N_LIN_GD, "Linear regression (batch GD)", N_LIN_NE), (N_POLY_GD, f"Polynomial d={DEG_STAR} (Adam)", N_POLY_NE)]):
    h = HISTORY[nm]
    for c, (met, lab) in enumerate([("mse", "MSE"), ("rmse", "RMSE (mg/m³)"), ("r2", "R²")]):
        a = axes[r, c]
        a.plot(h[f"train_{met}"], h[f"train_{met}"], color=BLUE, label="train"); a.plot(h[f"epoch"], h[f"test_{met}"], color=ORANGE, label="test")
        a.axhline(RESULTS[opt][f"train_{met}"], color=BLUE, ls=":", lw=1); a.axhline(RESULTS[opt][f"test_{met}"], color=ORANGE, ls=":", lw=1)
        a.set_xscale("symlog", linthresh=1); a.set_xlim(0, h[f"epoch"].max()); a.set_xlabel("epoch"); a.set_ylabel(lab); a.set_title(f"{ttl}: {lab}")
        if met == "r2": a.set_ylim(-3, 1)
        else: a.set_yscale("log")
        if r == 0 and c == 0: a.legend()
plt.suptitle("Metrics vs. epoch (dotted = closed-form optimum)", y=1.01, fontweight="bold"); plt.tight_layout(); plt.show()

chk = [0, 1, 5, 10, 50, 100, 200, 500, 1000, 3000]
rows = []
for nm in (N_LIN_GD, N_POLY_GD):
    h = HISTORY[nm].set_index("epoch")
    for e in [c for c in chk if c in h.index]:
        rows.append({"model": nm, "epoch": e, **{f"{s}_{m}": h.loc[e, f"{s}_{m}"]} for s in ("train", "test") for m in ("mse", "rmse", "r2")})
display(pd.DataFrame(rows).set_index(["model", "epoch"]).style.format(precision=4))
conv = pd.DataFrame({nm: {"epochs to 1% of final train MSE": epochs_to_tol(HISTORY[nm], tol=0.01), "epochs to 0.1%": epochs_to_tol(HISTORY[nm],
tol=0.001),
                        "epoch of minimum TEST RMSE": int(HISTORY[nm].loc[HISTORY[nm][f"test_rmse"].idxmin(), "epoch"]), "final epoch":
int(HISTORY[nm][f"epoch"].max()),
                        "final test RMSE": HISTORY[nm][f"test_rmse"].iloc[-1], "minimum test RMSE": HISTORY[nm][f"test_rmse"].min()} for nm in
(N_LIN_GD, N_POLY_GD)}).T
display(conv.style.format(precision=4))


# effect of the learning rate on the per-epoch curves (training metrics, linear model)
fig, ax = plt.subplots(1, 2, figsize=(14, 4))
for lr, c in zip([0.001, 0.01, 0.05, 0.1, 0.15], [GREY, BLUE, GREEN, ORANGE, RED]):
    h = lr_hist[lr]; ax[0].plot(h[f"epoch"], h[f"train_mse"], color=c, label=f"lr={lr}"); ax[1].plot(h[f"epoch"], h[f"train_r2"], color=c, label=f"lr=
{lr}")
for a, lab in zip(ax, ["train MSE", "train R²"]): a.set_xscale("symlog", linthresh=1); a.set_xlim(0, EPOCHS_LIN); a.set_xlabel("epoch");
a.set_ylabel(lab)
ax[0].set_yscale("log"); ax[1].set_ylim(-3, 1); ax[0].legend(); ax[0].set_title("Learning rate vs. convergence speed (MSE)"); ax[1].set_title("... and
R²")
plt.tight_layout(); plt.show()
```

Metrics vs. epoch (dotted = closed-form optimum)



		train_mse	train_rmse	train_r2	test_mse	test_rmse	test_r2
model epoch							
Linear scratch (GD)	0	6.6191	2.5728	-2.1936	6.5178	2.5530	-2.2181
	1	2.7233	1.6502	-0.3139	2.7770	1.6664	-0.3711
	5	0.4351	0.6596	0.7901	0.4497	0.6706	0.7779
	10	0.2891	0.5377	0.8605	0.3043	0.5516	0.8497
	50	0.2477	0.4977	0.8805	0.2656	0.5154	0.8689
	100	0.2355	0.4853	0.8864	0.2517	0.5017	0.8757
	200	0.2308	0.4804	0.8887	0.2448	0.4948	0.8791
	500	0.2303	0.4799	0.8889	0.2434	0.4933	0.8798
	1000	0.2303	0.4799	0.8889	0.2433	0.4933	0.8799
	0	6.6191	2.5728	-2.1936	6.5178	2.5530	-2.2181
	1	5.7345	2.3947	-1.7668	5.7554	2.3991	-1.8417
Poly(d=2) scratch (Adam GD)	5	4.6004	2.1449	-1.2196	4.6010	2.1450	-1.2717
	10	3.8278	1.9565	-0.8469	3.9022	1.9754	-0.9267
	50	0.8548	0.9246	0.5876	0.8875	0.9421	0.5618
	100	0.2196	0.4686	0.8940	0.2453	0.4953	0.8789
	200	0.1847	0.4297	0.9109	0.2103	0.4585	0.8962
	500	0.1725	0.4153	0.9168	0.2018	0.4492	0.9004
	1000	0.1693	0.4115	0.9183	0.2003	0.4475	0.9011
	3000	0.1691	0.4112	0.9184	0.1998	0.4470	0.9013

	epochs to 1% of final train MSE	epochs to 0.1%	epoch of minimum TEST RMSE	final epoch	final test RMSE	minimum test RMSE
Linear scratch (GD)	135.0000	233.0000	1000.0000	1000.0000	0.4933	0.4933
Poly(d=2) scratch (Adam GD)	638.0000	1065.0000	2794.0000	3000.0000	0.4470	0.4469



Observations — metrics vs. epoch.

- Linear (batch GD):** MSE falls from 6.62 to within 1 % of its final value in **135 epochs** (0.1 % in 233) and then follows the closed-form optimum exactly (train RMSE 0.4799, test 0.4933). Train and test curves move together at every epoch, so there is **no over-training** — the test minimum is at the last epoch and early stopping would not help.
- Polynomial (Adam):** much slower start (R^2 still -0.85 after 10 epochs; it needs ≈ 100 epochs to reach $R^2 \approx 0.88$) and ≈ 640 epochs to settle within 1 %; it ends at the closed-form optimum (test RMSE 0.4470). The train–test gap is slightly larger than

for the linear model, in line with its extra capacity. The small transient spikes after epoch ~1,000 are typical of Adam's adaptive steps; the loss recovers immediately.


- The learning-rate plot shows the expected trade-off: $\text{lr} = 0.001$ has barely begun to converge after 1,000 epochs while $\text{lr} = 0.15$ converges in about 100.

3.4 Performance metrics vs. tree depth

Train, CV and test **MSE**, **RMSE** and **R²** for depths 1–20 and None (plotted at its effective depth, 28). The four depths required by the assignment (3, 5, 10, None) are marked.

```
TS = TREE_SWEEP.copy(); x = TS["depth_actual"]
TS["cv_mse"] = TS["cv_rmse"] ** 2 # (mean-of-RMSE)^2: a close proxy for CV MSE
fig, axes = plt.subplots(1, 3, figsize=(19, 4.8))
for a, (met, lab) in zip(axes, [("mse", "MSE"), ("rmse", "RMSE (mg/m³)"), ("r2", "R²")]):
    a.plot(x, TS[f"train_{met}"], color=BLUE, marker="o", ms=4, label="train")
    if met != "mse": a.plot(x, TS[f"cv_{met}"], color=RED, marker="o", ms=4, label="CV (validation weeks)")
    a.plot(x, TS[f"test_{met}"], color=ORANGE, marker="o", ms=4, label="test")
    for d in (3, 5, 10, 28): a.axvline(d, color=GREY, ls=":", lw=1)
    a.set_xticks([1, 3, 5, 10, 15, 20, 28]); a.set_xticklabels(["1", "3", "5", "10", "15", "20", "None\n(28)"]); a.set_xlabel("max_depth");
a.set_ylabel(lab); a.set_title(f"{lab} vs. tree depth")
axes[1].legend(); plt.tight_layout(); plt.show()

req_tbl = TS[TS["max_depth"].astype(str).isin(["3", "5", "10", "None"])].set_index("max_depth")[["depth_actual", "n_leaves", "train_mse", "train_rmse",
"train_r2", "test_mse", "test_rmse", "test_r2", "cv_rmse", "cv_r2"]]
display(req_tbl.style.format(precision=4).set_caption("Required depths"))
opt_cv = TS.loc[TS["cv_rmse"].idxmin()]; opt_te = TS.loc[TS["test_rmse"].idxmin()]
print(f"CV-optimal depth: {int(opt_cv['depth_actual'])} (CV RMSE {opt_cv['cv_rmse']:.4f}, test RMSE {opt_cv['test_rmse']:.4f}) | "
      f"test-optimal depth (hindsight only): {int(opt_te['depth_actual'])} (test RMSE {opt_te['test_rmse']:.4f})")
```



Required depths

	depth_actual	n_leaves	train_mse	train_rmse	train_r2	test_mse	test_rmse	test_r2	cv_rmse	cv_r2
max_depth										
3	3	8	0.3042	0.5515	0.8532	0.3427	0.5854	0.8308	0.5773	0.8347
5	5	32	0.1938	0.4403	0.9065	0.2633	0.5131	0.8700	0.5092	0.8730
10	10	693	0.0678	0.2603	0.9673	0.2961	0.5442	0.8538	0.5981	0.8236
None	28	4245	0.0000	0.0000	1.0000	0.3716	0.6096	0.8165	0.6593	0.7859

CV-optimal depth: 6 (CV RMSE 0.4998, test RMSE 0.5049) | test-optimal depth (hindsight only): 7 (test RMSE 0.5014)

Observations — metrics vs. tree depth.

- Training error decreases monotonically to exactly 0 ($R^2 = 1$) at None ; test/CV error is **U-shaped**: it drops quickly until depth ≈ 5 –6 and then rises as the tree starts to memorise noise.
- Required depths: **3** $\rightarrow$ **under-fit** (test RMSE 0.585), **5** $\rightarrow$ **good** (0.513, R^2 0.870), **10** $\rightarrow$ **over-fit** (0.544, train RMSE 0.260), **None** $\rightarrow$ **worst** (0.610, 4,245 leaves).
- The CV-optimal depth is 6 (CV RMSE 0.500, test 0.505); the test-optimal depth in hindsight is 7 (0.501), but the difference between depths 5–8 on test is negligible, so the choice of ≈ 5 –6 is robust. The CV curve is a little sharper than the test curve because validation folds are whole weeks and therefore harder than interleaved test hours.


3.5 Residual diagnostics

For the three main model families (Linear, Polynomial $d=2$, Tree depth 5) on the **test set**: predicted vs. actual, residuals vs. fitted,

residual distribution and Q-Q plot, and residuals over time. Residual = actual – predicted (positive ⇒ under-prediction).

```
DIAG = [(N_LIN_NE, "Linear"), (N_POLY_NE, f"Poly d={DEG_STAR}"), (tree_names[5], "Tree depth 5")]
reg_col = dict(zip(bands_labels, [BLUE, GREEN, ORANGE, RED]))
bt = pd.cut(y_test, bands_edges, labels=bands_labels, right=False)
fig, axes = plt.subplots(4, 3, figsize=(17, 17))
for j, (nm, ttl) in enumerate(DIAG):
    p, res = PRED[nm]["test"], yt_all - PRED[nm]["test"]
    for b in bands_labels:
        m_ = (bt == b).values; axes[0, j].scatter(yt_all[m_], p[m_], s=7, alpha=.5, color=reg_col[b], label=b)
    axes[0, j].plot([0, 9], [0, 9], "k--", lw=1); axes[0, j].set_xlabel("actual CO(GT)"); axes[0, j].set_ylabel("predicted"); axes[0, j].set_title(f"{ttl}: predicted vs actual")
    axes[1, j].scatter(p, res, s=6, alpha=.4, color=BLUE); axes[1, j].axhline(0, color="k", lw=1)
    q = pd.qcut(p, 12, duplicates="drop"); g = pd.DataFrame({"p": p, "r": res}).groupby(q, observed=True).agg(p=("p", "mean"), m=("r", "mean"), s=("r", "std"))
    axes[1, j].plot(g["p"], g["m"], color=RED, lw=2, label="binned mean"); axes[1, j].fill_between(g["p"], g["m"] - g["s"], g["m"] + g["s"], color=RED, alpha=.15, label="± 1 sd")
    axes[1, j].set_xlabel("fitted value"); axes[1, j].set_ylabel("residual"); axes[1, j].set_title(f"{ttl}: residuals vs fitted")
    sns.histplot(res, bins=40, stat="density", color=BLUE, ax=axes[2, j], edgecolor="white")
    xs = np.linspace(res.min(), res.max(), 200); axes[2, j].plot(xs, stats.norm.pdf(xs, res.mean(), res.std()), color=RED); axes[2, j].set_title(f"{ttl}: residual distribution (red = normal fit)")
    axes[3, j].scatter(y_test.index, res, s=6, alpha=.5, color=BLUE); axes[3, j].axhline(0, color="k", lw=1); axes[3, j].set_title(f"{ttl}: residuals over time (test hours)"); axes[3, j].tick_params(axis="x", rotation=45)
axes[0, 0].legend(fontsize=7, markerscale=2); axes[1, 0].legend(fontsize=8)
plt.tight_layout(); plt.show()

rows = []
for nm, ttl in DIAG:
    p, res = PRED[nm]["test"], yt_all - PRED[nm]["test"]; terc = pd.qcut(p, 3, labels=["low", "mid", "high"], duplicates="drop")
    sd_t = pd.Series(res).groupby(terc, observed=True).std()
    rows.append({"model": ttl, "mean residual": res.mean(), "sd": res.std(), "skew": stats.skew(res), "excess kurtosis": stats.kurtosis(res),
                "Shapiro p": stats.shapiro(res)[1], "p(|res|, fitted)": stats.spearmanr(np.abs(res), p)[0],
                "sd low/mid/high fitted": " / ".join(f"{v:.2f}" for v in sd_t.values), "max |res|": np.abs(res).max()})
RESID_STATS = pd.DataFrame(rows).set_index("model"); display(RESID_STATS.style.format({c: "{:.3f}" for c in RESID_STATS.columns if c not in ("sd low/mid/high fitted", "Shapiro p")} | {"Shapiro p": "{:.1e}"}))
```



	mean residual	sd	skew	excess kurtosis	Shapiro p	$p(res , \text{fitted})$	sd low/mid/high fitted	max res
model								
Linear	0.016	0.493	0.483	5.407	4.6e-28	0.241	0.33 / 0.37 / 0.68	2.583
Poly d=2	0.016	0.447	0.259	7.146	3.4e-32	0.383	0.28 / 0.37 / 0.62	2.559
Tree depth 5	0.016	0.513	0.563	4.665	2.1e-27	0.296	0.35 / 0.41 / 0.72	2.649

Observations — residual diagnostics.


- Linear model:** the residuals-vs-fitted plot shows a **curved (U-shaped) bias** — positive residuals (under-prediction) at low *and* high fitted values, negative in the middle. This is direct evidence of **non-linearity** that a straight line cannot capture; the polynomial model's binned mean residual is essentially flat around zero.
- Tree:** predictions form horizontal bands (a piecewise-constant function, only 32 distinct values), so it cannot reproduce the smooth relation and its predicted-vs-actual plot is visibly stair-stepped; it also **flattens at $\approx 6.8 \text{ mg/m}^3$** , never predicting the highest values.
- All models are heteroscedastic:** the residual spread roughly doubles from low to high fitted values (e.g. linear sd 0.33 $\rightarrow$ 0.68), and $|residual|$ correlates positively with the fitted value ($\rho = 0.24\text{--}0.38$).
- Residuals are **heavy-tailed and right-skewed** (excess kurtosis 4.7–7.1, skew 0.26–0.56; the Shapiro test rejects normality), so confidence statements that assume Gaussian errors would be unreliable — one more reason for using the bootstrap above.
- Mean residuals are +0.016 for all models (i.e. unbiased overall); the problem is in the tails, not the centre.
- Residuals over time show clusters of large errors in winter (Nov–Feb) rather than a drift.

3.6 Where do the errors occur? — CO regime, hour of day and month

The brief only restricts the *features*, so hour-of-day is not a model input; plotting the residuals against it is a diagnostic that tells us what structure the sensor/weather features do **not** capture.

```
# (a) per-regime test RMSE and bias (from Stage 2)
fig, ax = plt.subplots(1, 2, figsize=(16, 4.2)); w = 0.8 / len(FINAL)
pr = REGIME.pivot(index="model", columns="regime", values="rmse")[bands_labels].loc[FINAL]; pb = REGIME.pivot(index="model", columns="regime",
values="bias (pred-true)")[bands_labels].loc[FINAL]
pal = sns.color_palette("tab10", len(FINAL))
for i, nm in enumerate(FINAL):
    ax[0].bar(np.arange(4) + i * w, pr.loc[nm], w, color=pal[i], label=nm); ax[1].bar(np.arange(4) + i * w, pb.loc[nm], w, color=pal[i])
for a, t in zip(ax, ["Test RMSE by CO regime", "Mean error (pred - true): negative = under-prediction"]): a.set_xticks(np.arange(4) + 0.4 - w / 2);
a.set_xticklabels(bands_labels); a.set_title(t)
ax[1].axhline(0, color="k", lw=.8); ax[0].legend(fontsize=6, ncol=2); plt.tight_layout(); plt.show()

# (b) residuals by hour of day and by month (test set)
sel = [(N_LIN_NE, "Linear", BLUE), (N_POLY_NE, f"Poly d={DEG_STAR}", ORANGE), (tree_names[5], "Tree depth 5", GREEN)]
fig, ax = plt.subplots(1, 3, figsize=(18, 4.2))
for nm, ttl, c in sel:
    r_ = pd.Series(yt_all - PRED[nm]["test"], index=y_test.index)
    ax[0].plot(r_.groupby(r_.index.hour).mean(), color=c, marker="o", ms=3, label=ttl)
    ax[1].plot(r_.groupby(r_.index.hour).apply(lambda s: np.sqrt((s ** 2).mean()))), color=c, marker="o", ms=3, label=ttl)
    mo = r_.groupby(r_.index.to_period("M")).apply(lambda s: np.sqrt((s ** 2).mean())); ax[2].plot(mo.index.astype(str), mo.values, color=c,
marker="o", label=ttl)
ax[0].axhline(0, color="k", lw=.8); ax[0].set_title("Mean residual by hour (actual - predicted)"); ax[1].set_title("RMSE by hour of day");
ax[2].set_title("RMSE by month"); ax[2].tick_params(axis="x", rotation=60)
for a in ax[:2]: a.set_xlabel("hour")
ax[0].legend(); plt.tight_layout(); plt.show()
hr = pd.DataFrame({ttl: pd.Series(yt_all - PRED[nm]["test"], index=y_test.index).groupby(y_test.index.hour).mean() for nm, ttl, _ in sel})
print("Largest mean residual by hour (rush hours) ->", hr.abs().mean(axis=1).sort_values(ascending=False).head(5).round(3).to_dict())
```



Largest mean residual by hour (rush hours) -> {8: 0.326, 20: 0.322, 19: 0.283, 7: 0.272, 6: 0.231}

Observations.

- Regime:** the RMSE rises from ≈ 0.33 – 0.40 (Low) to ≈ 0.95 – 1.33 (Very high), and every model has a large negative bias for *Very high* hours (under-prediction) — see Stage 2; the bias is -0.93 for linear versus about -0.56 for the polynomial and tuned tree (the deeper trees are less biased there, -0.41 to -0.50 , at the cost of overfitting elsewhere).
- Hour of day:** all three models share the same systematic pattern — they **under-predict in the evening (18–21 h; mean residual up to $+0.3$ mg/m³)** and **over-predict in the early morning rush (06–08 h; down to -0.3)**. Because hour-of-day is not a feature, this is structure the sensor/weather variables cannot explain (for example lagged sensor response or traffic that the sensors see only later). RMSE peaks at 19 h (0.55 – 0.72) and is smallest at 04–05 h and 11 h.
- Month:** RMSE is 2–3 $\times$ larger in **Nov–Jan (0.65 – 0.77)** than in summer (0.21 – 0.36), tracking the higher and more variable CO of the cold season; the polynomial model is the most stable.



- Together with the heteroscedasticity above, this suggests that an hour-of-day feature or a model of the error variance would be the first extensions (outside the scope of the brief).

3.7 Evaluation summary — accuracy vs. complexity vs. cost

```
SUMMARY = MASTER[["family", "complexity", "fit_time_s", "test_rmse", "test_r2", "cv_rmse", "rmse_gap (test-train)"]].astype(
    {"complexity": float, "fit_time_s": float, "test_rmse": float, "test_r2": float, "cv_rmse": float, "rmse_gap (test-train)": float},
    errors="ignore").copy()
SUMMARY["RMSE 95% CI"] = [f"[{BOOT_TBL.loc[n, 'RMSE 95% CI low']:.3f}, {BOOT_TBL.loc[n, 'RMSE 95% CI high']:.3f}]" for n in SUMMARY.index]
SUMMARY["rank (test RMSE)"] = SUMMARY["test_rmse"].rank(method="min").astype(int)
display(SUMMARY.sort_values("test_rmse").style.format(precision=4, na_rep="-"))

fig, ax = plt.subplots(1, 2, figsize=(15, 4.8))
s_ = SUMMARY.drop(index=N_MEAN)
for nm, r in s_.iterrows(): ax[0].scatter(r["complexity"], r["test_rmse"], s=60, color=fam_col[r["family"]]); ax[0].annotate(nm.split(" | ")[-1] if
r["family"] != "Decision Tree" else nm.replace("Tree | ", ""), (r["complexity"], r["test_rmse"]), fontsize=6.5, xytext=(4, 3), textcoords="offset
points")
ax[0].set_xscale("log"); ax[0].set_xlabel("number of parameters / leaves (log)"); ax[0].set_ylabel("test RMSE"); ax[0].set_title("Accuracy vs. model
complexity")
for fam, c in [("Linear", BLUE), ("Polynomial", ORANGE), ("Decision Tree", GREEN)]: ax[0].scatter([], [], color=c, label=fam)
ax[0].legend()
ax[1].bar(range(len(s_)), s_["rmse_gap (test-train)"], color=[fam_col[f] for f in s_.index]); ax[1].set_xticks(range(len(s_));
ax[1].set_xticklabels([n.replace(" | ", "\n") for n in s_.index], rotation=70, fontsize=6.5)
ax[1].set_ylabel("test RMSE - train RMSE"); ax[1].set_title("Generalisation gap (overfitting indicator)"); plt.tight_layout(); plt.show()
```

	family	complexity	fit_time_s	test_rmse	test_r2	cv_rmse	rmse_gap (test-train)	RMSE 95% CI	rank (test RMSE)
Poly(d=2) scratch (normal eq.)	Polynomial	45.0000	0.0050	0.4470	0.9013	0.4530	0.0359	[0.368, 0.532]	1
Poly(d=2) sklearn	Polynomial	45.0000	0.0056	0.4470	0.9013	–	0.0359	[0.368, 0.532]	2
Poly(d=2) scratch (Adam GD)	Polynomial	45.0000	0.6554	0.4470	0.9013	–	0.0359	[0.368, 0.532]	3
Poly(d=2)+Ridge(α=1) scratch	Polynomial+Ridge	45.0000	0.0020	0.4474	0.9012	0.4516	0.0360	[0.369, 0.533]	4
Poly(d=2)+Ridge(α=1) sklearn	Polynomial+Ridge	45.0000	0.0031	0.4474	0.9012	–	0.0360	[0.369, 0.533]	5
Linear scratch (normal eq.)	Linear	9.0000	0.0005	0.4933	0.8799	0.4882	0.0134	[0.415, 0.577]	6
Linear sklearn	Linear	9.0000	0.0011	0.4933	0.8799	–	0.0134	[0.415, 0.577]	7
Linear scratch (GD)	Linear	9.0000	0.0634	0.4933	0.8799	–	0.0134	[0.415, 0.577]	8
Tree max_depth=5	Decision Tree	32.0000	0.0113	0.5131	0.8700	0.5092	0.0728	[0.440, 0.591]	9
Tree tuned (CV, 1-SE)	Decision Tree	29.0000	0.0103	0.5157	0.8687	0.5134	0.0459	[0.444, 0.592]	10
Tree pruned (ccp, CV)	Decision Tree	21.0000	0.0490	0.5304	0.8611	0.5212	0.0679	[0.461, 0.605]	11
Tree max_depth=10	Decision Tree	693.0000	0.0166	0.5442	0.8538	0.5981	0.2838	[0.473, 0.616]	12
Tree max_depth=3	Decision Tree	8.0000	0.0070	0.5854	0.8308	0.5773	0.0339	[0.510, 0.669]	13
Tree max_depth=None	Decision Tree	4245.0000	0.0217	0.6096	0.8165	0.6593	0.6096	[0.523, 0.699]	14
Baseline mean predictor	Baseline	1.0000	0.0002	1.4232	-0.0001	1.4378	-0.0165	[1.312, 1.534]	15



Observations.

- **Accuracy vs. complexity:** the polynomial d=2 model is the most accurate with only 45 parameters; the linear model is next with 9 parameters; the best tree needs ≈ 30 leaves and is still behind. Beyond ≈ 30 leaves, extra complexity (693 and 4,245 leaves) *hurts*.
- **Generalisation gap:** linear 0.013, polynomial 0.036, trees 0.034–0.610 — it grows with tree depth.
- **Cost:** all closed-form fits take milliseconds; the from-scratch gradient-based fits cost ≈ 0.08 s (linear GD) and ≈ 0.66 s (polynomial Adam), i.e. roughly 150× slower than their closed forms for identical results, which is why the closed form is the practical choice here.

- **Ranking by test RMSE:** Poly d=2 (0.447) < Linear (0.493) < Tree depth 5 (0.513) $\approx$ tuned tree (0.516) < pruned tree (0.530) < depth 10 (0.544) < depth 3 (0.585) < None (0.610) $\ll$ mean predictor (1.423).

✅ **Stage 3 complete** — required metrics for all models, epoch and depth curves, uncertainty and residual analysis.

Stage 4 — Model Visualization

Section	Content
4.1	Decision tree plot for <code>max_depth = 3</code> (thresholds shown in original sensor units), rule list and leaf summary
4.2	Feature importances from four independent methods and their agreement
4.3	Grouped permutation importance (needed because the sensors are collinear — Stage 1)
4.4	Partial-dependence curves: how do the models use the top features?
4.5	A pollution episode reconstructed by the models

4.1 The decision tree for `max_depth = 3`


Trees are scale-invariant (verified in Stage 2), so for readability the depth-3 tree is re-fitted on the **unscaled** features: its split thresholds are then in real sensor units. The assertion confirms that it makes exactly the same predictions as the tree trained on standardised data.

```
from sklearn.tree import plot_tree, export_text
tree3 = DecisionTreeRegressor(max_depth=3, random_state=SEED).fit(X_train, y_train)
assert np.allclose(tree3.predict(X_test), PRED[treenames[3]]["test"]); print("Raw-unit tree predictions == standardised-data tree ✓ | leaves:",
tree3.get_n_leaves(), "| depth:", tree3.get_depth())

fig, ax = plt.subplots(figsize=(24, 9))
plot_tree(tree3, feature_names=FEATURES, filled=True, rounded=True, precision=2, fontsize=11, impurity=True, proportion=False, ax=ax)
ax.set_title("Decision Tree Regressor, max_depth = 3 (value = mean CO(GT) mg/m³ of the node; squared_error = node MSE; darker = higher CO)",
fontsize=13, fontweight="bold")
plt.tight_layout(); plt.show()
print(export_text(tree3, feature_names=FEATURES, decimals=1))
```

Raw-unit tree predictions == standardised-data tree ✓ | leaves: 8 | depth: 3

Decision Tree Regressor, max_depth = 3 (value = mean CO(GT) mg/m³ of the node; squared_error = node MSE; darker = higher CO)



```

|--- PT08.S2(NMHC) <= 1120.5
| |--- PT08.S2(NMHC) <= 846.5
| | |--- PT08.S2(NMHC) <= 703.5
| | | |--- value: [0.7]
| | |--- PT08.S2(NMHC) > 703.5
| | | |--- value: [1.2]
| |--- PT08.S2(NMHC) > 846.5
| | |--- PT08.S2(NMHC) <= 993.5
| | | |--- value: [1.8]
| | |--- PT08.S2(NMHC) > 993.5
| | | |--- value: [2.4]
|--- PT08.S2(NMHC) > 1120.5
| |--- PT08.S2(NMHC) <= 1400.0
| | |--- PT08.S2(NMHC) <= 1224.5
| | | |--- value: [3.1]
| | |--- PT08.S2(NMHC) > 1224.5
| | | |--- value: [3.9]
| |--- PT08.S2(NMHC) > 1400.0
| | |--- PT08.S2(NMHC) <= 1587.0
| | | |--- value: [5.2]
| | |--- PT08.S2(NMHC) > 1587.0
| | | |--- value: [7.1]
  
```

Observations — the depth-3 tree.

- All 7 splits use a single feature, **PT08.S2(NMHC)** ; no other sensor or weather variable enters a depth-3 tree. The tree is simply a *step function of one sensor*.
- The thresholds (in sensor units) are 704, 846, 994, 1,121, 1,225, 1,400 and 1,587, and the predicted CO rises monotonically along them from 0.72 to 7.11 mg/m³ — a single dominant, monotone dose–response.
- The root split ($S2 \leq 1,120.5$) already separates the data into a low-CO group (mean 1.49, MSE 0.57) and a high-CO group (mean 3.99, MSE 1.75).


```
def leaf_rules(t, names):
    tr, out = t.tree_, {}
    def walk(n, conds):
        if tr.children_left[n] == -1: out[n] = " & ".join(conds) if conds else "(root)"; return
        f, th = names[tr.feature[n]], tr.threshold[n]
        walk(tr.children_left[n], conds + [f"{f} <= {th:.0f}" if th > 50 else f"{f} <= {th:.2f}"]); walk(tr.children_right[n], conds + [f"{f} > {th:.0f}" if th > 50 else f"{f} > {th:.2f}"])
        walk(0, []); return out

rules = leaf_rules(tree3, FEATURES); leaf_tr, leaf_te = tree3.apply(X_train), tree3.apply(X_test); rows = []
for lf, rule in rules.items():
    mt, me = leaf_tr == lf, leaf_te == lf
    rows.append({"leaf": lf, "rule": rule, "n_train": int(mt.sum()), "mean CO (train)": y_train[mt].mean(), "sd CO (train)": y_train[mt].std(),
                "n_test": int(me.sum()), "test RMSE in leaf": rmse(y_test[me], np.full(me.sum(), tree3.tree_.value[lf].ravel()[0])) if me.sum() else
                np.nan})
LEAVES = pd.DataFrame(rows).set_index("leaf").sort_values("mean CO (train)"); display(LEAVES.style.format(precision=3))
used = pd.Series(tree3.tree_.feature[tree3.tree_.feature >= 0]).map(lambda i: FEATURES[i]).value_counts()
print("Features used for splitting in the depth-3 tree:", used.to_dict())

# the partition of the dominant sensor axis
fig, ax = plt.subplots(figsize=(11, 4.5)); sidx = FEATURES.index("PT08.S2(NMHC)")
ax.scatter(X_train["PT08.S2(NMHC)"], y_train, s=4, alpha=.25, color=GREY)
thr = sorted(set(tree3.tree_.threshold[tree3.tree_.feature == sidx]))
for th in thr: ax.axvline(th, color=RED, ls="--", lw=1.2)
gx = np.linspace(X_train["PT08.S2(NMHC)"].min(), X_train["PT08.S2(NMHC)"].max(), 400); gdf = pd.DataFrame(np.tile(X_train.mean().values, (400, 1)),
columns=FEATURES); gdf["PT08.S2(NMHC)"] = gx
ax.step(gx, tree3.predict(gdf), color=BLUE, lw=2.5, where="post", label="depth-3 tree prediction (other features at mean)")
ax.set_xlabel("PT08.S2(NMHC) sensor signal"); ax.set_ylabel("CO(GT) mg/m³"); ax.set_title("The depth-3 tree is a 8-step function of PT08.S2 (red = split thresholds)"); ax.legend(); plt.tight_layout(); plt.show()
```

	rule	n_train	mean CO (train)	sd CO (train)	n_test	test RMSE in leaf
leaf						
3	PT08.S2(NMHC) <= 1120 & PT08.S2(NMHC) <= 846 & PT08.S2(NMHC) <= 704	1167	0.724	0.321	294	0.367
4	PT08.S2(NMHC) <= 1120 & PT08.S2(NMHC) <= 846 & PT08.S2(NMHC) > 704	1164	1.226	0.438	290	0.452
6	PT08.S2(NMHC) <= 1120 & PT08.S2(NMHC) > 846 & PT08.S2(NMHC) <= 994	1185	1.821	0.463	314	0.483
7	PT08.S2(NMHC) <= 1120 & PT08.S2(NMHC) > 846 & PT08.S2(NMHC) > 994	845	2.435	0.541	201	0.684
10	PT08.S2(NMHC) > 1120 & PT08.S2(NMHC) <= 1400 & PT08.S2(NMHC) <= 1224	594	3.083	0.653	144	0.681
11	PT08.S2(NMHC) > 1120 & PT08.S2(NMHC) <= 1400 & PT08.S2(NMHC) > 1224	581	3.935	0.727	148	0.820
13	PT08.S2(NMHC) > 1120 & PT08.S2(NMHC) > 1400 & PT08.S2(NMHC) <= 1587	250	5.162	0.792	60	0.895
14	PT08.S2(NMHC) > 1120 & PT08.S2(NMHC) > 1400 & PT08.S2(NMHC) > 1587	89	7.109	1.612	18	1.132

Features used for splitting in the depth-3 tree: {'PT08.S2(NMHC)': 7}



Observations — leaf summary.

- The eight leaves partition the training hours into groups of 89 to 1,185 samples with mean CO from 0.72 to 7.11 mg/m³; the test RMSE inside each leaf grows with CO (0.37 in the lowest leaf up to 1.13 in the highest), mirroring the heteroscedasticity found in Stage 3.
- The highest leaf (S2 > 1,587) contains only **89 training hours (1.5 %)** with a large spread (sd 1.61): it is the tree's weak spot, because a single constant has to represent the rare extreme episodes.

4.2 Feature importances — which sensor / environment variable contributes most?

Four independent methods are compared, because each has known blind spots (impurity importance is biased toward features used for early splits; permutation importance splits credit among correlated features; coefficients are only meaningful for linear models; drop-one retraining is expensive but direct):

1. **Impurity (MDI) importance** of the trees (depth 5 and tuned);
2. **Permutation importance** on the test set (increase in RMSE when a feature is shuffled) — tuned tree, linear, polynomial;
3. **Absolute standardised coefficients** of the linear model (with the bootstrap CIs from Stage 2);
4. **Drop-one ablation**: increase in CV RMSE when the feature is removed and the model re-trained (Stage 2).

```
fm = FEATURES; pk = f"Poly(d={DEG_STAR})"
IMP = pd.DataFrame({
    "MDI - tree depth 5": IMP_TREE[tree_names[5]], "MDI - tree tuned": IMP_TREE[N_TREE_T],
    "Permutation - tree tuned": PERM[N_TREE_T]["mean"], "Permutation - linear": PERM[N_LIN_NE]["mean"], f"Permutation - poly d={DEG_STAR}":
    PERM[N_POLY_NE]["mean"],
    "|std. coef| - linear": COEF_BOOT["coef"].abs(),
    "Drop-one ΔCV RMSE - linear": pd.Series({f: tbl.loc[f"Drop {f}", "Linear"] - base["Linear"] for f in fm}),
    "Drop-one ΔCV RMSE - tree": pd.Series({f: tbl.loc[f"Drop {f}", "Tree (tuned)"] - base["Tree (tuned)"] for f in fm}))
IMP_N = IMP.clip(lower=0); IMP_N = IMP_N / IMP_N.sum() # normalised to shares of the total
RANK = IMP.rank(ascending=False, method="min").astype(int)
print("Raw importance values:"); display(IMP.style.format(precision=3))
print("Normalised importance (share of the total, negatives clipped to 0):"); display(IMP_N.style.format("{:.1%}").background_gradient(cmap="Blues",
axis=0))
print("Rank per method (1 = most important):"); display(RANK.assign(**{"mean rank": RANK.mean(axis=1)}).sort_values("mean
rank").style.format(precision=1))
rho = IMP.corr(method="spearman"); print("Spearman rank agreement between methods: min = %.2f, median = %.2f" % (rho.values[np.triu_indices(len(rho),
1)].min(), np.median(rho.values[np.triu_indices(len(rho), 1)])))
```

Raw importance values:

	MDI - tree depth 5	MDI - tree tuned	Permutation - tree tuned	Permutation - linear	Permutation - poly d=2	std. coef - linear	Drop-one ΔCV RMSE - linear	Drop-one ΔCV RMSE - tree
PT08.S1(CO)	0.003	0.002	0.003	0.247	0.418	0.394	0.022	-0.009
PT08.S2(NMHC)	0.958	0.970	1.418	1.452	2.373	1.366	0.119	0.151
PT08.S3(NOx)	0.007	0.000	0.000	0.075	0.088	0.199	0.007	-0.003
PT08.S4(NO2)	0.004	0.000	0.000	0.067	1.073	0.174	0.001	0.000
PT08.S5(O3)	0.000	0.000	0.000	0.010	0.063	0.077	-0.002	-0.001
T	0.027	0.028	0.092	0.119	0.953	0.244	0.004	0.017
RH	0.000	0.000	0.000	0.003	0.833	0.030	-0.000	0.000
AH	0.000	0.000	0.000	0.010	0.918	0.087	-0.001	0.000

Normalised importance (share of the total, negatives clipped to 0):

	MDI - tree depth 5	MDI - tree tuned	Permutation - tree tuned	Permutation - linear	Permutation - poly d=2	std. coef - linear	Drop-one ΔCV RMSE - linear	Drop-one ΔCV RMSE - tree
PT08.S1(CO)	0.3%	0.2%	0.2%	12.5%	6.2%	15.3%	14.3%	0.0%
PT08.S2(NMHC)	95.8%	97.0%	93.7%	73.2%	35.3%	53.1%	77.6%	89.8%
PT08.S3(NOx)	0.7%	0.0%	0.0%	3.8%	1.3%	7.7%	4.4%	0.0%
PT08.S4(NO2)	0.4%	0.0%	0.0%	3.4%	16.0%	6.8%	0.9%	0.0%
PT08.S5(O3)	0.0%	0.0%	0.0%	0.5%	0.9%	3.0%	0.0%	0.0%
T	2.7%	2.8%	6.1%	6.0%	14.2%	9.5%	2.8%	10.2%
RH	0.0%	0.0%	0.0%	0.1%	12.4%	1.2%	0.0%	0.0%
AH	0.0%	0.0%	0.0%	0.5%	13.7%	3.4%	0.0%	0.0%

Rank per method (1 = most important):


	MDI - tree depth 5	MDI - tree tuned	Permutation - tree tuned	Permutation - linear	Permutation - poly d=2	std. coef - linear	Drop-one ΔCV RMSE - linear	Drop-one ΔCV RMSE - tree	mean rank
PT08.S2(NMHC)	1	1	1	1	1	1	1	1	1.0
T	2	2	2	3	3	3	4	2	2.6
PT08.S1(CO)	5	3	3	2	6	2	2	8	3.9
PT08.S4(NO2)	4	4	4	5	2	5	5	4	4.1
PT08.S3(NOx)	3	4	4	4	7	4	3	7	4.5
AH	6	4	4	7	4	6	7	4	5.2
RH	6	4	4	8	5	8	6	3	5.5
PT08.S5(O3)	6	4	4	6	8	7	8	6	6.1

Spearman rank agreement between methods: min = 0.07, median = 0.73

Observations — feature importance (answer to "which variable contributes most?").

- **PT08.S2(NMHC)** is rank 1 in all 8 methods, with a mean share of **77 %** of total importance (tree methods 94–97 %, linear methods 53–78 %).
- The runner-ups are **T (mean share 6.8 %, rank 2–4)** and **PT08.S1(CO)** (**6.1 %**); PT08.S1 matters for the linear model (rank 2 in coefficients, permutation and drop-one) but **not for the trees** ($\approx 0 %$), since it is redundant with PT08.S2 ($r = 0.89$).
- PT08.S4(NO₂), PT08.S3(NO_x), AH, RH and PT08.S5(O₃) contribute little ($\leq 3.4 %$ on average).
- **Agreement:** the median Spearman rank correlation between methods is 0.73, but the minimum is 0.07. The disagreement is almost entirely about the *minor* features; the top feature is unanimous.
- **A caveat on the polynomial permutation column:** it assigns surprisingly large importance to RH, AH, T and PT08.S4 (0.8–1.1 mg/m³ Δ RMSE, more than the linear model gives to anything except S2), although retraining without any of them (drop-one, Stage 2) costs the polynomial *nothing*. Shuffling a single weather column creates physically impossible combinations (e.g. a humidity that does not match the temperature), and the degree-2 model — with its many cross-terms between collinear inputs — reacts violently to such unrealistic inputs. That is a sign of **fragility off the data manifold**, not of real predictive reliance, and it is why drop-one and the on-data plots below are the more trustworthy evidence.

```
order = IMP_N.mean(axis=1).sort_values().index
fig, ax = plt.subplots(1, 3, figsize=(21, 5.2), gridspec_kw={"width_ratios": [1.1, 1.3, 1]})
sns.heatmap(IMP_N.loc[order[:~1]], annot=True, fmt=".0%", cmap="Blues", cbar=False, ax=ax[0], annot_kws={"size": 8}); ax[0].set_title("Importance
share by method"); ax[0].tick_params(axis="x", rotation=60)
pm = pd.DataFrame({"tree depth 3 (MDI)": IMP_TREE[tree_names[3]], "tree tuned (MDI)": IMP_TREE[N_TREE_T], "linear |coef| (norm.)": IMP_N["|std. coef| -
linear"],
                    f"poly perm. (norm.)": IMP_N[f"Permutation - poly d={DEG_STAR}"]}).loc[order]
pm.plot.barh(ax=ax[1], color=[BLUE, GREEN, ORANGE, RED], width=.8); ax[1].set_title("Normalised importance: tree vs. linear vs. polynomial");
ax[1].set_xlabel("share of total"); ax[1].legend(fontsize=8)
cb = COEF_BOOT.loc[order]; ax[2].barh(cb.index, cb["coef"], xerr=[cb["coef"] - cb["ci_low"], cb["ci_high"] - cb["coef"]], color=[BLUE if c > 0 else
ORANGE for c in cb["coef"]]); ax[2].axvline(0, color="k", lw=.8)
ax[2].set_title("Linear coefficients (std.)  $\pm$  95% bootstrap CI\nblue = raises CO, orange = lowers CO"); plt.tight_layout(); plt.show()
top = IMP_N.mean(axis=1).sort_values(ascending=False); print("Mean share across the 8 methods:", {k: f"{v:.1%}" for k, v in top.items()})
```



Mean share across the 8 methods: {'PT08.S2(NMHC)': '77.0%', 'T': '6.8%', 'PT08.S1(CO)': '6.1%', 'PT08.S4(NO₂)': '3.4%', 'PT08.S3(NO_x)': '2.2%', 'AH': '2.2%', 'RH': '1.7%', 'PT08.S5(O₃)': '0.6%'}

4.3 Grouped permutation importance (handling collinearity)

When correlated features are permuted one at a time the model can still lean on their twins, so single-feature importance *understates* groups of redundant sensors. Here, columns of a group are shuffled **jointly** (same permutation) on the test set and the RMSE increase is recorded (30 repeats).


```

S1, S2, S3, S4, S5 = SENSOR_COLS
GROUPS = {"PT08.S2 (NMHC sensor) alone": [S2], "PT08.S1 + S5 (CO / O3 sensors)": [S1, S5], "PT08.S1 + S2 + S5 (correlated cluster)": [S1, S2, S5],
          "PT08.S3 + S4 (NOx / NO2 sensors)": [S3, S4], "All 5 sensors": SENSOR_COLS, "Temperature T": [T], "Humidity (RH + AH)": [RH, AH], "All
weather (T, RH, AH)": WEATHER_COLS}
def grouped_perm(model, cols, n_repeats=30, seed=SEED):
    rg = np.random.default_rng(seed); b0 = rmse(y_test, model.predict(X_test_s)); d = []
    for _ in range(n_repeats):
        Xp = X_test_s.copy(); Xp[cols] = X_test_s[cols].values[rg.permutation(len(Xp))]; d.append(rmse(y_test, model.predict(Xp)) - b0)
    return np.mean(d), np.std(d)
GP_MODELS = {"Linear": N_LIN_NE, f"Poly d={DEG_STAR}": N_POLY_NE, "Tree tuned": N_TREE_T}
GPERM = pd.DataFrame({m_: {g: grouped_perm(FITTED[nm], c)[0] for g, c in GROUPS.items()} for m_, nm in GP_MODELS.items()})
GPERM_SD = pd.DataFrame({m_: {g: grouped_perm(FITTED[nm], c)[1] for g, c in GROUPS.items()} for m_, nm in GP_MODELS.items()})
display(GPERM.style.format(precision=3).background_gradient(cmap="Reds", axis=None).set_caption("Increase in test RMSE (mg/m³) when the group is
shuffled jointly"))
fig, ax = plt.subplots(figsize=(11, 4.8)); w = 0.27
for i, (m_, c) in enumerate(zip(GPERM.columns, [BLUE, ORANGE, GREEN])): ax.barh(np.arange(len(GPERM)) + (i - 1) * w, GPERM[m_], w, xerr=GPERM_SD[m_],
color=c, label=m_)
ax.set_yticks(range(len(GPERM))); ax.set_yticklabels(GPERM.index); ax.invert_yaxis(); ax.set_xlabel("ΔRMSE when shuffled (mg/m³)");
ax.set_title("Grouped permutation importance (test set)"); ax.legend(); plt.tight_layout(); plt.show()

```

Increase in test RMSE (mg/m³) when the group is shuffled jointly

	Linear	Poly d=2	Tree tuned
PT08.S2 (NMHC sensor) alone	1.456	2.394	1.434
PT08.S1 + S5 (CO / O3 sensors)	0.177	0.242	0.003
PT08.S1 + S2 + S5 (correlated cluster)	1.855	2.242	1.442
PT08.S3 + S4 (NOx / NO2 sensors)	0.188	1.397	0.000
All 5 sensors	1.468	1.544	1.442
Temperature T	0.118	0.953	0.092
Humidity (RH + AH)	0.011	1.264	0.000
All weather (T, RH, AH)	0.063	0.238	0.092



Observations — grouped permutation.

- Sensors drive everything.** Shuffling all five sensors jointly costs the linear model +1.47 mg/m³ RMSE (tree +1.44) while shuffling all weather variables jointly costs only +0.06 (linear) / +0.09 (tree).
- For the linear model the correlated cluster S1 + S2 + S5 (+1.86) is clearly more important than S2 alone (+1.46): its twins add information that single-feature permutation hides. For the tree the cluster adds nothing beyond S2 (+1.44 vs. +1.42), because the tree only ever uses S2.
- T (+0.12 linear, +0.09 tree) is the only weather variable with a measurable effect; humidity (RH + AH) is negligible for the linear model and the tree (+0.011 and 0.000).
- The polynomial model again shows large values for non-S2 groups (e.g. S3 + S4 +1.40, humidity +1.26): the same off-manifold fragility discussed above.

4.4 Partial dependence — how do the models *use* the important features?


Average model prediction as one feature is swept across its observed range (2nd–98th percentile) with all other features left at their observed values. This visualises the **shape** of each learnt relationship: a straight line (linear), a smooth curve (polynomial) or a staircase (tree). Axes are in original units; the shaded histogram shows where the data lie.

```
top4 = list(top.index[:2]) + [f for f in ["PT08.S1(CO)", "PT08.S3(NOx)"] if f not in top.index[:2]]
top4 = ["PT08.S2(NMHC)", "PT08.S1(CO)", "T", "PT08.S3(NOx)"]
mu, sc_ = dict(zip(FEATURES, scaler.mean_)), dict(zip(FEATURES, scaler.scale_))
PDP_MODELS = {"Linear": (N_LIN_NE, BLUE), f"Poly d={DEG_STAR3}": (N_POLY_NE, ORANGE), "Tree depth 3": (tree_names[3], GREY), "Tree tuned": (N_TREE_T, GREEN)}

def pdp_1d(model, feat, grid_s):
    out = []
    for g in grid_s:
        Xg = X_train_s.copy(); Xg[feat] = g; out.append(model.predict(Xg).mean())
    return np.array(out)

fig, axes = plt.subplots(1, 4, figsize=(21, 4.4)); PDP = {}
for a, f in zip(axes, top4):
    lo, hi = np.percentile(X_train_s[f], [2, 98]); gs_ = np.linspace(lo, hi, 40); xo = gs_ * sc_[f] + mu[f]
    for lbl, (nm, c) in PDP_MODELS.items():
        PDP[(f, lbl)] = pdp_1d(FITTED[nm], f, gs_); a.plot(xo, PDP[(f, lbl)], color=c, lw=2.2, label=lbl, drawstyle="steps-mid" if "Tree" in lbl else "default")
        a2 = a.twinx(); a2.hist(X_train[f], bins=40, color=GREY, alpha=.15); a2.set_yticks([]); a2.spines[["top", "right"]].set_visible(False)
        a.set_xlabel(f); a.set_ylabel("average predicted CO(GT)"); a.set_title(f"PDP: {f}")
axes[0].legend(fontsize=8); plt.tight_layout(); plt.show()

# strength of each feature's average effect (range of the PDP)
display(pd.DataFrame({lbl: {f: np.ptp(PDP[(f, lbl)]) for f in top4} for lbl in PDP_MODELS}).style.format(precision=3).set_caption("Range of the partial-dependence curve (mg/m³ of average CO change across the feature's 2–98 % range)"))
```




Range of the partial-dependence curve (mg/m³ of average CO change across the feature's 2–98 % range)

	Linear	Poly d=2	Tree depth 3	Tree tuned
PT08.S2(NMHC)	5.355	4.535	4.439	4.812
PT08.S1(CO)	1.533	1.037	0.000	0.048
T	0.953	3.565	0.000	0.417
PT08.S3(NOx)	0.825	0.666	0.000	0.000

Observations — partial dependence.

- PT08.S2** : the linear model gives a straight line (≈ 5.4 mg/m³ rise across the central 96 % of the sensor range) and the trees follow the *same* line as a staircase (rise 4.4–4.8): the average effect of the dominant sensor is essentially **linear**, which is why the tree cannot beat the line.
- PT08.S1** , **PT08.S3** : the tree ignores them (flat curves) while the linear model (and the polynomial, with some curvature) assigns a positive effect to both.
- T** : the linear model gives a mild negative effect (-0.95 mg/m³ over the range), the tuned tree a small step down, depth 3 none.
- Polynomial model — a warning**: its partial-dependence curves are concave and even negative (about -2 mg/m³ at low S2), which is physically implausible. A PDP evaluates the model at feature combinations that never occur (e.g. low S2 together with the observed values of the correlated S1, S5), where the degree-2 cross-terms do not cancel as they do on real data. The next plot shows that **on the actual data the polynomial model is the most accurate**, so this is a limitation of the PDP *and* a sign that the polynomial model should not be trusted for inputs that break the sensor correlations.


```
# On-data view (binned means, a.k.a. M-plot): average ACTUAL and PREDICTED CO per quantile bin of each feature on the test set.
# Unlike the PDP it never evaluates a model at unrealistic feature combinations, so it shows how each model behaves where the data actually are.
fig, axes = plt.subplots(1, 4, figsize=(21, 4.2)); MPLLOT = {}
for a, f in zip(axes, top4):
    bins = pd.qcut(X_test[f], 12, duplicates="drop"); xm = X_test[f].groupby(bins, observed=True).mean()
    act = y_test.groupby(bins, observed=True).mean(); a.plot(xm, act, "k-o", lw=2.2, ms=4, label="actual", zorder=5); MPLLOT[(f, "actual")] = act.values
    for lbl, (nm, c) in PDP_MODELS.items():
        pm_ = pd.Series(PRED[nm]["test"], index=y_test.index).groupby(bins, observed=True).mean(); MPLLOT[(f, lbl)] = pm_.values
        a.plot(xm, pm_, color=c, lw=1.6, marker="s", ms=3, label=lbl)
    a.set_xlabel(f); a.set_ylabel("mean CO(GT) in bin"); a.set_title(f"On-data binned means: {f}")
axes[0].legend(fontsize=8); plt.tight_layout(); plt.show()
print("Max |binned mean prediction - binned mean actual| per model (mg/m³), averaged over the 4 features:")
print({lbl: round(float(np.mean([np.abs(MPLLOT[(f, lbl)] - MPLLOT[(f, "actual")]).max() for f in top4])), 3) for lbl in PDP_MODELS})
```



Max |binned mean prediction - binned mean actual| per model (mg/m³), averaged over the 4 features:
{'Linear': 0.227, 'Poly d=2': 0.09, 'Tree depth 3': 0.264, 'Tree tuned': 0.13}

Observations — on-data binned means. This view uses only realistic inputs. All models follow the actual mean CO closely along PT08.S2 and PT08.S1 (monotone, nearly straight) and PT08.S3 (a convex decay, which the straight line misses at both ends). Averaged over the four features, the largest bin-wise deviation from the actual means is **0.09 mg/m³ for the polynomial**, **0.13 for the tuned tree**, **0.23 for the linear model** and **0.26 for the depth-3 tree**. The linear model's errors are concentrated at the extremes of the sensor ranges, where it **under-estimates at both ends** (e.g. lowest PT08.S2 bin 0.3 vs. 0.6 actual, highest 4.9 vs. 5.2), which is the curvature already seen in the Stage 3 residuals. For T, the actual relation is irregular (non-monotone), and only the polynomial and tuned tree follow its wiggles; the depth-3 tree, which ignores T, is the farthest off.

```
# 2-D partial dependence of the two most important drivers: does a model capture their interaction?
f1, f2 = "PT08.S2(NMHC)", "T"; bg = X_train_s.sample(800, random_state=SEED)
g1 = np.linspace(*np.percentile(X_train_s[f1], [2, 98]), 28); g2 = np.linspace(*np.percentile(X_train_s[f2], [2, 98]), 28)
def pdp_2d(model):
    Z = np.zeros((len(g2), len(g1)))
    for i, b in enumerate(g2):
        for j, a_ in enumerate(g1):
            Xg = bg.copy(); Xg[f1], Xg[f2] = a_, b; Z[i, j] = model.predict(Xg).mean()
    return Z
fig, axes = plt.subplots(1, 3, figsize=(18, 4.6)); Zs = {l: pdp_2d(FITTED[PDP_MODELS[l][0]]) for l in ["Linear", f"Poly d={DEG_STAR}", "Tree tuned"]}
vmin, vmax = min(z.min() for z in Zs.values()), max(z.max() for z in Zs.values())
for a, (l, Z) in zip(axes, Zs.items()):
    im = a.pcolormesh(g1 * sc_[f1] + mu[f1], g2 * sc_[f2] + mu[f2], Z, cmap="viridis", vmin=vmin, vmax=vmax, shading="auto"); a.set_xlabel(f1);
    a.set_ylabel(f2); a.set_title(f"{l}: predicted CO(GT)")
fig.colorbar(im, ax=axes, label="mg/m³", shrink=.9); plt.show()
Zl = Zs["Linear"]; add_resid = {l: np.abs(Z - (Z.mean(0, keepdims=True) + Z.mean(1, keepdims=True) - Z.mean())) for l, Z in Zs.items()}
print("Mean |non-additive part| of each model's 2-D surface (0 = purely additive, no interaction):", {k: round(v, 4) for k, v in add_resid.items()})
```




Mean |non-additive part| of each model's 2-D surface ($\emptyset$ = purely additive, no interaction): {'Linear': np.float64(0.0), 'Poly d=2': np.float64(0.1954), 'Tree tuned': np.float64(0.1311)}

Observations. The linear surface is a plain tilted plane (no interaction by construction: non-additive part 0.000). The polynomial and the tuned tree both contain interactions (non-additive part 0.195 and 0.131 mg/m^3), with the tree producing blocky regions where T modulates the effect of PT08.S2 (higher CO at low temperature for high S2 values). The polynomial surface is smooth but, as noted above, partly extrapolates to unrealistic sensor/weather combinations (e.g. strongly negative predictions at low S2 and low T), so its absolute shape should not be over-interpreted.

4.5 A pollution episode as seen by the models

```
X_all_s = pd.DataFrame(scaler.transform(data[FEATURES]), columns=FEATURES, index=data.index)
peak = y_test.idxmax(); win = slice(peak - pd.Timedelta(days=4), peak + pd.Timedelta(days=4)); W = data.loc[win]; Xw = X_all_s.loc[win]
is_test = W.index.isin(y_test.index)
fig, ax = plt.subplots(figsize=(17, 5)); ax.plot(W.index, W[TARGET], color="k", lw=1.6, label="actual CO(GT)")
for lbl, (nm, c) in {"Linear": (N_LIN_NE, BLUE), "Poly d={DEG_STAR}": (N_POLY_NE, ORANGE), "Tree tuned": (N_TREE_T, GREEN)}.items():
    p = FITTED[nm].predict(Xw); ax.plot(W.index, p, color=c, lw=1.4, alpha=.9, label=lbl); ax.scatter(W.index[is_test], p[is_test], s=22, color=c,
zorder=3)
ax.scatter(W.index[is_test], W[TARGET][is_test], s=26, color="k", zorder=4, label="held-out (test) hours")
ax.set_ylabel("CO(GT) mg/m³"); ax.set_title(f"Eight days around the highest test CO value ({peak:%d %b %Y}); dots = hours not used for training");
ax.legend(ncol=5, fontsize=9); plt.tight_layout(); plt.show()
```



Observations. The models reproduce the timing and shape of the daily CO cycles and of the multi-hour pollution episodes well, including on held-out hours. At the highest peak of the test set (8.7 mg/m^3 on 17 Jan 2005) **all models under-predict** (polynomial 7.3 is closest, tree 7.1, linear 6.4), in line with the under-prediction of extremes found in Stages 2–3. They also miss abrupt drops that are not visible in the sensors (e.g. the last hours on 21 Jan).

✅ **Stage 4 complete** — depth-3 tree plot, rules and leaf summary, consolidated feature-importance analysis, grouped permutation importance, partial dependence and an episode reconstruction.

Stage 5 — Baseline Comparison: Linear Regression vs. Decision Tree

Linear regression is the **baseline**; the question is whether a Decision Tree (and, as an extra reference, the degree-2 polynomial) captures structure the line misses. The comparison combines accuracy, uncertainty, *where* each model wins, extrapolation behaviour and an overview of every ablation from Stage 2.

Section	Content
5.1	Head-to-head: Linear vs. every Decision Tree configuration, with paired uncertainty
5.2	Is there non-linear structure? Statistical test, polynomial terms, a hybrid "linear + tree" probe
5.3	What each model family can represent (synthetic demonstration, clearly labelled)
5.4	Where does each model win? (by sensor level) and extrapolation to unseen conditions
5.5	Ablation ledger — every ablation in one table
5.6	Scorecard



5.1 Head-to-head: Linear vs. Decision Tree

```
LIN = N_LIN_NE; TREES = [tree_names[3], tree_names[5], tree_names[10], tree_names[None], N_TREE_T, N_TREE_P]
rows = []
for nm in TREES:
    d = BOOT[nm][0] - BOOT[LIN][0]; lo, hi = ci(d)
    fold_wins = int((FOLD_RMSE.iloc[:5][["Linear"]] < FOLD_RMSE.iloc[:5][{tree_names[3]: "Tree depth 3", tree_names[5]: "Tree depth 5", tree_names[10]:
"Tree depth 10", tree_names[None]: "Tree depth None", N_TREE_T: "Tree (tuned)"}]).sum()) if nm != N_TREE_P else np.nan
    rows.append({"Decision Tree": nm.replace("Tree | ", ""), "leaves": RESULTS[nm]["n_leaves"], "train RMSE": RESULTS[nm]["train_rmse"], "test RMSE":
RESULTS[nm]["test_rmse"], "test R²": RESULTS[nm]["test_r2"],
                "CV RMSE": RESULTS[nm]["cv_rmse"], "Δ test RMSE vs Linear": RESULTS[nm]["test_rmse"] - RESULTS[LIN]["test_rmse"], "95% CI low": lo,
"95% CI high": hi,
                "P(tree worse)": np.mean(d > 0), "folds Linear wins (of 5)": fold_wins})
H2H = pd.DataFrame(rows).set_index("Decision Tree")
print(f"Reference — Linear regression: train RMSE {RESULTS[LIN]['train_rmse']:.4f}, test RMSE {RESULTS[LIN]['test_rmse']:.4f}, test R² {RESULTS[LIN]
['test_r2']:.4f}, CV RMSE {RESULTS[LIN]['cv_rmse']:.4f}, 9 parameters")
display(H2H.style.format(precision=4, na_rep="-").background_gradient(subset=["Δ test RMSE vs Linear"], cmap="Reds"))

fig, ax = plt.subplots(1, 2, figsize=(16, 4.6)); names = ["Linear"] + [t.replace("Tree | ", "").replace("max_depth=", "depth ") for t in TREES]; vals =
[LIN] + TREES; x = np.arange(len(vals)); w = .38
ax[0].bar(x - w / 2, [RESULTS[n]["train_rmse"] for n in vals], w, color=BLUE, label="train"); ax[0].bar(x + w / 2, [RESULTS[n]["test_rmse"] for n in
vals], w, color=ORANGE, label="test")
ax[0].set_xticks(x); ax[0].set_xticklabels(names, rotation=25, ha="right"); ax[0].set_ylabel("RMSE (mg/m³)"); ax[0].set_title("Linear vs. Decision
Trees: train and test RMSE"); ax[0].legend()
ax[0].axhline(RESULTS[LIN]["test_rmse"], color=ORANGE, ls=":", lw=1)
for i, nm in enumerate(TREES):
    d = BOOT[nm][0] - BOOT[LIN][0]; lo, hi = ci(d); ax[1].errorbar(d.mean(), i, xerr=[d.mean() - lo, [hi - d.mean()]], fmt="o", color=GREEN,
capsize=3)
ax[1].axvline(0, color="k", lw=1); ax[1].set_yticks(range(len(TREES))); ax[1].set_yticklabels([t.replace("Tree | ", "") for t in TREES]);
ax[1].invert_yaxis()
ax[1].set_xlabel("test RMSE(tree) - test RMSE(linear) (> 0: linear better)"); ax[1].set_title("Paired week-bootstrap difference (95% CI)");
plt.tight_layout(); plt.show()
```

Reference – Linear regression: train RMSE 0.4799, test RMSE 0.4933, test R^2 0.8799, CV RMSE 0.4882, 9 parameters

	leaves	train RMSE	test RMSE	test R^2	CV RMSE	Δ test RMSE vs Linear	95% CI low	95% CI high	P(tree worse)	fold Linear wins (of 5)
Decision Tree										
max_depth=3	8	0.5515	0.5854	0.8308	0.5773	0.0922	0.0528	0.1279	1.0000	5.0000
max_depth=5	32	0.4403	0.5131	0.8700	0.5092	0.0198	-0.0006	0.0393	0.9730	5.0000
max_depth=10	693	0.2603	0.5442	0.8538	0.5981	0.0509	0.0168	0.0872	1.0000	5.0000
max_depth=None	4245	0.0000	0.6096	0.8165	0.6593	0.1163	0.0776	0.1537	1.0000	5.0000
tuned (CV, 1-SE)	29	0.4698	0.5157	0.8687	0.5134	0.0224	0.0001	0.0443	0.9750	4.0000
pruned (ccp, CV)	21	0.4625	0.5304	0.8611	0.5212	0.0371	0.0144	0.0604	1.0000	–



Observations – Linear vs. Decision Tree.

- **Linear regression is more accurate than every Decision Tree configuration**, on every measure of the held-out error: test RMSE 0.493 (R^2 0.880) versus 0.513–0.610 (R^2 0.817–0.870).
- The gap is *small for the best tree and large for the others*: depth 5 is +0.020 RMSE worse (CI [–0.001, +0.039], borderline), the tuned tree +0.022 (CI [+0.0001, +0.044]), the pruned tree +0.037, depth 10 +0.051, depth 3 +0.092 and None +0.116 (the last four CIs exclude 0).
- Linear wins in **5 of 5 CV folds** against every fixed-depth tree and 4 of 5 against the tuned tree, so the ordering does not depend on the test split.
- The trees also generalise worse: train–test RMSE gap 0.046–0.610 versus 0.013 for linear, while the tree needs 21–4,245 leaves against 9 coefficients.

5.2 Is there non-linear structure for the Decision Tree to exploit?

Three probes: (a) a nested-model **F-test** of linear vs. degree-2 polynomial (training data), together with the biggest quadratic/interaction terms; (b) the **single-feature vs. all-feature** behaviour of the linear and tree models (Stage 2 ablation) — a tree can model a non-linear curve of *one* variable, but is inefficient at adding up *several* variables linearly; (c) a **hybrid** that fits the tree to the *residuals of the linear model* — if the tree finds structure there, the linear model is leaving non-linear signal on the table.

```
# (a) F-test: does the degree-2 expansion explain significantly more than the linear model? (training data)
n_ = len(y_train); sse1 = np.sum((y_train.values - PRED[N_LIN_NE] ["train"]) ** 2); sse2 = np.sum((y_train.values - PRED[N_POLY_NE] ["train"]) ** 2); p1,
p2 = len(FEATURES) + 1, nterms(DEG_STAR)
F = ((sse1 - sse2) / (p2 - p1)) / (sse2 / (n_ - p2)); pval = stats.f.sf(F, p2 - p1, n_ - p2)
print(f"Nested F-test linear (p={p1}) vs degree-2 (p={p2}): F = {F:.1f}, p = {pval:.1e} | R² train {RESULTS[N_LIN_NE] ['train_r2']:.4f} ->
{RESULTS[N_POLY_NE] ['train_r2']:.4f}, "
      f"test {RESULTS[N_LIN_NE] ['test_r2']:.4f} -> {RESULTS[N_POLY_NE] ['test_r2']:.4f}")
print("(p-value is optimistic because hourly residuals are autocorrelated; the held-out CV/test gains in Stages 2–3 are the more reliable evidence)")
pm_ = FITTED[N_POLY_NE]; tn = pm_.poly_.feature_names(FEATURES)
terms = pd.DataFrame({"term": tn, "std. coefficient": pm_.coef_}); terms["|coef|"] = terms["std. coefficient"].abs()
terms["order"] = terms["term"].map(lambda s: s.count("*") + 1); TOTTERMS = terms.sort_values("|coef|", ascending=False).head(10).set_index("term")
display(TOTTERMS.style.format(precision=3).set_caption("Ten largest terms of the degree-2 model (collinear terms trade off against each other: do not
read individually)"))
q = terms[terms["order"] == 2]; print(f"Share of total |coef| carried by degree-2 terms: {q["|coef|"].sum() / terms["|coef|"].sum():.1%}")

# (b) one feature vs. all features
sf = ABL_FEAT[ABL_FEAT["feature set"].isin(["Best single sensor (PT08.S2)", "All 8 features"])] .pivot(index="feature set", columns="model",
values="cv_rmse")
sf = sf.loc[["Best single sensor (PT08.S2)", "All 8 features"]]; sf.loc["relative gain from adding the other 7 features"] = (sf.iloc[0] - sf.iloc[1]) /
sf.iloc[0]
display(sf.style.format(precision=4).set_caption("CV RMSE with PT08.S2 only vs. all features (last row: relative improvement)"))
```

Nested F-test linear (p=9) vs degree-2 (p=45): $F = 58.7$, $p = 0.0e+00$ | R^2 train 0.8889 → 0.9184, test 0.8799 → 0.9013
(p-value is optimistic because hourly residuals are autocorrelated; the held-out CV/test gains in Stages 2–3 are the more reliable evidence)

Ten largest terms of the degree-2 model (collinear terms trade off against each other: do not read individually)

term	std. coefficient	coef	order
PT08.S2(NMHC)	1.378	1.378	1
PT08.S2(NMHC)*PT08.S4(NO2)	1.214	1.214	2
T	0.912	0.912	1
AH	-0.761	0.761	1
PT08.S1(CO)*PT08.S2(NMHC)	0.755	0.755	2
RH	0.712	0.712	1
PT08.S2(NMHC)*PT08.S2(NMHC)	-0.686	0.686	2
PT08.S2(NMHC)*AH	-0.656	0.656	2
PT08.S4(NO2)*PT08.S4(NO2)	-0.592	0.592	2
PT08.S4(NO2)*AH	0.452	0.452	2

Share of total |coef| carried by degree-2 terms: 65.1%

CV RMSE with PT08.S2 only vs. all features (last row: relative improvement)

feature set	model	Linear	Poly(d=2)	Tree (tuned)
Best single sensor (PT08.S2)		0.5750	0.5196	0.5341
All 8 features		0.4882	0.4530	0.5134
relative gain from adding the other 7 features		0.1511	0.1282	0.0389

Observations — is there non-linear structure?

- **Yes, but it is modest and smooth.** The nested F-test strongly rejects the linear model in favour of degree 2 ($F = 58.7$; the p-value underflows to 0, and is over-optimistic because of autocorrelation). The held-out evidence is more convincing: test R^2 $0.880 \rightarrow 0.901$ and CV RMSE $0.488 \rightarrow 0.453$. Degree-2 terms carry 65 % of the total absolute coefficient mass (largest: $S2 \times S4$, $S1 \times S2$, $S2^2$, $S2 \times AH$), but these terms are collinear, so individual coefficients should not be interpreted.
- **A tree is good at one-variable non-linearity and weak at adding variables up.** With PT08.S2 alone, the tuned tree (CV RMSE 0.534) is better than the straight line (0.575) — it captures the curvature of that single relation. But when the other seven features are added, the linear model improves by 15.1 % ($0.575 \rightarrow 0.488$) and the polynomial by 12.8 %, whereas the tree improves by only **3.9 %** ($0.534 \rightarrow 0.513$): a staircase function needs very many splits to represent a weighted sum of several correlated sensors, and each split uses only one feature.
- **The hybrid probe confirms that the line leaves non-linear signal:** a shallow tree fitted on the linear residuals lowers CV RMSE from 0.488 to 0.471 and test RMSE from 0.493 to 0.470, so *combining* a linear backbone with tree corrections beats either alone — yet it still does not reach the smooth degree-2 polynomial (test 0.447). The remaining non-linearity is thus better described by smooth quadratic terms than by axis-aligned steps.

```
# (c) hybrid probe: linear model + shallow tree fitted on the linear model's residuals (extension, NOT one of the assignment models)
class LinearPlusTree(RegressorMixin, BaseEstimator):
    def __init__(self, max_depth=3, min_samples_leaf=50): self.max_depth, self.min_samples_leaf = max_depth, min_samples_leaf
    def fit(self, X, y):
        self.lin_ = LinearRegressionScratch().fit(X, y); self.tree_ = DecisionTreeRegressor(max_depth=self.max_depth,
min_samples_leaf=self.min_samples_leaf, random_state=SEED).fit(X, np.asarray(y) - self.lin_.predict(X)); return self
    def predict(self, X): return self.lin_.predict(X) + self.tree_.predict(X)

rows = [{"hybrid depth": d, **cv_eval(LinearPlusTree(max_depth=d))} for d in (1, 2, 3, 4, 5)]
HYB = pd.DataFrame(rows).set_index("hybrid depth")[["cv_rmse", "cv_rmse_sd", "cv_r2", "cv_train_rmse"]]; d_h = int(HYB["cv_rmse"].idxmin())
display(HYB.style.format(precision=4).set_caption(f"Linear + tree-on-residuals, group-CV (linear alone: {RESULTS[N_LIN_NE]['cv_rmse']:.4f}; polynomial
d=2: {RESULTS[N_POLY_NE]['cv_rmse']:.4f})))
hyb = LinearPlusTree(max_depth=d_h).fit(X_train_s, y_train); ph = hyb.predict(X_test_s)
HYBRID = {"depth": d_h, "test_rmse": rmse(y_test, ph), "test_r2": r2(y_test, ph), "cv_rmse": HYB.loc[d_h, "cv_rmse"]}
print(f"Hybrid (depth {d_h}): test RMSE {HYBRID['test_rmse']:.4f}, R² {HYBRID['test_r2']:.4f} vs Linear {RESULTS[N_LIN_NE]['test_rmse']:.4f} / Tree
tuned {RESULTS[N_TREE_T]['test_rmse']:.4f} / Poly {RESULTS[N_POLY_NE]['test_rmse']:.4f}")
```

Linear + tree-on-residuals, group-CV (linear alone: 0.4882;

polynomial d=2: 0.4530)

	cv_rmse	cv_rmse_sd	cv_r2	cv_train_rmse
hybrid depth				
1	0.4852	0.0502	0.8841	0.4642
2	0.4757	0.0526	0.8884	0.4528
3	0.4718	0.0516	0.8903	0.4472
4	0.4717	0.0528	0.8903	0.4431
5	0.4707	0.0523	0.8908	0.4384

Hybrid (depth 5): test RMSE 0.4703, R^2 0.8908 vs Linear 0.4933 / Tree tuned 0.5157 / Poly 0.4470


5.3 What can each model family represent? — a synthetic demonstration

Synthetic data (not the air-quality data): one input $x \sim U(0, 1)$, noise sd 0.15, four known truths. Linear regression, the degree-2 polynomial and a depth-4 tree are fitted on 70 % and scored on the remaining 30 %. The second figure trains on $x \in [0, 1]$ and tests on $x \in [1, 1.5]$ to show **extrapolation**. This isolates the properties that explain the real-data results.

```
rs = np.random.default_rng(0)
truths = {"linear: 2x": lambda x: 2 * x, "quadratic: 4(x-½)²": lambda x: 4 * (x - .5) ** 2, "smooth wave: sin(2πx)": lambda x: np.sin(2 * np.pi * x),
"step: 1[x>½]": lambda x: (x > .5).astype(float)}
mk = {"Linear": lambda: LinearRegressionScratch(), "Poly d=2": lambda: PolyRegressionScratch(2), "Tree depth 4": lambda:
DecisionTreeRegressor(max_depth=4, min_samples_leaf=10, random_state=SEED)}
mc = {"Linear": BLUE, "Poly d=2": ORANGE, "Tree depth 4": GREEN}
fig, axes = plt.subplots(1, 4, figsize=(20, 4.2)); SYN = {}
for a, (tn_, f_) in zip(axes, truths.items()):
    x = rs.uniform(0, 1, 500); y = f_(x) + rs.normal(0, .15, 500); xtr, xte, ytr, yte = x[:350], x[350:], y[:350], y[350:]
    a.scatter(xte, yte, s=6, color=GREY, alpha=.5); gx = np.linspace(0, 1, 300); a.plot(gx, f_(gx), "k--", lw=1.2, label="truth")
    for m_, mk_ in mk.items():
        mod = mk_.fit(xtr[:, None], ytr); SYN[(tn_, m_)] = rmse(yte, mod.predict(xte[:, None])); a.plot(gx, mod.predict(gx[:, None]), color=mc[m_],
lw=2, label=f"{m_} (RMSE {SYN[(tn_, m_)]:.2f})")
    a.set_title(tn_); a.legend(fontsize=7)
plt.suptitle("Synthetic truths, test RMSE in legend (noise floor = 0.15)", y=1.02, fontweight="bold"); plt.tight_layout(); plt.show()

fig, axes = plt.subplots(1, 2, figsize=(12, 4)); EXT = {}
for a, (tn_, f_) in zip(axes, [{"linear: 2x", truths["linear: 2x"]}, {"quadratic: 4(x-½)²", truths["quadratic: 4(x-½)²"]}]):
    x = rs.uniform(0, 1, 400); y = f_(x) + rs.normal(0, .15, 400); gx = np.linspace(0, 1.5, 300); a.scatter(x, y, s=5, color=GREY, alpha=.5);
a.plot(gx, f_(gx), "k--", lw=1.2, label="truth"); a.axvspan(1, 1.5, color=RED, alpha=.07)
xe = rs.uniform(1, 1.5, 200); ye = f_(xe) + rs.normal(0, .15, 200)
for m_, mk_ in mk.items():
    mod = mk_.fit(x[:, None], y); EXT[(tn_, m_)] = rmse(ye, mod.predict(xe[:, None])); a.plot(gx, mod.predict(gx[:, None]), color=mc[m_], lw=2,
label=f"{m_} (extrap. RMSE {EXT[(tn_, m_)]:.2f})")
    a.set_title(f"Extrapolation beyond the training range (shaded): {tn_}"); a.legend(fontsize=7)
plt.tight_layout(); plt.show()
display(pd.concat({"in-range test RMSE": pd.Series(SYN).unstack(), "extrapolation RMSE": pd.Series(EXT).unstack()}, axis=1).style.format(precision=3))
```

Synthetic truths, test RMSE in legend (noise floor = 0.15)



	in-range test RMSE			extrapolation RMSE		
	Linear	Poly d=2	Tree depth 4	Linear	Poly d=2	Tree depth 4
linear: $2x$	0.139	0.138	0.149	0.160	0.159	0.668
quadratic: $4(x-\frac{1}{2})^2$	0.354	0.152	0.167	2.208	0.156	1.707
smooth wave: $\sin(2\pi x)$	0.470	0.469	0.193	nan	nan	nan
step: $1[x > \frac{1}{2}]$	0.290	0.289	0.158	nan	nan	nan

Observations (synthetic data, so the mechanisms can be seen in isolation).

- **Linear truth:** all three models recover it ($\text{RMSE} \approx 0.14\text{--}0.15 \approx$ the noise floor); the tree draws it as a staircase.
- **Quadratic truth:** the line fails completely (0.35), the degree-2 polynomial is exact (0.15) and the tree approximates it with steps (0.17).
- **Sine wave / step:** the polynomial (degree 2) cannot represent either and equals the straight line (0.47 and 0.29), whereas the **tree is best (0.19 and 0.16)** — trees excel at abrupt or irregular non-linear shapes.
- **Extrapolation:** outside the training range the tree **freezes at its last leaf value** (RMSE 0.67 for a linear truth, 1.71 for a quadratic truth), the straight line keeps going (exact for a linear truth, 2.21 error for a quadratic one) and only the polynomial that matches the truth extrapolates well.
- Linking back to the real data: its CO ↔ sensor relation looks like the *smooth, nearly linear/quadratic* case, not the abrupt-step case — consistent with linear/polynomial ≥ tree.


5.4 Where does each model win? — by sensor level, and under unseen conditions

(a) Test RMSE per decile of the dominant sensor PT08.S2 : does the tree beat the line anywhere? (b) **Extrapolation on real data:** using the chronological split (Stage 2), how large is the fraction of test hours with a feature outside the training range, and how do the models behave there?

```
dec = pd.qcut(X_test["PT08.S2(NMHC)"], 10, labels=[f"D{i}" for i in range(1, 11)])
cmp_models = {"Linear": N_LIN_NE, f"Poly d={DEG_STAR}": N_POLY_NE, "Tree depth 5": tree_names[5], "Tree tuned": N_TREE_T}
DECILE = pd.DataFrame({l: pd.Series(yt_all - PRED[nm]["test"], index=y_test.index).groupby(dec, observed=True).apply(lambda s: np.sqrt((s ** 2).mean())) for l, nm in cmp_models.items()})
DECILE["mean CO"] = y_test.groupby(dec, observed=True).mean(); DECILE["winner"] = DECILE[list(cmp_models)].idxmin(axis=1)
display(DECILE.style.format(precision=3).background_gradient(subset=list(cmp_models), cmap="RdYlGn_r", axis=1))
lt = DECILE[["Linear", "Tree tuned"]]; print("Deciles in which Tree tuned beats Linear:", list(lt.index[lt["Tree tuned"] < lt["Linear"]]), "| Linear beats tree in", int((lt["Linear"] < lt["Tree tuned"]).sum()), "of 10")
fig, ax = plt.subplots(figsize=(12, 4)); w = .2
for i, (l, c) in enumerate(zip(cmp_models, [BLUE, ORANGE, GREY, GREEN])): ax.bar(np.arange(10) + (i - 1.5) * w, DECILE[l], w, color=c, label=l)
ax.set_xticks(range(10)); ax.set_xticklabels(DECILE.index); ax.set_xlabel("decile of PT08.S2 (D1 = lowest signal)"); ax.set_ylabel("test RMSE");
ax.set_title("RMSE by level of the dominant sensor"); ax.legend(ncol=4); plt.tight_layout(); plt.show()
```

	Linear	Poly d=2	Tree depth 5	Tree tuned	mean CO	winner
PT08.S2(NMHC)						
D1	0.300	0.148	0.218	0.218	0.608	Poly d=2
D2	0.411	0.361	0.428	0.428	0.879	Poly d=2
D3	0.399	0.371	0.416	0.416	1.074	Poly d=2
D4	0.385	0.376	0.401	0.401	1.364	Poly d=2
D5	0.382	0.361	0.388	0.388	1.639	Poly d=2
D6	0.350	0.320	0.391	0.391	1.966	Poly d=2
D7	0.525	0.508	0.548	0.548	2.358	Poly d=2
D8	0.539	0.507	0.589	0.589	2.878	Poly d=2
D9	0.628	0.609	0.642	0.642	3.482	Poly d=2
D10	0.800	0.672	0.842	0.857	4.976	Poly d=2

Deciles in which Tree tuned beats Linear: ['D1'] | Linear beats tree in 9 of 10



Observations. The degree-2 polynomial has the lowest RMSE in **all ten deciles** of the dominant sensor. Between the two models the assignment asks about, Linear beats the tuned tree in **9 of 10** deciles; the tree wins only in the lowest decile (D1, 0.218 vs. 0.300), where the relation flattens. The linear model's disadvantage is largest at the extremes of the sensor range (D1 and D10, where the straight line misses the curvature), and the tree's disadvantage is largest in the top decile (0.857 vs. 0.800 for linear and 0.672 for the polynomial). The depth-5 and the tuned tree are identical in nine deciles because both split almost only on PT08.S2 .

```
oot_pipes = {l: Pipeline([("scale", StandardScaler()), ("m", mk_())]).fit(Xo_tr, yo_tr) for l, mk_ in {"Linear": make_linear, f"Poly d={DEG_STAR}":
make_poly, "Tree depth 5": lambda: DecisionTreeRegressor(max_depth=5, random_state=SEED), "Tree tuned": make_tree}.items()}
lo_, hi_ = Xo_tr.min(), Xo_tr.max(); out_f = (Xo_te < lo_) | (Xo_te > hi_); out_any = out_f.any(axis=1)
print("Share of out-of-time test hours outside the training range, per feature:", {k: f"{v:.1%}" for k, v in out_f.mean().items() if v > 0}, "| any
feature:", f"{out_any.mean():.1%}")
EXTRAP = pd.DataFrame({l: {"RMSE inside training range": rmse(yo_te[~out_any], p.predict(Xo_te[~out_any])), "RMSE outside range": rmse(yo_te[out_any],
p.predict(Xo_te[out_any])) if out_any.sum() else np.nan,
                        "n inside": int((~out_any).sum()), "n outside": int(out_any.sum())} for l, p in oot_pipes.items()}).T
display(EXTRAP.style.format({"RMSE inside training range": "{:.3f}", "RMSE outside range": "{:.3f}"})
```

Share of out-of-time test hours outside the training range, per feature: {'PT08.S2(NMHC)': '0.1%', 'PT08.S4(N02)': '1.7%', 'PT08.S5(03)': '0.4%', 'T': '2.6%', 'AH': '0.3%'} | any feature: 3.3%

	RMSE inside training range	RMSE outside range	n inside	n outside
Linear	0.489	0.281	1420.000000	49.000000
Poly d=2	0.451	0.261	1420.000000	49.000000
Tree depth 5	0.534	0.285	1420.000000	49.000000
Tree tuned	0.519	0.285	1420.000000	49.000000

Observations. Only **3.3 %** of the out-of-time test hours have any feature outside the range seen in training (mainly T 2.6 % and PT08.S4 1.7 %), and those hours are easy (RMSE 0.26–0.29 versus 0.45–0.53 inside the range — low CO, low variance). This dataset therefore provides **no empirical evidence of extrapolation failure**; the synthetic demonstration above shows the mechanism, and it would matter for deployment if the sensors drifted further (PT08.S4 already drifts, Stage 1): a tree would saturate while a linear model would follow the drift.

5.5 Ablation ledger — every ablation in one table

```
NOISE = 0.03      # differences below this are within the fold-to-fold noise seen in Stages 2-3

def verdict(delta):
    return "negligible (within noise)" if abs(delta) < NOISE else ("worse" if delta > 0 else "better")

L = []

def add(group, what, model, delta, unit="ΔCV RMSE", note=""): L.append({"ablation": group, "change": what, "model": model, unit: delta, "verdict": verdict(delta), "note": note})

for m_ in ["Linear", f"Poly(d={DEG_STAR})", "Tree (tuned)"]:
    add("Features", "sensors only (drop weather)", m_, tbl.loc["Sensors only (5)", m_] - base[m_])
    add("Features", "weather only (drop sensors)", m_, tbl.loc["Weather only (T, RH, AH)", m_] - base[m_])
    add("Features", "drop PT08.S2 only", m_, tbl.loc["Drop PT08.S2(NMHC)", m_] - base[m_])
    add("Features", "PT08.S2 as the only feature", m_, tbl.loc["Best single sensor (PT08.S2)", m_] - base[m_])

add("Scaling", "no standardisation (batch GD, 1000 epochs)", "Linear", ABL_SCALE.loc["raw features", "gap to OLS optimum"], "Δ train RMSE", "optimisation, not model quality")

for m_ in ABL_LOG.index: add("Target", "log1p-transformed target", m_, ABL_LOG.loc[m_, "CV RMSE (log1p target)" ] - ABL_LOG.loc[m_, "CV RMSE (raw target)"])

for c_ in ["absolute_error", "poisson"]: add("Tree criterion", c_ + " instead of squared_error", "Tree (tuned)", ABL_CRIT.loc[c_, "cv_rmse"] - ABL_CRIT.loc["squared_error", "cv_rmse"])

add("Regularisation", f"ridge (best α≈{RIDGE_CV['cv_rmse'].idxmin():.0f}) instead of OLS", "Linear", RIDGE_CV["cv_rmse"].min() - RIDGE_CV.iloc[0] ["cv_rmse"])

add("Regularisation", f"ridge α={ALPHA_R:g} on degree {DEG_R}", f"Poly(d={DEG_R})", RESULTS[N_PR_SC] ["cv_rmse"] - RESULTS[N_POLY_NE] ["cv_rmse"])

for d_ in (1, 3, 4, 5): add("Polynomial degree", f"degree {d_} instead of {DEG_STAR}", "Poly", unreg.loc[d_, "cv_rmse"] - unreg.loc[DEG_STAR, "cv_rmse"])

for d_ in ("3", "10", "None"): add("Tree depth", f"max_depth={d_} instead of 5", "Tree", TREE_SWEEP[TREE_SWEEP["max_depth"].astype(str) == d_] ["cv_rmse"].iloc[0] - RESULTS[tree_names[5]] ["cv_rmse"])

add("Tree regularisation", "cost-complexity pruning vs. tuned depth/leaf", "Tree", RESULTS[N_TREE_P] ["cv_rmse"] - RESULTS[N_TREE_T] ["cv_rmse"])

add("Tree regularisation", "min_samples_leaf=50 added to depth 5", "Tree", RESULTS[N_TREE_T] ["cv_rmse"] - RESULTS[tree_names[5]] ["cv_rmse"])

add("Hybrid extension", f"linear + depth-{HYBRID['depth']} tree on residuals", "Linear", HYBRID["cv_rmse"] - RESULTS[N_LIN_NE] ["cv_rmse"])

for m_, r_ in OOT.iterrows():
    if m_ in ("Linear", f"Poly(d={DEG_STAR})", "Tree (tuned)", "Tree max_depth=5"): add("Evaluation protocol", "chronological instead of shuffled split", m_, r_["degradation (oot - random)"], "Δ test RMSE")

LEDGER = pd.DataFrame(L); LEDGER["Δ"] = LEDGER[["ΔCV RMSE", "Δ train RMSE", "Δ test RMSE"]].bfill(axis=1).iloc[:, 0]

LEDGER = LEDGER[["ablation", "change", "model", "Δ", "verdict", "note"]].fillna({"note": ""})

display(LEDGER.style.format({"Δ": "{:+.4f}" }).map(lambda v: "color:#b00000;font-weight:600" if v == "worse" else ("color:#1a7f37;font-weight:600" if v == "better" else "" ), subset=["verdict"]).hide(axis="index"))

print("Positive Δ = the ablated variant is worse than the reference. Noise threshold used for the verdict:", NOISE)
```

ablation	change	model	Δ	verdict	note
Features	sensors only (drop weather)	Linear	+0.0143	negligible (within noise)	
Features	weather only (drop sensors)	Linear	+0.9541	worse	
Features	drop PT08.S2 only	Linear	+0.1191	worse	
Features	PT08.S2 as the only feature	Linear	+0.0869	worse	
Features	sensors only (drop weather)	Poly(d=2)	+0.0157	negligible (within noise)	
Features	weather only (drop sensors)	Poly(d=2)	+1.0027	worse	
Features	drop PT08.S2 only	Poly(d=2)	+0.0710	worse	
Features	PT08.S2 as the only feature	Poly(d=2)	+0.0666	worse	
Features	sensors only (drop weather)	Tree (tuned)	+0.0143	negligible (within noise)	
Features	weather only (drop sensors)	Tree (tuned)	+0.9575	worse	
Features	drop PT08.S2 only	Tree (tuned)	+0.1505	worse	
Features	PT08.S2 as the only feature	Tree (tuned)	+0.0208	negligible (within noise)	
Scaling	no standardisation (batch GD, 1000 epochs)	Linear	+0.1410	worse	optimisation, not model quality
Target	log1p-transformed target	Linear	+0.0355	worse	
Target	log1p-transformed target	Poly(d=2)	-0.0043	negligible (within noise)	
Target	log1p-transformed target	Tree (tuned)	-0.0012	negligible (within noise)	
Tree criterion	absolute_error instead of squared_error	Tree (tuned)	-0.0115	negligible (within noise)	
Tree criterion	poisson instead of squared_error	Tree (tuned)	-0.0073	negligible (within noise)	
Regularisation	ridge (best $\alpha=22$) instead of OLS	Linear	-0.0005	negligible (within noise)	
Regularisation	ridge $\alpha=1$ on degree 2	Poly(d=2)	-0.0015	negligible (within noise)	
Polynomial degree	degree 1 instead of 2	Poly	+0.0351	worse	
Polynomial degree	degree 3 instead of 2	Poly	+0.0469	worse	
Polynomial degree	degree 4 instead of 2	Poly	+0.7746	worse	
Polynomial degree	degree 5 instead of 2	Poly	+8.1184	worse	
Tree depth	max_depth=3 instead of 5	Tree	+0.0681	worse	
Tree depth	max_depth=10 instead of 5	Tree	+0.0889	worse	
Tree depth	max_depth=None instead of 5	Tree	+0.1501	worse	
Tree regularisation	cost-complexity pruning vs. tuned depth/leaf	Tree	+0.0078	negligible (within noise)	
Tree regularisation	min_samples_leaf=50 added to depth 5	Tree	+0.0042	negligible (within noise)	
Hybrid extension	linear + depth-5 tree on residuals	Linear	-0.0175	negligible (within noise)	
Evaluation protocol	chronological instead of shuffled split	Linear	-0.0097	negligible (within noise)	
Evaluation protocol	chronological instead of shuffled split	Poly(d=2)	-0.0011	negligible (within noise)	
Evaluation protocol	chronological instead of shuffled split	Tree max_depth=5	+0.0145	negligible (within noise)	
Evaluation protocol	chronological instead of shuffled split	Tree (tuned)	-0.0031	negligible (within noise)	

Positive Δ = the ablated variant is worse than the reference. Noise threshold used for the verdict: 0.03

Observations — what the ablations tell us.

- **Features:** sensors alone are essentially as good as everything (+0.014 to +0.016, within noise); weather alone is as bad as the mean predictor (+0.95 to +1.00); PT08.S2 is the one indispensable feature (dropping it costs +0.07 to +0.15) and the only one for which a single-feature model is already competitive.
- **Model complexity:** the polynomial degree must be exactly 2 (degree 1: +0.035, degree 3: +0.047, degree 4: +0.77, degree 5: +8.1); for trees depth 5 is best (3: +0.068, 10: +0.089, None : +0.150).
- **Preprocessing:** standardisation matters for gradient-based optimisation (+0.14 train RMSE without it) and log-transforming the target hurts the linear model (+0.036) while being neutral for the others.
- **Negligible (within noise):** ridge for the linear model or the polynomial, the tree split criterion, cost-complexity pruning vs. the tuned tree, the hybrid correction (−0.018, in the right direction), and the chronological split ($|\Delta| \leq 0.015$). The conclusions are therefore robust to these choices.

5.6 Scorecard — Linear vs. Decision Tree vs. Polynomial

```
SC_M = {"Linear": N_LIN_NE, "Tree depth 5": tree_names[5], "Tree tuned": N_TREE_T, f"Poly d={DEG_STAR}": N_POLY_NE}
reg = REGIME.set_index(["model", "regime"]); vh = bands_labels[-1]; hi_r = bands_labels[-2]
fold_col = {"Linear": "Linear", "Tree depth 5": "Tree depth 5", "Tree tuned": "Tree (tuned)", f"Poly d={DEG_STAR}": f"Poly(d={DEG_STAR})"}
oot_row = {"Linear": "Linear", "Tree depth 5": "Tree max_depth=5", "Tree tuned": "Tree (tuned)", f"Poly d={DEG_STAR}": f"Poly(d={DEG_STAR})"}
SCORE = pd.DataFrame({l: {"test RMSE": RESULTS[nm]["test_rmse"], "test R²": RESULTS[nm]["test_r2"], "CV RMSE": RESULTS[nm]["cv_rmse"], "CV fold-to-fold
sd": FOLD_RMSE.loc["sd", fold_col[l]],
                        "out-of-time test RMSE": OOT.loc[oot_row[l], "oot_test_rmse"], "RMSE gap (test - train)": RESULTS[nm]["rmse_gap (test-
train)"],
                        f"RMSE, {hi_r}": reg.loc[(nm, hi_r), "rmse"], f"RMSE, {vh}": reg.loc[(nm, vh), "rmse"], f"bias, {vh}": reg.loc[(nm, vh),
"bias (pred-true)"],
                        "parameters / leaves": RESULTS[nm]["complexity"], "fit time (s)": RESULTS[nm]["fit_time_s"]} for l, nm in SC_M.items()})
lower_better = [r for r in SCORE.index if r not in ("test R²", f"bias, {vh}")]
SCORE["best"] = [SCORE.loc[r].idxmin() if r in lower_better else (SCORE.loc[r].idxmax() if r == "test R²" else SCORE.loc[r].abs().idxmin()) for r in
SCORE.index]
display(SCORE.style.format({c: "{:.4f}" for c in SCORE.columns if c != "best"}))
static = pd.DataFrame({"Linear": ["yes (straight line)", "no", "9 signed coefficients with CIs", "needs scaling for GD", "yes"],
                       "Decision Tree": ["no - constant beyond the training range", "yes", "readable rules / leaves, but unstable across data subsets",
"none needed", "no (staircase)"],
                       f"Poly d={DEG_STAR}": ["yes, but quadratic growth", "mild (degree 2)", "45 correlated terms: hard to read", "needs scaling",
"yes"]},
                      index=["extrapolates beyond the training range", "captures non-linear / interaction effects", "interpretability", "preprocessing
sensitivity", "smooth predictions"])
display(static)
```

	Linear	Tree depth 5	Tree tuned	Poly d=2	best
test RMSE	0.4933	0.5131	0.5157	0.4470	Poly d=2
test R²	0.8799	0.8700	0.8687	0.9013	Poly d=2
CV RMSE	0.4882	0.5092	0.5134	0.4530	Poly d=2
CV fold-to-fold sd	0.0610	0.0572	0.0534	0.0538	Tree tuned
out-of-time test RMSE	0.4836	0.5276	0.5126	0.4460	Poly d=2
RMSE gap (test - train)	0.0134	0.0728	0.0459	0.0359	Linear
RMSE, High (3-5)	0.5193	0.6460	0.6529	0.5594	Linear
RMSE, Very high (>=5)	1.1729	1.1518	1.1608	0.9462	Poly d=2
bias, Very high (>=5)	-0.9270	-0.6597	-0.5561	-0.5550	Poly d=2
parameters / leaves	9.0000	32.0000	29.0000	45.0000	Linear
fit time (s)	0.0005	0.0113	0.0103	0.0050	Linear

	Linear	Decision Tree	Poly d=2
extrapolates beyond the training range	yes (straight line)	no — constant beyond the training range	yes, but quadratic growth
captures non-linear / interaction effects	no	yes	mild (degree 2)
interpretability	9 signed coefficients with CIs	readable rules / leaves, but unstable across d...	45 correlated terms: hard to read
preprocessing sensitivity	needs scaling for GD	none needed	needs scaling
smooth predictions	yes	no (staircase)	yes

Observations.

- **Accuracy:** Polynomial d=2 is best on test RMSE/R², CV RMSE and out-of-time RMSE; Linear is second; the trees are third. In the High regime Linear (0.519) is far better than the trees (0.646–0.653); in the Very-high regime the models are close (tree 1.15–1.16, linear 1.17, polynomial 0.95), with the trees slightly less biased than linear (–0.56 to –0.66 vs. –0.93).
- **Simplicity and stability:** Linear has the fewest parameters, the smallest generalisation gap and the fastest fit; the differences in fold-to-fold sd (0.053–0.061) are not meaningful.
- **Qualitative:** trees need no scaling and give readable rules, but cannot extrapolate and give staircase predictions; the polynomial is accurate but its 45 collinear terms are hard to interpret and it is fragile on unrealistic inputs (Stage 4).

 **Stage 5 complete** — head-to-head comparison with uncertainty, evidence on linear vs. non-linear structure, a synthetic illustration, regime-wise wins, extrapolation behaviour, the ablation ledger and the scorecard.

Stage 6 — Discussion

Short written answers to the two questions in the brief, followed by limitations, suggested next steps and a checklist of the deliverables. **All numbers in the answers below are computed from the results above** (they are inserted programmatically, so the text cannot drift out of sync with the analysis).

Q1. Which features were most important for predicting CO(GT)?

```
from IPython.display import Markdown

shr = IMP_N.mean(axis=1).sort_values(ascending=False); f1_, f2_, f3_ = shr.index[:3]
co_ = COEF_BOOT.loc["PT08.S2(NMHC)"]
dS2 = {m_: tbl.loc["Drop PT08.S2(NMHC)", m_] - base[m_] for m_ in ["Linear", f"Poly(d={DEG_STAR})", "Tree (tuned)"]}
sens_only = tbl.loc["Sensors only (5)", "Linear"] - base["Linear"]; weather_only = tbl.loc["Weather only (T, RH, AH)", "Linear"]

Markdown(f"""
**Answer.** The {f1_} sensor is by far the most important feature; the other sensors and the weather variables add only small corrections.

- **Consensus across methods.** {f1_} ranks first in all {IMP.shape[1]} importance methods (impurity, permutation, standardised coefficients, drop-one retraining) with a mean share of {shr.iloc[0]:.0%} of the total importance. The second and third are {f2_} ({shr.iloc[1]:.1%}) and {f3_} ({shr.iloc[2]:.1%}).

- **Linear view.** Its standardised coefficient is {co_['coef']:.2f} (95 % bootstrap CI {co_['ci_low']:.2f} to {co_['ci_high']:.2f}), about {co_['coef'] / COEF_BOOT['coef'].drop('PT08.S2(NMHC)').abs().max():.1f} times the next largest effect: one standard deviation more PT08.S2 signal goes with about {co_['coef']:.1f} mg/m3 more CO.

- **Tree view.** The depth-3 tree splits only on PT08.S2 (all {int(used.sum())} splits); the tuned tree assigns it {IMP_TREE.loc['PT08.S2(NMHC)', N_TREE_T]:.0%} of its impurity importance.

- **Drop-one test (retraining without the feature).** Removing PT08.S2 raises the CV RMSE by {dS2['Linear']:.3f} (linear), {dS2[f'Poly(d={DEG_STAR})']:.3f} (polynomial) and {dS2['Tree (tuned)']:.3f} (tree), whereas removing any other single feature changes it by no more than 0.022 for the linear and tree models.

- **Sensors vs. weather.** Using only the five sensors costs just {sens_only:+.3f} RMSE relative to all eight features, while weather alone is no better than predicting the mean (CV RMSE {weather_only:.2f} vs. {RESULTS[N_MEAN]['cv_rmse']:.2f}). The weather variables help only slightly, mainly temperature, which ranks second or third (mean share {shr.get('T', 0):.1%}; linear coefficient {COEF_BOOT.loc['T', 'coef']:.2f}).

- **Collinearity caveat.** PT08.S1(CO) and PT08.S5(O3) are strongly correlated with PT08.S2 ( $r \approx 0.85\text{--}0.89$ ). They matter for the linear model (PT08.S1 is its second most important feature) but not for the tree, which needs only one representative of the sensor cluster; importance of redundant sensors is therefore shared arbitrarily and should be read as "the sensor cluster, led by PT08.S2".
""")
```

Answer. The **PT08.S2(NMHC)** sensor is by far the most important feature; the other sensors and the weather variables add only small corrections.

- **Consensus across methods.** **PT08.S2(NMHC)** ranks first in all 8 importance methods (impurity, permutation, standardised coefficients, drop-one retraining) with a mean share of **77%** of the total importance. The second and third are **T** (6.8%) and **PT08.S1(CO)** (6.1%).
- **Linear view.** Its standardised coefficient is **1.37** (95 % bootstrap CI **1.24** to **1.48**), about **3.5** times the next largest effect: one standard deviation more **PT08.S2** signal goes with about **1.4** mg/m³ more CO.
- **Tree view.** The depth-3 tree splits **only** on **PT08.S2** (all **7** splits); the tuned tree assigns it **97%** of its impurity importance.
- **Drop-one test (retraining without the feature).** Removing **PT08.S2** raises the CV RMSE by **0.119** (linear), **0.071** (polynomial) and **0.151** (tree), whereas removing any other single feature changes it by no more than **0.022** for the linear and tree models.
- **Sensors vs. weather.** Using only the five sensors costs just **+0.014** RMSE relative to all eight features, while weather alone is no better than predicting the mean (CV RMSE **1.44** vs. **1.44**). The weather variables help only slightly, mainly **temperature**, which ranks second or third (mean share **6.8%**; linear coefficient **-0.24**).
- **Collinearity caveat.** **PT08.S1(CO)** and **PT08.S5(O3)** are strongly correlated with **PT08.S2** ($r \approx 0.85\text{--}0.89$). They matter for the linear model (**PT08.S1** is its second most important feature) but not for the tree, which needs only one representative of the sensor cluster; importance of redundant sensors is therefore shared arbitrarily and should be read as "the sensor cluster, led by **PT08.S2**".

Q2. Which performed better overall: Decision Tree or Linear Regression? Why?

```
t5, tt_ = RESULTS[tree_names[5]], RESULTS[N_TREE_T]; ln_ = RESULTS[N_LIN_NE]; pl_ = RESULTS[N_POLY_NE]
d5 = H2H.loc["max_depth=5"]; best_tree_gap = d5["Δ test RMSE vs Linear"]; wins_all = int((FOLD_RMSE.iloc[:5]["Linear"] < FOLD_RMSE.iloc[:5]["Tree depth
5"]).sum())
Markdown(f"""
**Answer.** **Linear Regression performed better overall**, although the margin over the best-sized tree is modest. The degree-2 polynomial regression
was the best model of all.

| | Linear | Best Decision Tree (depth 5) | Tuned tree | Polynomial d=2 |
|---|---|---|---|---|
| Test RMSE (mg/m³) | **{ln_['test_rmse']:.3f}** | {t5['test_rmse']:.3f} | {tt_['test_rmse']:.3f} | {pl_['test_rmse']:.3f} |
| Test R² | **{ln_['test_r2']:.3f}** | {t5['test_r2']:.3f} | {tt_['test_r2']:.3f} | {pl_['test_r2']:.3f} |
| CV RMSE (week-grouped) | **{ln_['cv_rmse']:.3f}** | {t5['cv_rmse']:.3f} | {tt_['cv_rmse']:.3f} | {pl_['cv_rmse']:.3f} |
| Train-test RMSE gap | **{ln_['rmse_gap (test-train)']:.3f}** | {t5['rmse_gap (test-train)']:.3f} | {tt_['rmse_gap (test-train)']:.3f} |
{pl_['rmse_gap (test-train)']:.3f} |
| Parameters / leaves | 9 | {int(t5['n_leaves'])} | {int(tt_['n_leaves'])} | 45 |

- **How strong is the evidence?* The best tree is {best_tree_gap:.3f} RMSE worse than linear (95 % week-bootstrap CI {d5['95% CI low']:+.3f} to
{d5['95% CI high']:+.3f}, linear better in {d5['P(tree worse)']:.0%} of resamples), and linear wins {wins_all} of 5 CV folds against it; against all
other tree depths the gap is larger and its CI excludes zero. So the preference for linear is consistent but, for the best tree, small.

- **Why linear is better here.**
  1. *The signal is largely a smooth, additive dose-response.* C0 rises smoothly and almost linearly with `PT08.S2`, with extra additive contributions
from other sensors and temperature. A line represents this with 9 numbers; a tree must approximate it with a staircase of constant steps (the depth-3
tree is an 8-step function of one sensor), which loses precision *inside* every step.
  2. *Trees add variables up inefficiently.* With `PT08.S2` alone the tuned tree is actually better than the line (CV RMSE {sf.loc['Best single sensor
(PT08.S2)', 'Tree (tuned)']:.3f} vs. {sf.loc['Best single sensor (PT08.S2)', 'Linear']:.3f}), but adding the other seven features improves the linear
model by {sf.loc['relative gain from adding the other 7 features', 'Linear']:.0%} and the tree by only {sf.loc['relative gain from adding the other 7
features', 'Tree (tuned)']:.0%}, because each split uses a single feature while the truth combines several correlated sensors.
  3. *Bias-variance.* Shallow trees under-fit (depth 3: RMSE {RESULTS[tree_names[3]]['test_rmse']:.3f}); deeper trees fit the noise (depth 10:
{RESULTS[tree_names[10]]['test_rmse']:.3f}; unrestricted: {RESULTS[tree_names[None]]['test_rmse']:.3f}, training R² = 1.00). The linear model has low
variance (gap {ln_['rmse_gap (test-train)']:.3f}) and, given the near-linear signal, little bias.
  4. *Where trees do win.* The tree beats the line only in the lowest sensor decile (RMSE {DECILE.loc['D1', 'Tree tuned']:.2f} vs. {DECILE.loc['D1',
'Linear']:.2f}), where the relation flattens, and on abrupt shapes (synthetic demonstration). They need no scaling and give readable rules.

- **Linear vs. non-linear patterns.** The residuals of the linear model show a slight U-shaped (non-linear) pattern, and a quadratic term gives a real
gain (test RMSE {ln_['test_rmse']:.3f} → {pl_['test_rmse']:.3f}); a tree fitted to the linear residuals also helps ({HYBRID['test_rmse']:.3f}). The
non-linearity is therefore real but **smooth and mild** – it is captured best by the degree-2 polynomial, second best by a linear backbone with tree
corrections, and least efficiently by a stand-alone staircase.

- **Where both fail.** All models under-predict the rare very-high C0 episodes (bias {reg.loc[(N_LIN_NE, vh), 'bias (pred-true)']:.2f} for linear,
{reg.loc[(tree_names[5], vh), 'bias (pred-true)']:.2f} for the depth-5 tree, in mg/m³) and have errors that grow with the C0 level.
""")
```

Answer. Linear Regression performed better overall, although the margin over the best-sized tree is modest. The degree-2 polynomial regression was the best model of all.

	Linear	Best Decision Tree (depth 5)	Tuned tree	Polynomial d=2
Test RMSE (mg/m ³)	0.493	0.513	0.516	0.447
Test R ²	0.880	0.870	0.869	0.901
CV RMSE (week-grouped)	0.488	0.509	0.513	0.453
Train-test RMSE gap	0.013	0.073	0.046	0.036
Parameters / leaves	9	32	29	45

- **How strong is the evidence?** The best tree is 0.020 RMSE worse than linear (95 % week-bootstrap CI -0.001 to +0.039, linear better in 97% of resamples), and linear wins 5 of 5 CV folds against it; against all other tree depths the gap is larger and its CI excludes zero. So the preference for linear is consistent but, for the best tree, small.
- **Why linear is better here.**
 1. *The signal is largely a smooth, additive dose-response.* CO rises smoothly and almost linearly with PT08.S2, with extra additive contributions from other sensors and temperature. A line represents this with 9 numbers; a tree must approximate it with a staircase of constant steps (the depth-3 tree is an 8-step function of one sensor), which loses precision *inside* every step.
 2. *Trees add variables up inefficiently.* With PT08.S2 alone the tuned tree is actually better than the line (CV RMSE 0.534 vs. 0.575), but adding the other seven features improves the linear model by 15% and the tree by only 4%, because each split uses a single feature while the truth combines several correlated sensors.
 3. *Bias-variance.* Shallow trees under-fit (depth 3: RMSE 0.585); deeper trees fit the noise (depth 10: 0.544; unrestricted: 0.610, training R² = 1.00). The linear model has low variance (gap 0.013) and, given the near-linear signal, little bias.
 4. *Where trees do win.* The tree beats the line only in the lowest sensor decile (RMSE 0.22 vs. 0.30), where the relation flattens, and on abrupt shapes (synthetic demonstration). They need no scaling and give readable rules.
- **Linear vs. non-linear patterns.** The residuals of the linear model show a slight U-shaped (non-linear) pattern, and a quadratic term gives a real gain (test RMSE 0.493 → 0.447); a tree fitted to the linear residuals also helps (0.470). The non-linearity is therefore real but **smooth and mild** — it is captured best by the degree-2 polynomial, second best by a linear backbone with tree corrections, and least efficiently by a stand-alone staircase.
- **Where both fail.** All models under-predict the rare very-high CO episodes (bias -0.93 for linear, -0.66 for the depth-5 tree, in mg/m³) and have errors that grow with the CO level.

Limitations and suggested next steps

Limitations

- **Optimistic test protocol.** The assigned 80/20 split is random; hourly samples are strongly autocorrelated, so the test scores are somewhat optimistic. This was mitigated by week-grouped CV for every decision, week-clustered bootstrap intervals, and a chronological robustness check (which showed no material degradation in RMSE — but it is a single split).
- **One site, ~13 months, one sensor array.** Seasonal coverage is a single cycle; the PT08.S4 sensor drifts over time, so a deployed model would need recalibration.
- **Differences among the top models are small relative to uncertainty** (95 % intervals of $\approx \pm 0.08$ RMSE; paired comparisons are more informative than individual intervals).
- **Interpretability under collinearity.** Individual coefficients and single-feature importances of correlated sensors are not reliable; the polynomial model is also fragile on unrealistic feature combinations (Stage 4).
- **Explanations marked "plausible, not tested"** in the notebook (e.g. the log-target result) are hypotheses, not findings.

Next steps (outside the brief) — add hour-of-day / day-of-week features (the residuals show a systematic evening under-prediction and morning over-prediction); try ensembles (random forest, gradient boosting) and quantile / heteroscedastic regression for the extreme episodes; use time-aware nested CV for a forecasting-style evaluation; monitor sensor drift.

Deliverables checklist

```
CHECK = pd.DataFrame([
    ("Data preprocessing and handling of missing values", "$1.2 – $1.12", "-200 sentinel → NaN, outage rows dropped, target never imputed, standardisation fitted on train only"),
    ("Load the data; target and feature selection; 80/20 split", "$1.1, $1.5, $1.12", "CO(GT); 8 sensor + weather features; stratified 80/20 split (seed 42)"),
    ("Linear and polynomial regression from scratch + degree choice + sklearn cross-check", "$2.2 – $2.3, $2.5", "NumPy-only; degree 2 chosen by week-grouped CV with the 1-SE rule; identical to sklearn"),
    ("Decision Tree training with different depths (3, 5, 10, None)", "$2.4", "plus sweep 1–20, CV tuning, pruning"),
    ("MSE, RMSE, R² for all models", "$2.5, $3.1", "train and test, plus MAE and CV; bootstrap CIs in §3.2"),
    ("Performance metrics vs. epoch (regression)", "$3.3", "linear (GD) and polynomial (Adam)"),
    ("Performance metric vs. tree depth", "$3.4", "MSE, RMSE, R² for train / CV / test"),
    ("Decision tree plot for a small depth (max_depth = 3)", "$4.1", "thresholds in sensor units, rules and leaf table"),
    ("Feature importances (which variable contributes most?)", "$4.2 – $4.3", "four methods + grouped permutation"),
    ("Baseline comparison: Linear vs. Decision Tree; linear vs. non-linear patterns", "$5.1 – $5.6", "paired statistics, synthetic demo, ablation ledger, scorecard"),
    ("Short discussion (important features; Tree vs. Linear and why)", "$6", "answers Q1 and Q2 above"),
    ("Neat, commented code", "all", "functions documented, fixed seeds, verified against sklearn"),
], columns=["Requirement", "Where", "Notes"]).set_index("Requirement")
display(CHECK.style.set_properties(**{"text-align": "left"}))

import sys, platform
print(f"Reproducibility: Python {platform.python_version()} | numpy {np.__version__} | pandas {pd.__version__} | scikit-learn {sklearn.__version__} | random seed {SEED}")
```

	Requirement	Where	Notes
	Data preprocessing and handling of missing values	\$1.2 – \$1.12	–200 sentinel → NaN, outage rows dropped, target never imputed, standardisation fitted on train only
	Load the data; target and feature selection; 80/20 split	\$1.1, \$1.5, \$1.12	CO(GT); 8 sensor + weather features; stratified 80/20 split (seed 42)
Linear and polynomial regression from scratch + degree choice + sklearn cross-check		\$2.2 – \$2.3, \$2.5	NumPy-only; degree 2 chosen by week-grouped CV with the 1-SE rule; identical to sklearn
	Decision Tree training with different depths (3, 5, 10, None)	\$2.4	plus sweep 1–20, CV tuning, pruning
	MSE, RMSE, R² for all models	\$2.5, \$3.1	train and test, plus MAE and CV; bootstrap CIs in §3.2
	Performance metrics vs. epoch (regression)	\$3.3	linear (GD) and polynomial (Adam)
	Performance metric vs. tree depth	\$3.4	MSE, RMSE, R² for train / CV / test
	Decision tree plot for a small depth (max_depth = 3)	\$4.1	thresholds in sensor units, rules and leaf table
	Feature importances (which variable contributes most?)	\$4.2 – \$4.3	four methods + grouped permutation
Baseline comparison: Linear vs. Decision Tree; linear vs. non-linear patterns		\$5.1 – \$5.6	paired statistics, synthetic demo, ablation ledger, scorecard
	Short discussion (important features; Tree vs. Linear and why)	\$6	answers Q1 and Q2 above
	Neat, commented code	all	functions documented, fixed seeds, verified against sklearn

Reproducibility: Python 3.11.9 | numpy 2.4.6 | pandas 3.0.6 | scikit-learn 1.9.1 | random seed 42

 **All stages complete.** The notebook runs top to bottom from the raw CSV; the PDF report in `docs/report/` is generated from it.